

# From Perceptron to Deep Learning

*AI Workshop, 2-4 December 2019*

Somnuk Phon-Amnuaisuk  
School of Computing and Informatics  
Centre for Innovative Engineering  
Universiti Teknologi Brunei



# meme

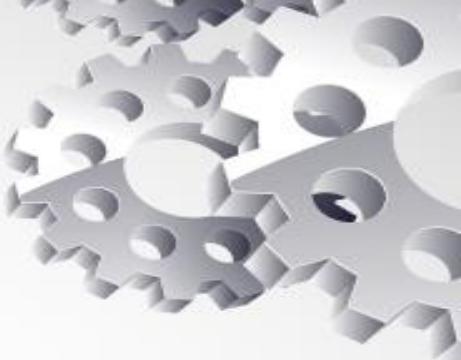
A meme (mēm), a neologism coined by Richard Dawkins, is "an idea, behavior, or style that spreads from person to person within a culture". A meme acts as a unit for carrying cultural ideas, symbols, or practices that can be transmitted from one mind to another through writing, speech, gestures, rituals, or other imitable phenomena with a mimicked theme.



# Disclaimer

This lecture is compiled from my lectures as well as materials gathering from lectures found in public domains.

# Outline



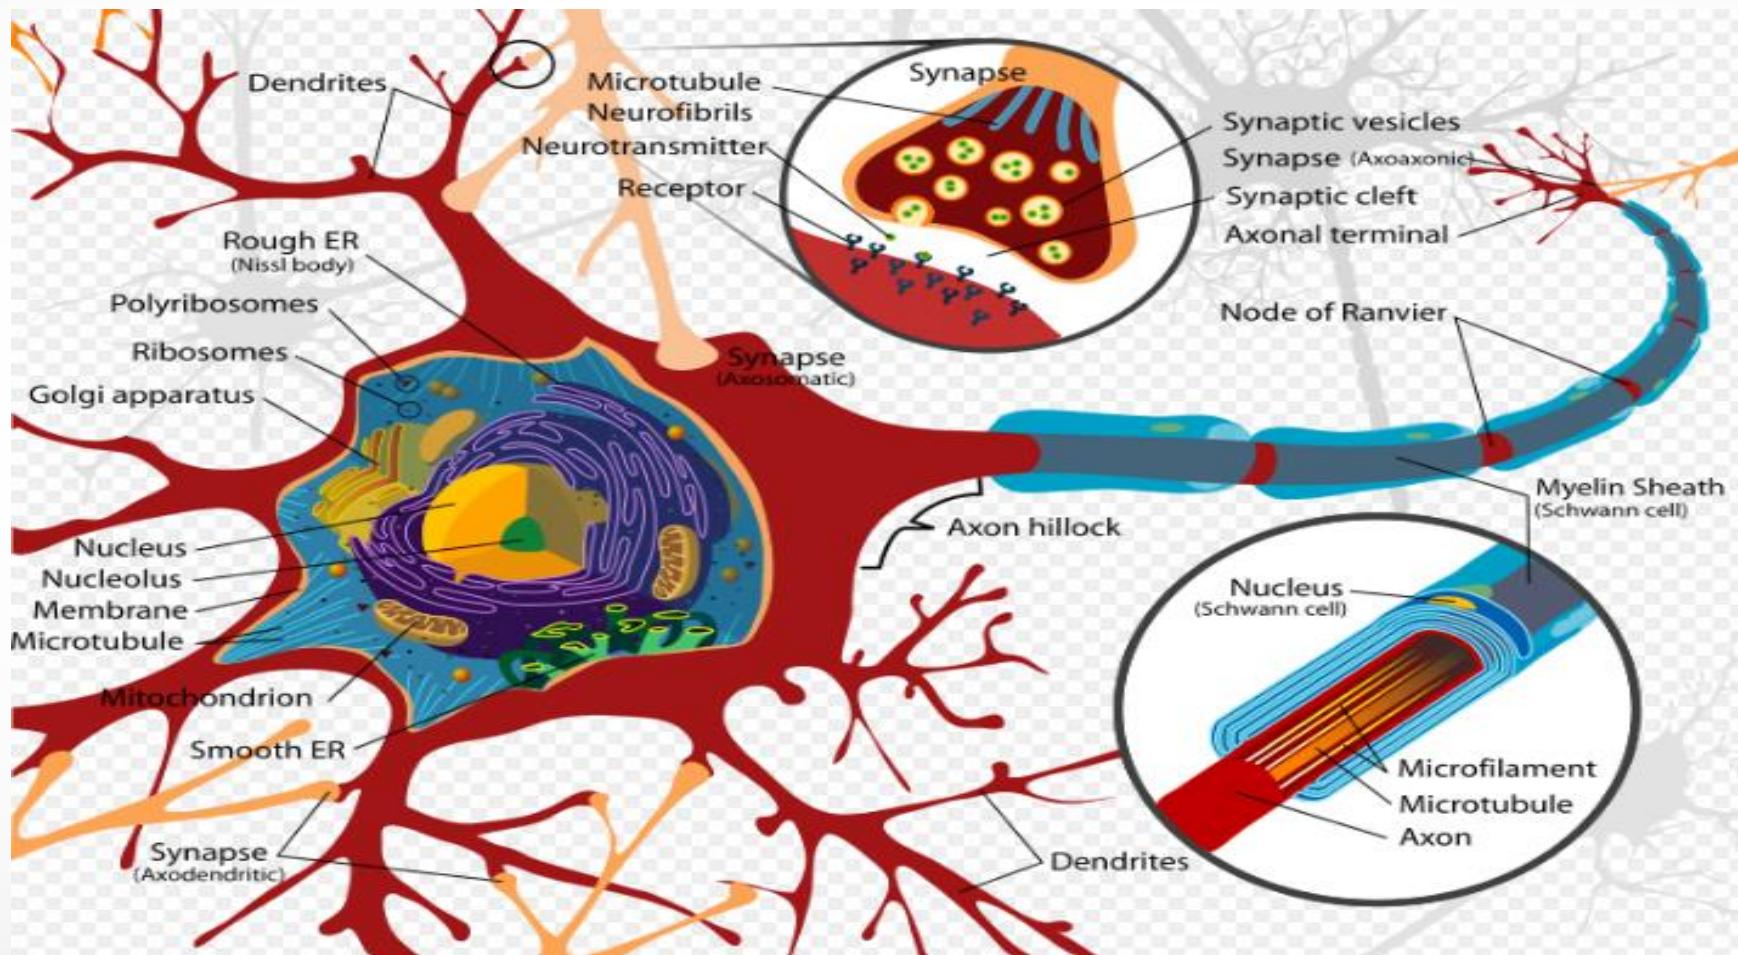
- Theoretical Foundation
  - Perceptron
  - Multi-Layer Perceptron
  - Deep Learning: CNN Based Computer Vision



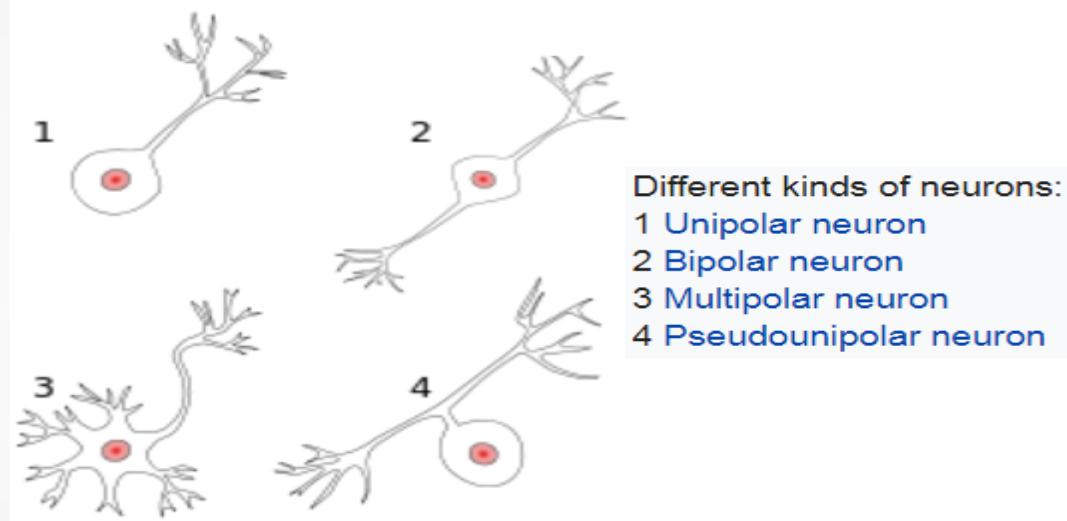
# Perceptron

AI workshop. Pre - Coding Conquest 2019 event

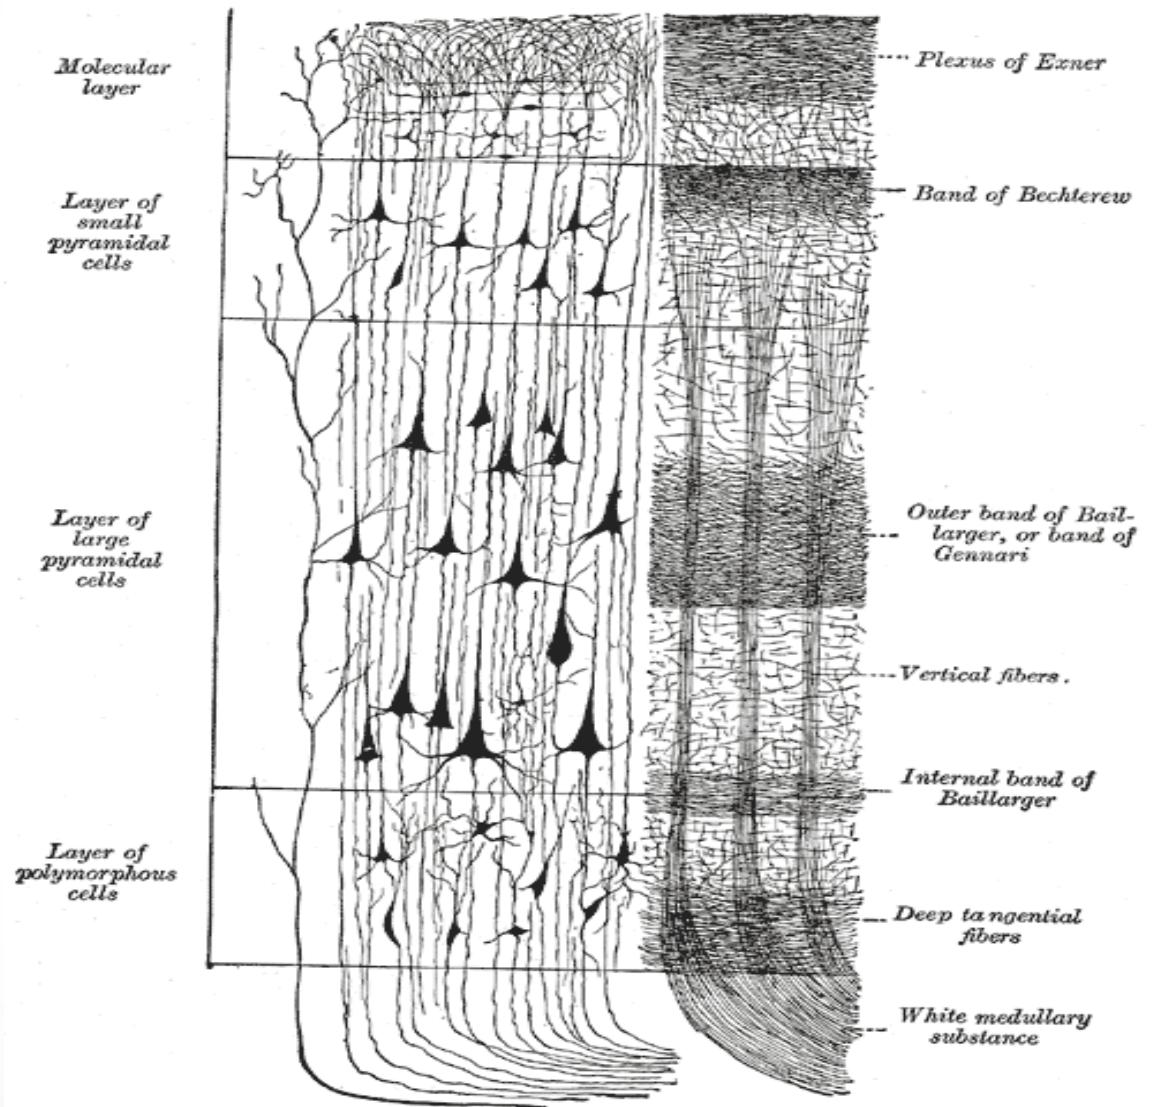
# Neurons



# Neurons

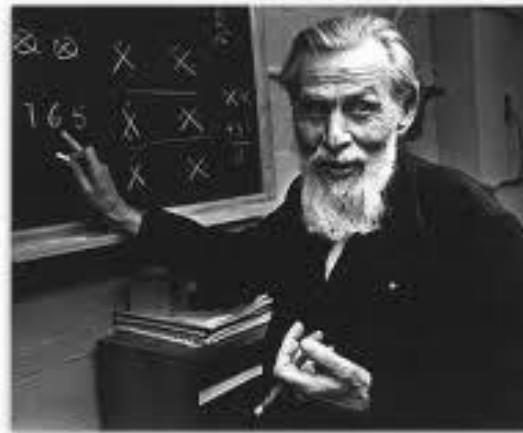
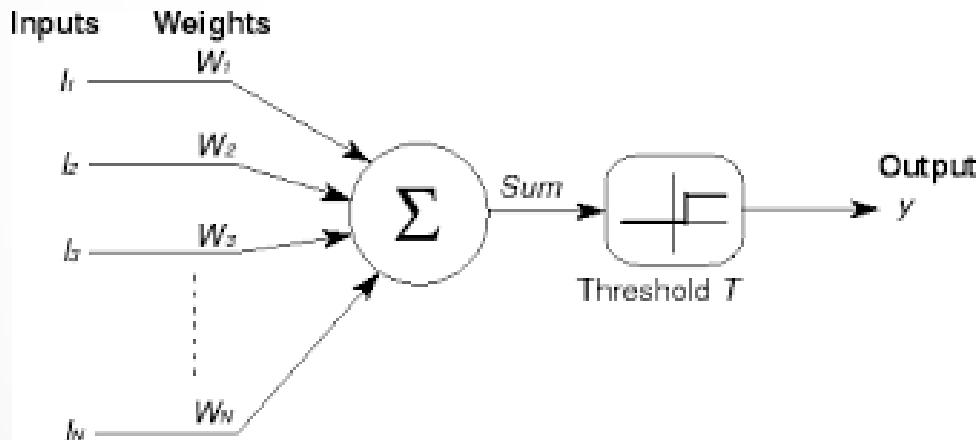


wikipedia

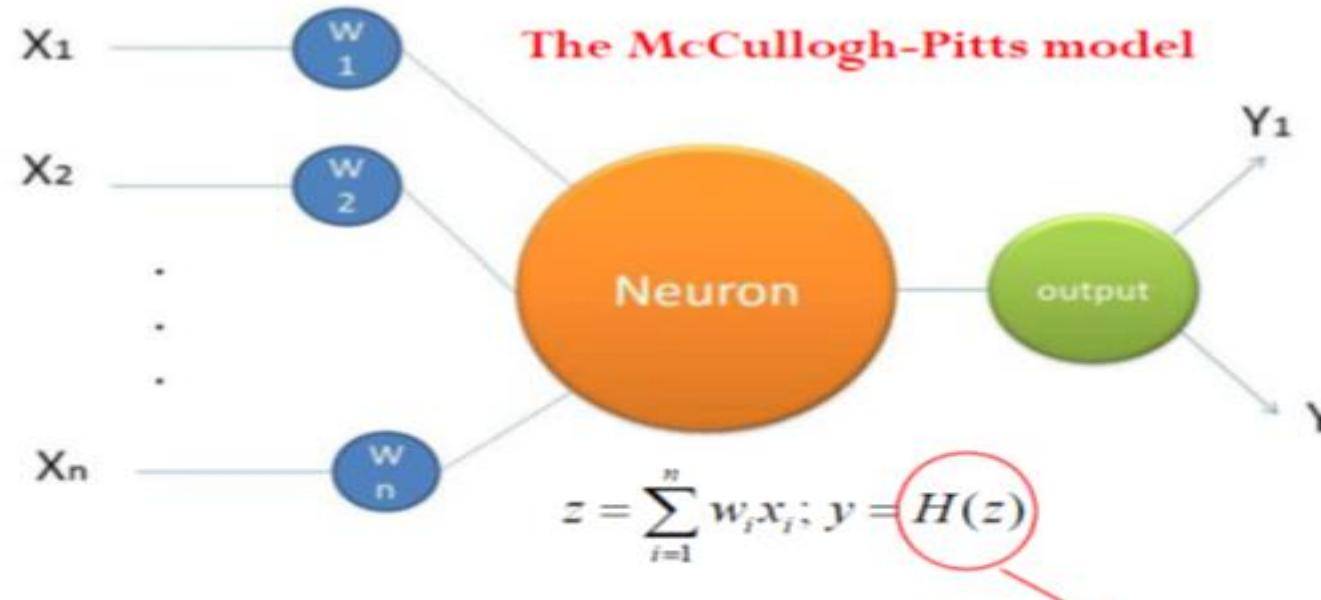


# Threshold Logic Unit

- The first artificial neuron was the Threshold Logic Unit (TLU), or Linear Threshold Unit, first proposed by Warren McCulloch and Walter Pitts in 1943.



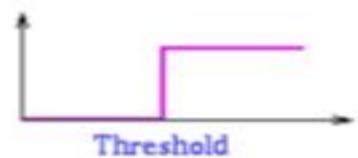
# Perceptrons



Wires : axon & dendrites

Connection weights: Synapse

Threshold function: activity in soma



# Implementing Perceptrons



- How to solve for W  $f(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + b > 0 \\ 0 & \text{otherwise} \end{cases}$

$\mathbf{w} \cdot \mathbf{x}$  is  $\sum_{i=1}^m w_i x_i$

- Randomly, manually
- Analytically
- Optimization algorithm e.g., gradient descent

# Perceptrons



- The perceptron is an algorithm for learning a classifier function.

$$f(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

$\mathbf{w} \cdot \mathbf{x}$  is  $\sum_{i=1}^m w_i x_i$

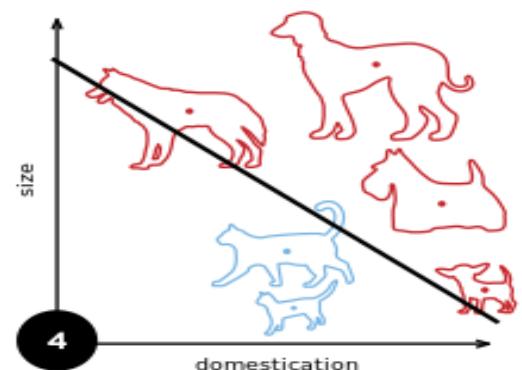
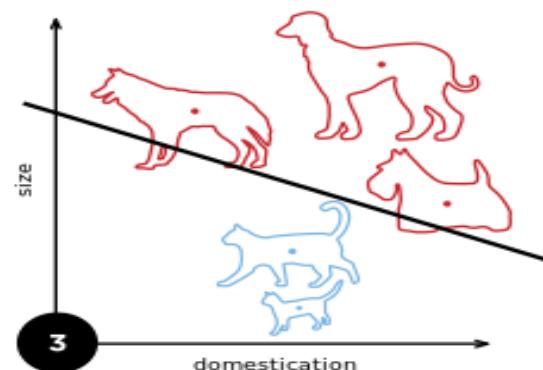
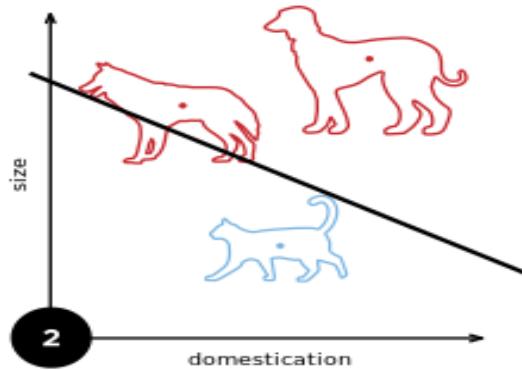
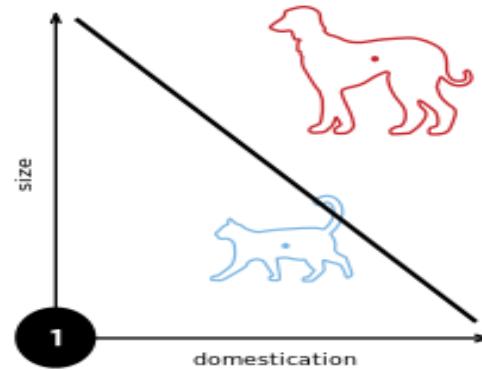
- The perceptron is an artificial neuron using the Heaviside step function as the activation function.

# Perceptron Learning Rule



- Rosenblatt: Perceptron

$$w_i(t + 1) = w_i(t) + \alpha(d_j - y_j(t))x_{j,i}$$

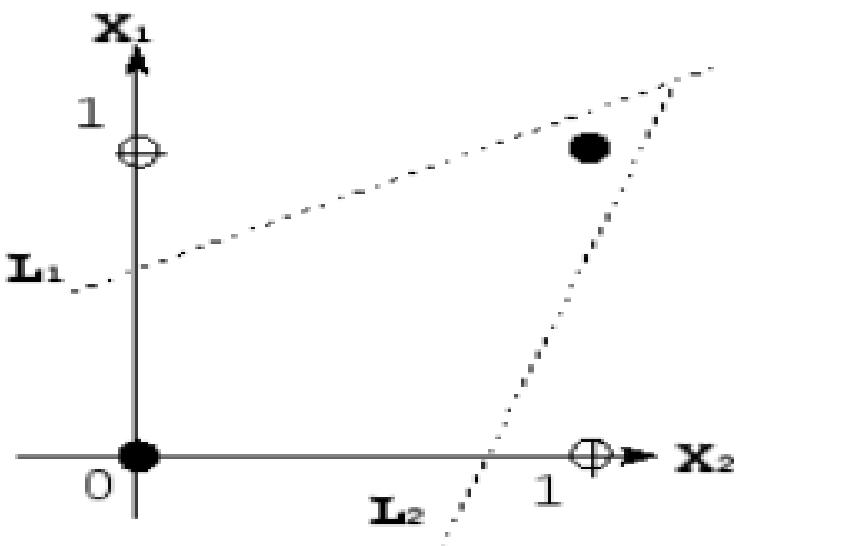


# Limitation of Perceptrons



- The perceptron can distinguish a linearly separable function. Hence a single perceptron cannot deal with an xor function.

$x_1$	$x_2$	$y$
0	0	0
0	1	1
1	0	1
1	1	0

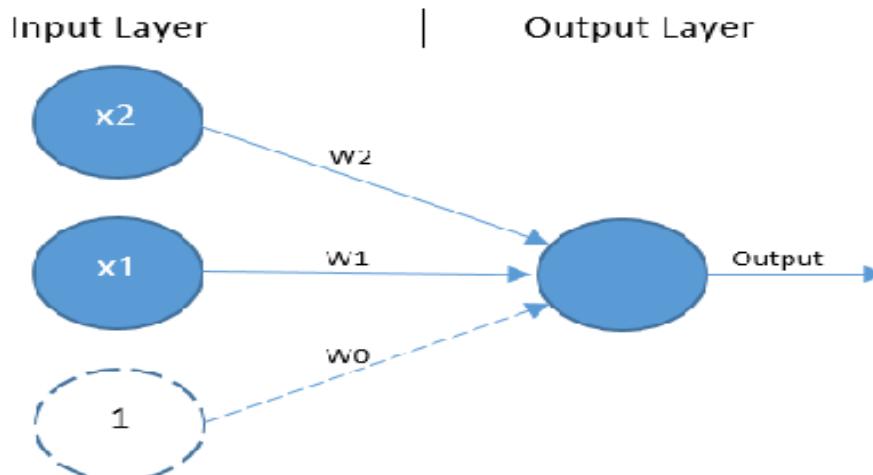
$$y = x_1 \oplus x_2$$


# Limitation of Perceptrons

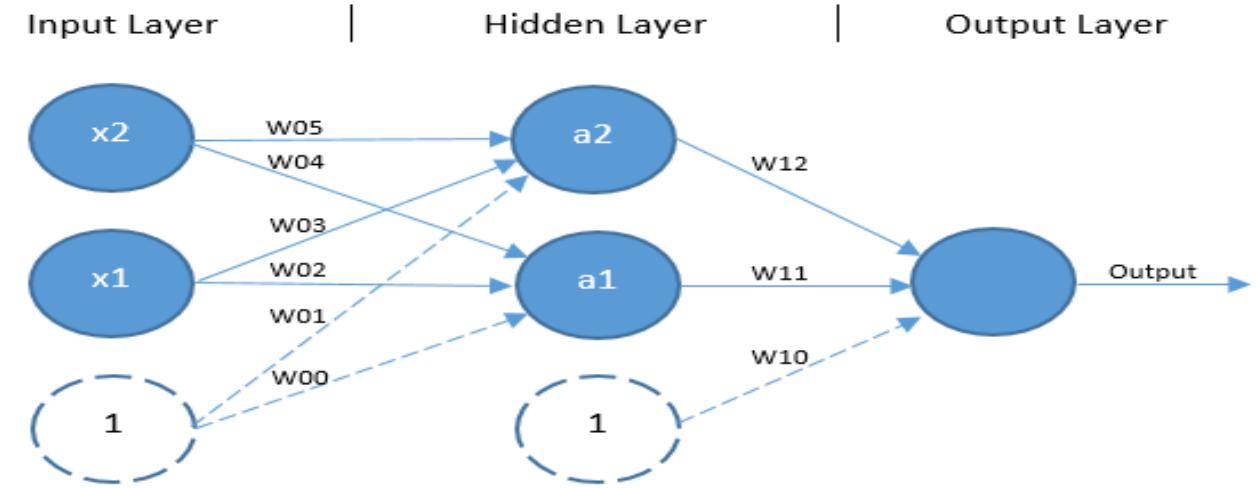


- Adding two perceptrons in a hidden layer, then it can handle an xor function.

perceptron



MLP

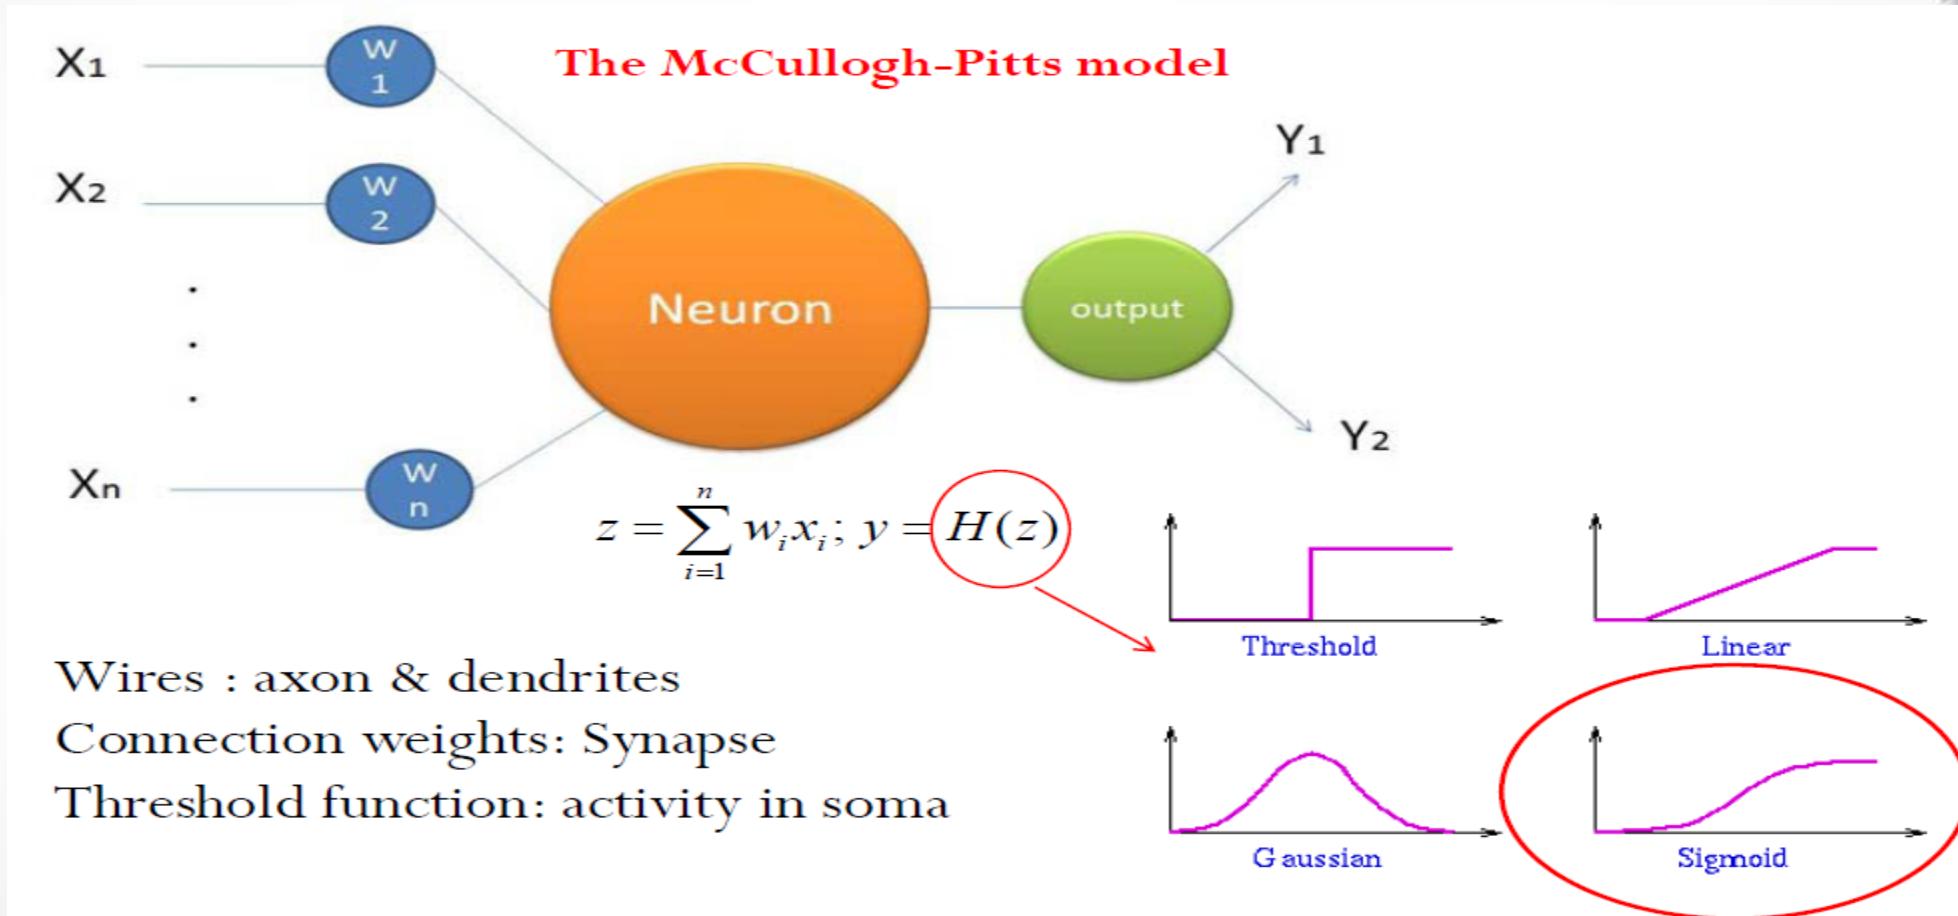
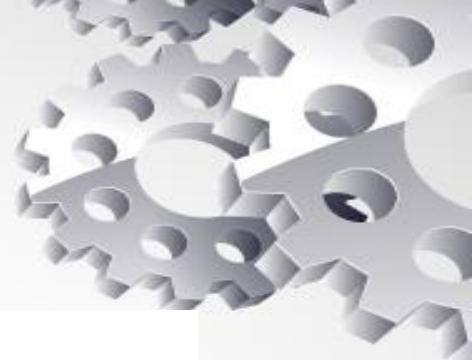




# Multi-Layer Perceptron

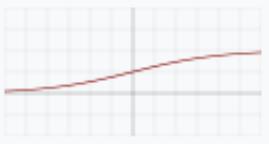
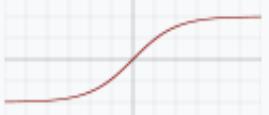
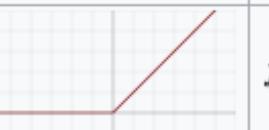
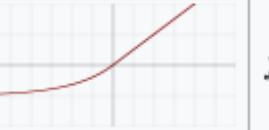
AI workshop. Pre - Coding Conquest 2019 event

# Activation Functions



# Activation Functions



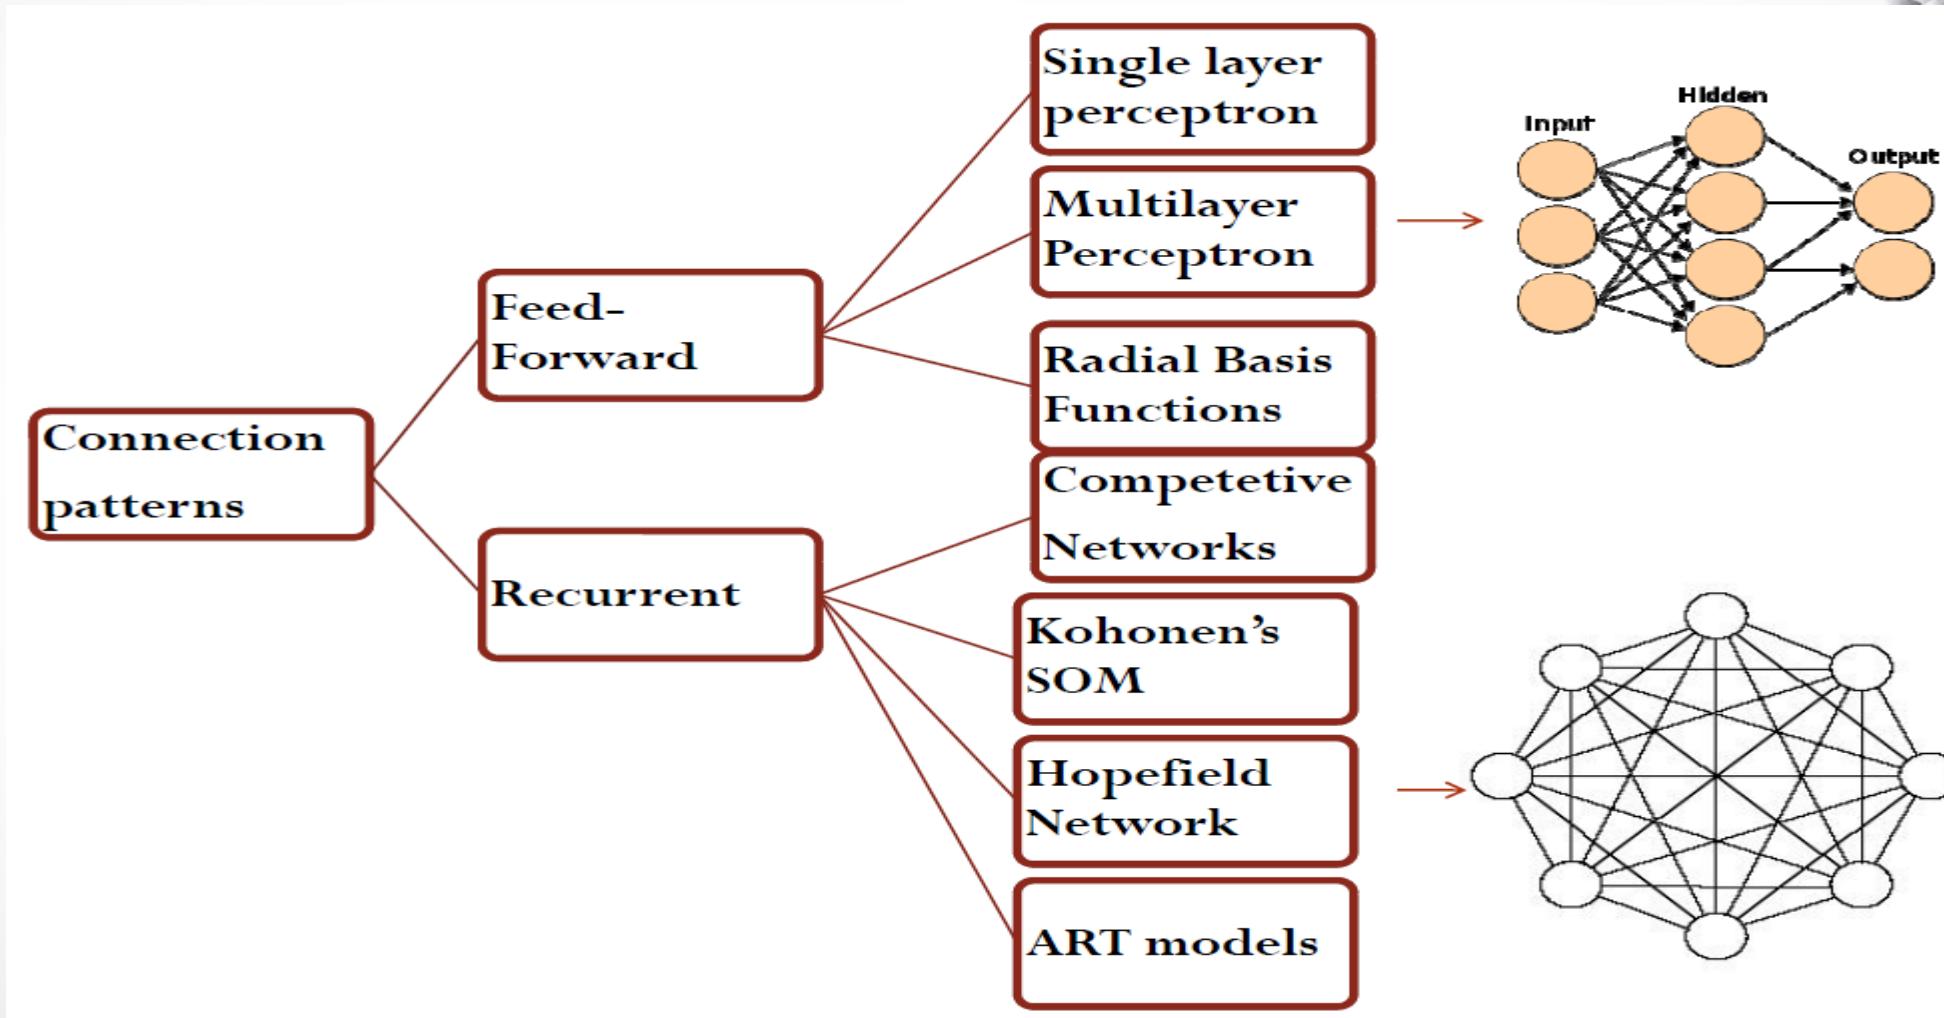
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a. Sigmoid or Soft step)		$f(x) = \sigma(x) = \frac{1}{1 + e^{-x}}^{[1]}$	$f'(x) = f(x)(1 - f(x))$
TanH		$f(x) = \tanh(x) = \frac{(e^x - e^{-x})}{(e^x + e^{-x})}$	$f'(x) = 1 - f(x)^2$
Rectified linear unit (ReLU) <sup>[15]</sup>		$f(x) = \begin{cases} 0 & \text{for } x \leq 0 \\ x & \text{for } x > 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \leq 0 \\ 1 & \text{for } x > 0 \end{cases}$
Exponential linear unit (ELU) <sup>[20]</sup>		$f(\alpha, x) = \begin{cases} \alpha(e^x - 1) & \text{for } x \leq 0 \\ x & \text{for } x > 0 \end{cases}$	$f'(\alpha, x) = \begin{cases} f(\alpha, x) + \alpha & \text{for } x \leq 0 \\ 1 & \text{for } x > 0 \end{cases}$

# Artificial Neural Networks

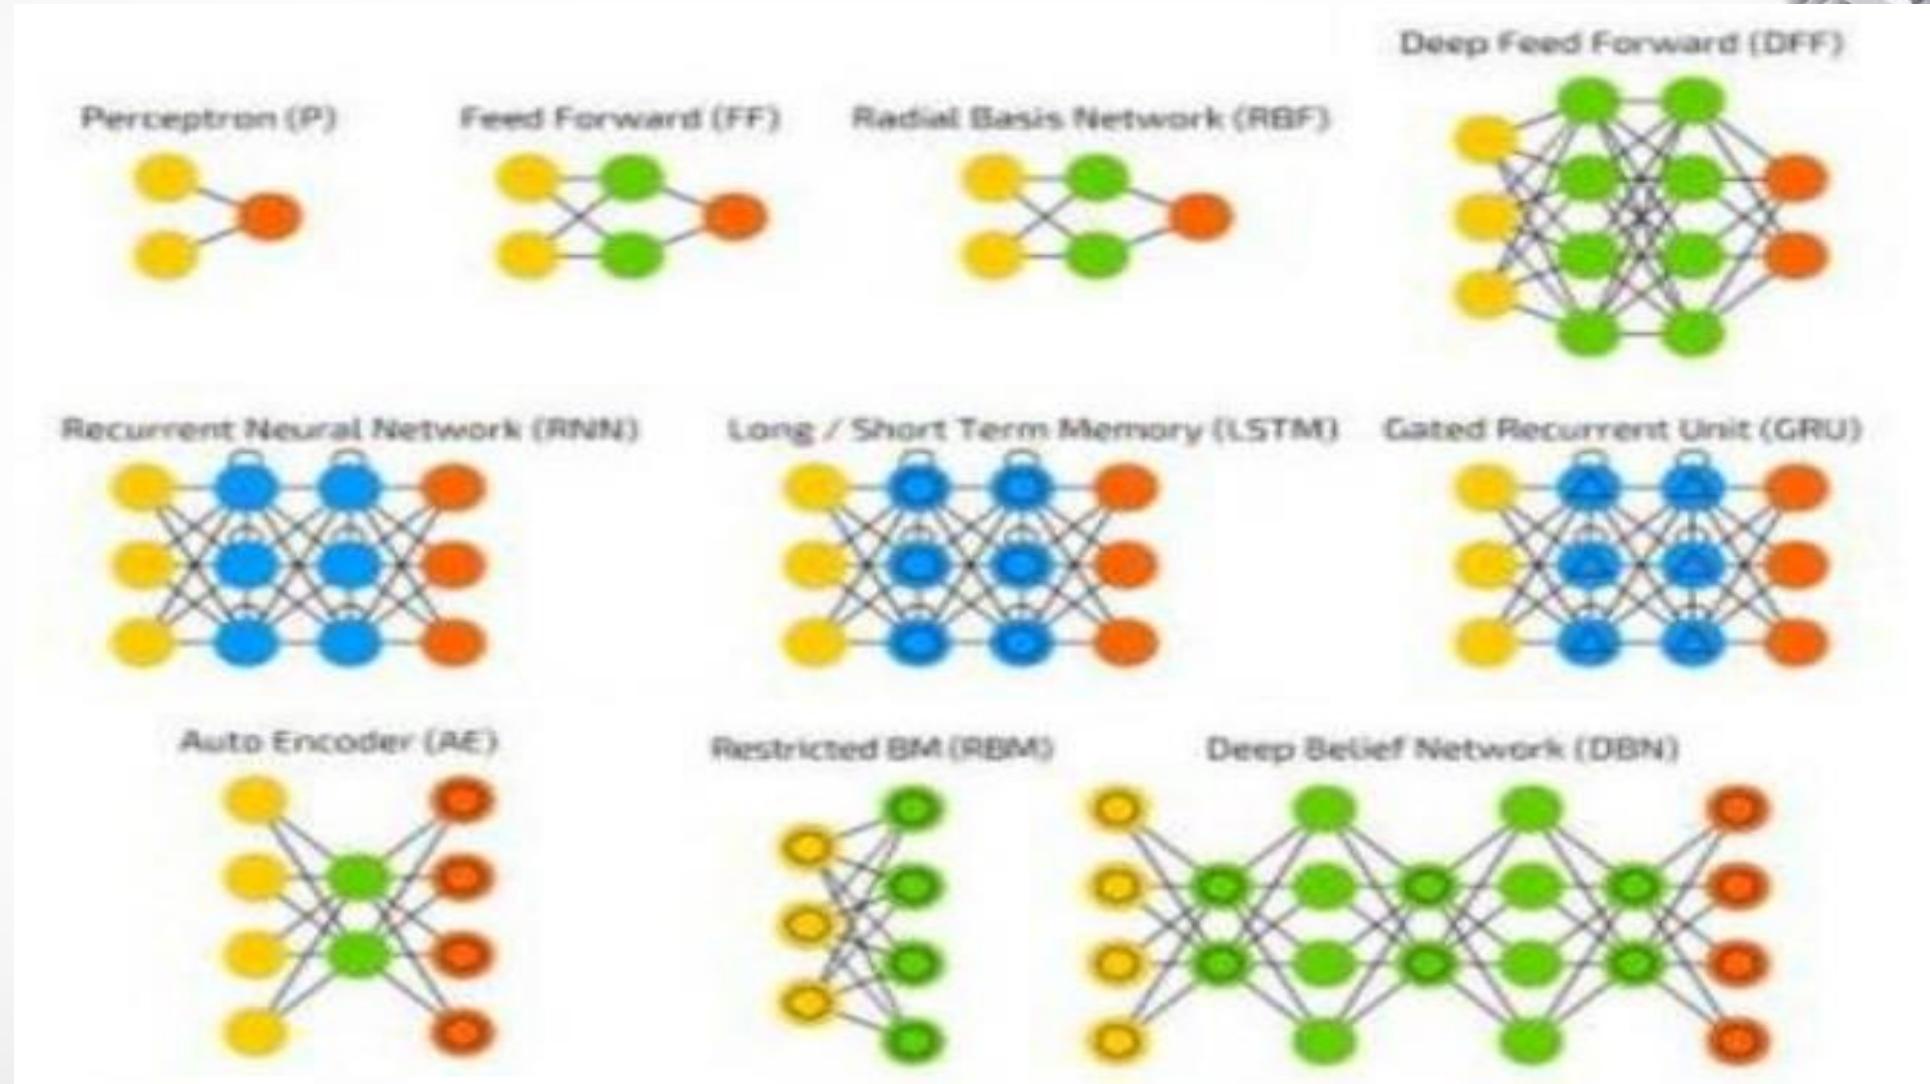


- 1940: McCulloch and Pitts
  - A logical calculus of the ideas immanent in nervous activity
- 1957: Frank Rosenblatt
  - The perceptron learning algorithm, convergence theorem
- 1960: Minsky and Papert
  - Perceptrons cannot deal with XOR
- 1980: Werbos and Rumelhart
  - Back-propagation learning algorithm
- 1980: Hopfield
  - Hopfield's energy approach (recurrent neural network)

# Artificial Neural Networks

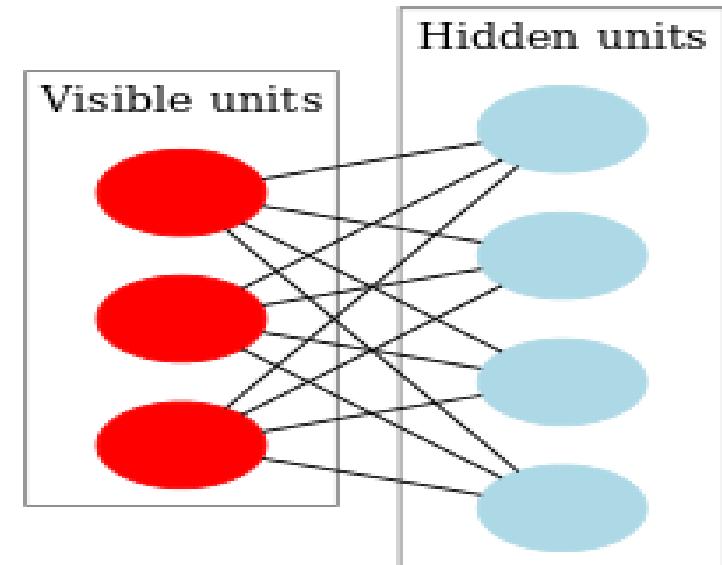
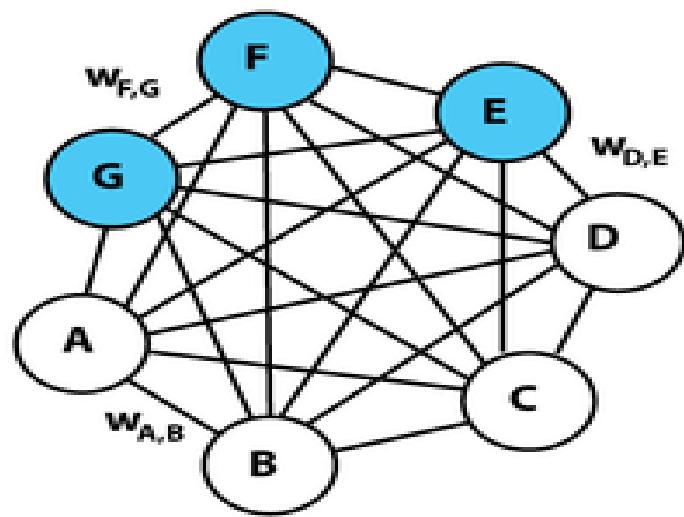


# Artificial Neural Networks

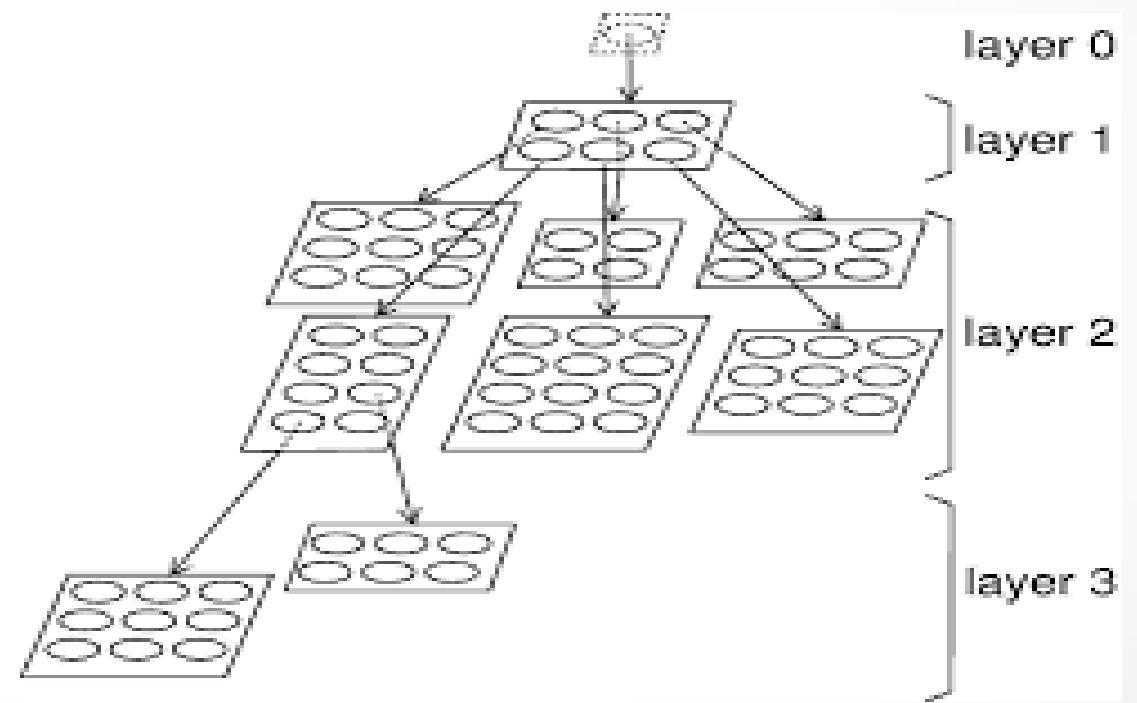
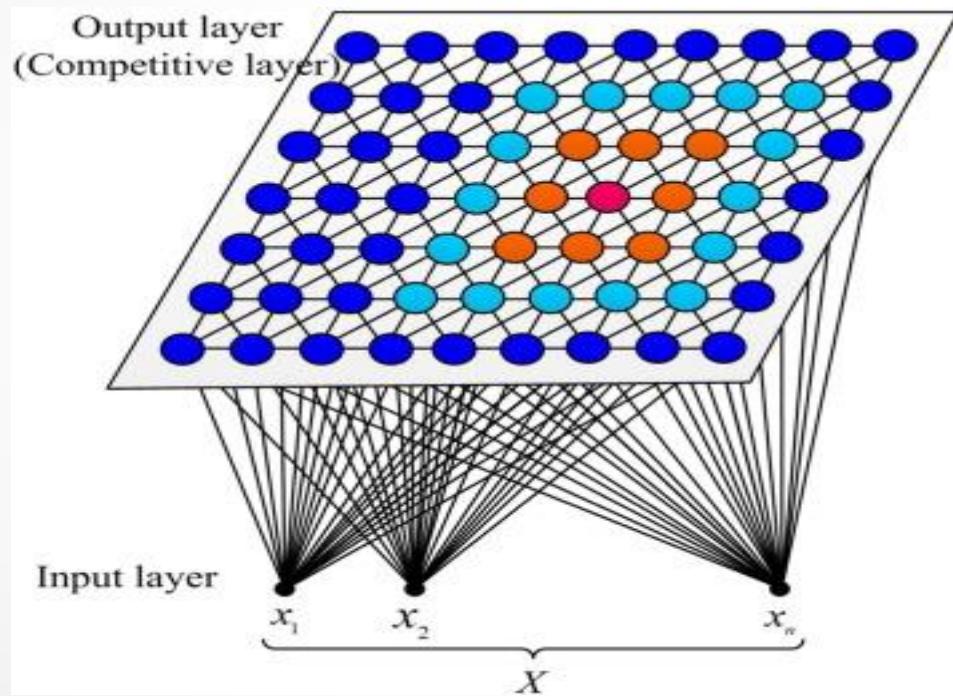


Credit: pinterest

# Artificial Neural Networks

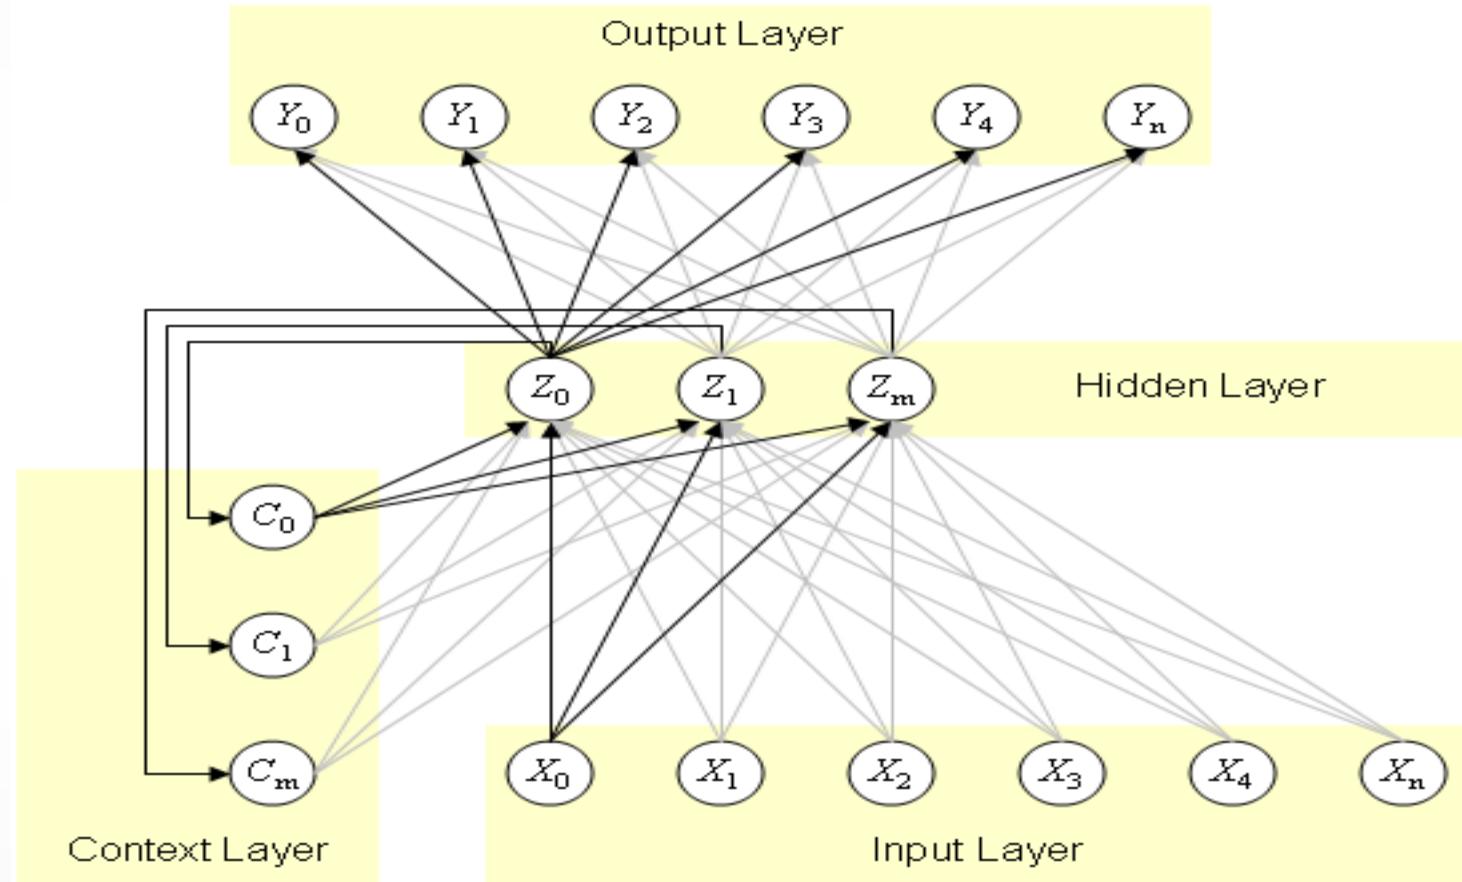


# Artificial Neural Networks



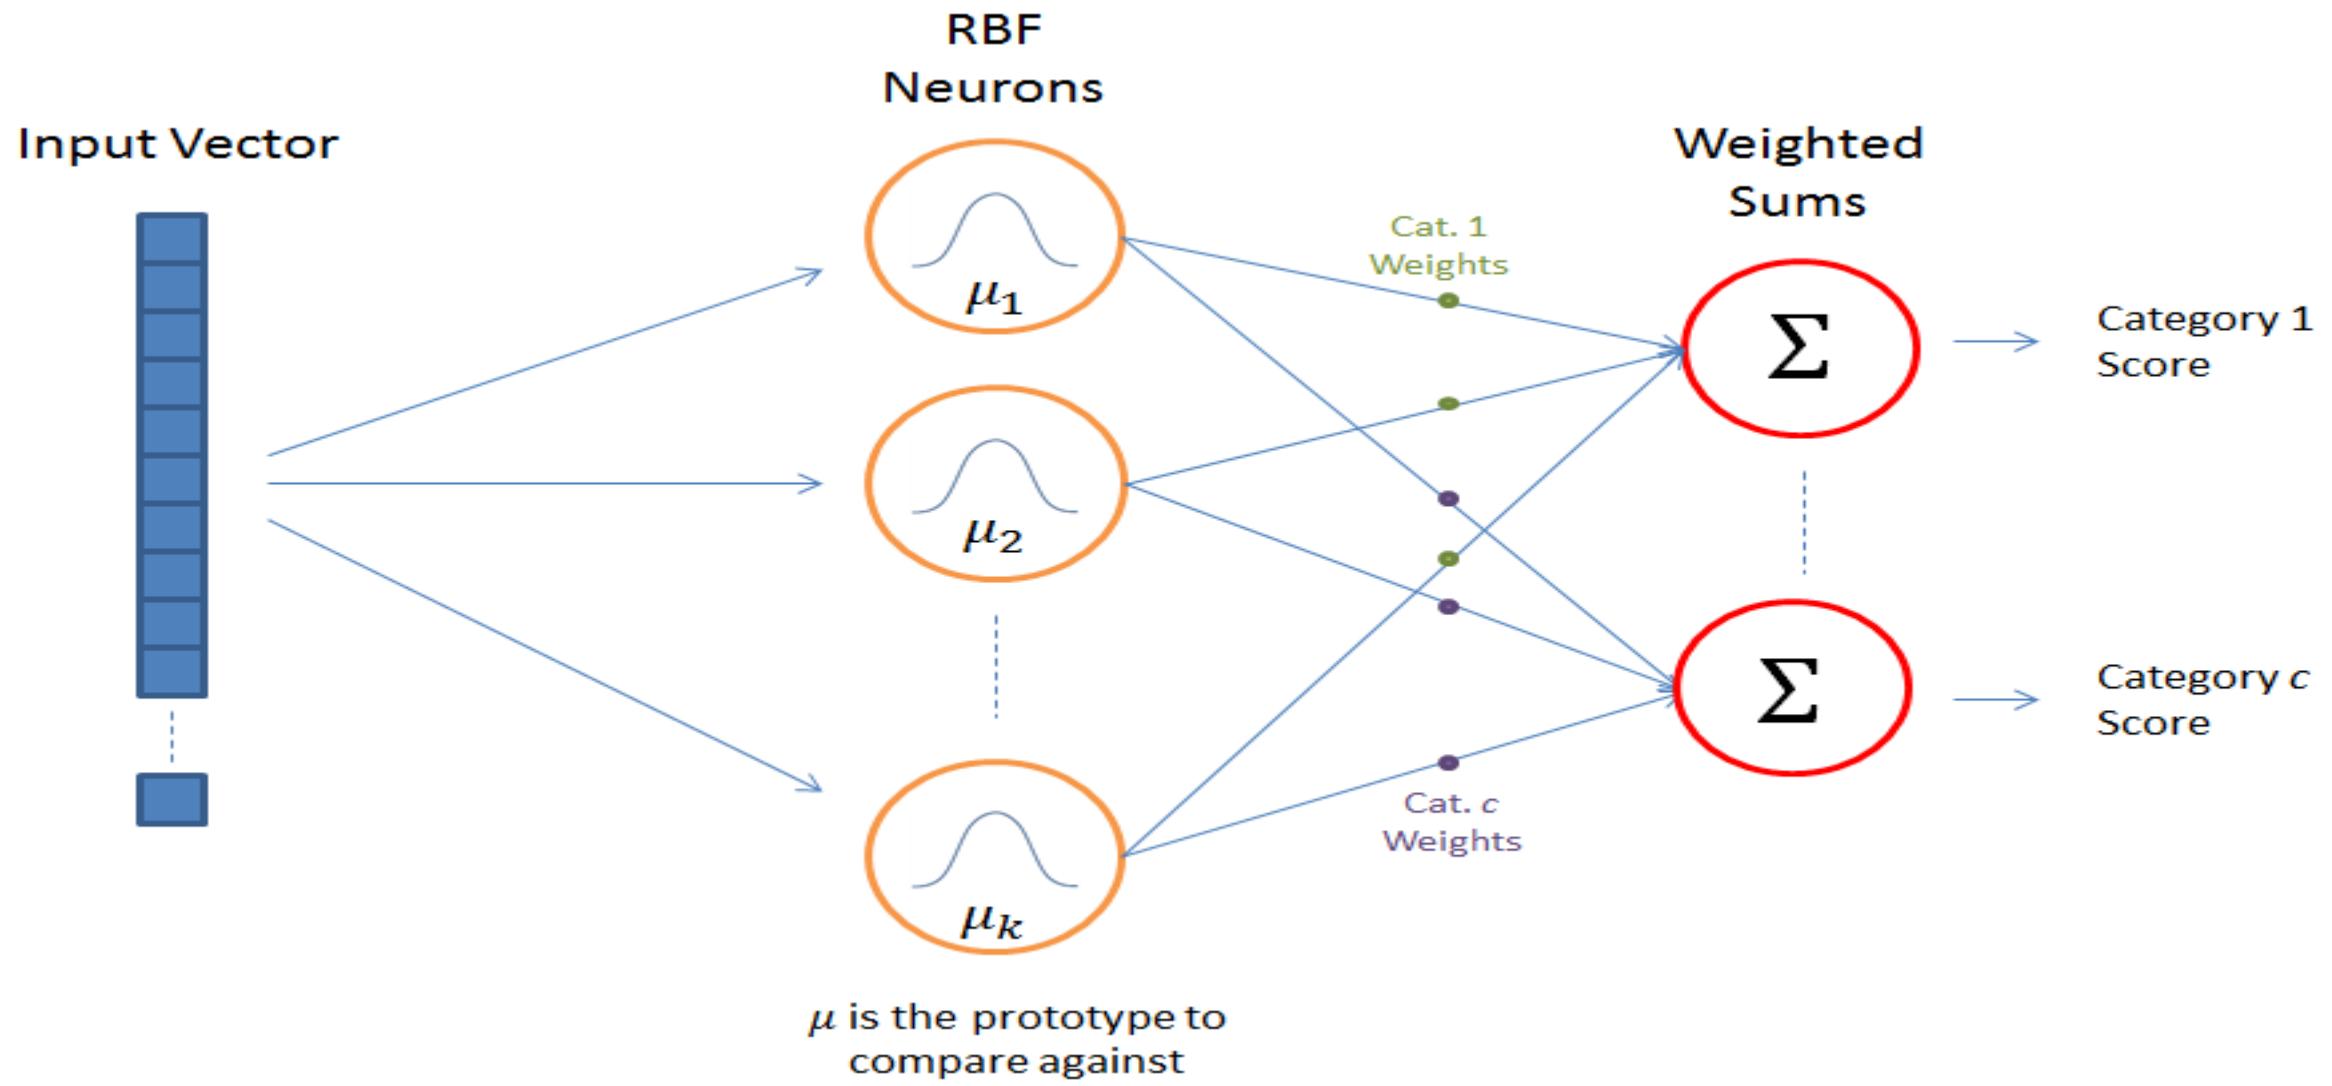
Credit: google

# Artificial Neural Networks



Credit: google

# Artificial Neural Networks

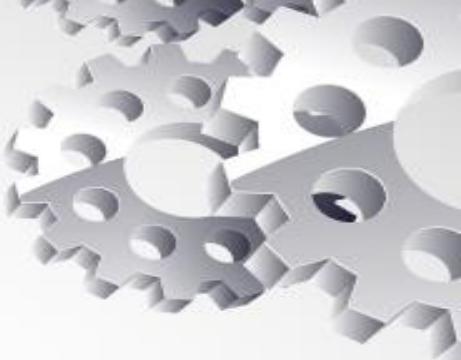


# Learning

- Learning Paradigms
  - Supervised
  - Unsupervised
  - Hybrid
- Learning Approaches
  - Hebbian rules: cells that fire together wire together
  - Competitive learning: increasing specialization of cells
  - Hopfield networks: associative memory
  - Error correction: backpropagation

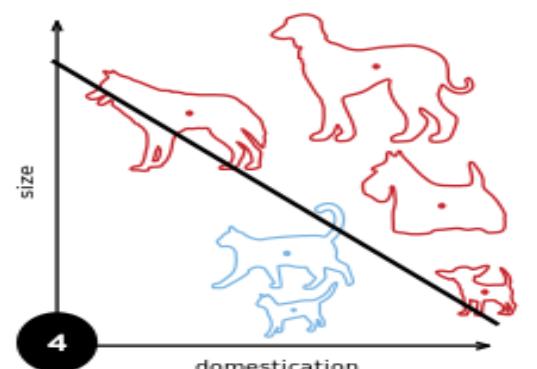
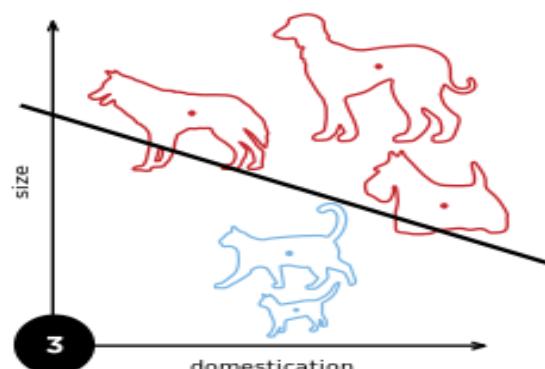
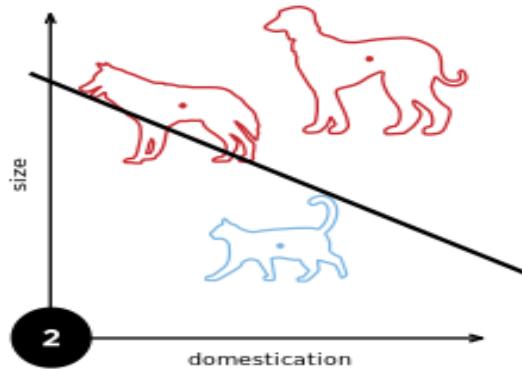
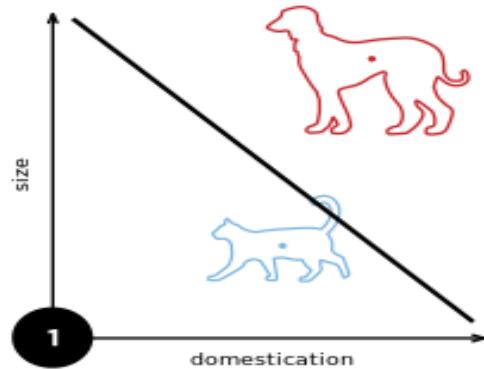


# Learning: Perceptron



- Rosenblatt: Perceptron

$$w_i(t + 1) = w_i(t) + \alpha(d_j - y_j(t))x_{j,i}$$



# Backpropagation (Rumelhart 1986)



$$E = \frac{1}{2}(t - y)^2 \quad \frac{\partial E}{\partial w_{ij}} = \frac{\partial E}{\partial o_j} \frac{\partial o_j}{\partial \text{net}_j} \frac{\partial \text{net}_j}{\partial w_{ij}}$$

$$o_j = \varphi(\text{net}_j) = \varphi \left( \sum_{i=1}^n w_{ij} o_i \right)$$

$$\varphi(z) = \frac{1}{1 + e^{-z}} \quad \frac{d\varphi}{dz}(z) = \varphi(z)(1 - \varphi(z))$$

$$\frac{\partial E}{\partial o_j} = \frac{\partial E}{\partial y} = \frac{\partial}{\partial y} \frac{1}{2}(t - y)^2 = y - t$$

$$\frac{\partial o_j}{\partial \text{net}_j} = \frac{\partial}{\partial \text{net}_j} \varphi(\text{net}_j) = \varphi(\text{net}_j)(1 - \varphi(\text{net}_j))$$

$$\frac{\partial \text{net}_j}{\partial w_{ij}} = \frac{\partial}{\partial w_{ij}} \left( \sum_{i=1}^n w_{ij} o_i \right) = \frac{\partial}{\partial w_{ij}} w_{ij} o_i = o_i$$

$$\frac{\partial E}{\partial w_{ij}} = (o_j - t) o_j (1 - o_j) o_i$$

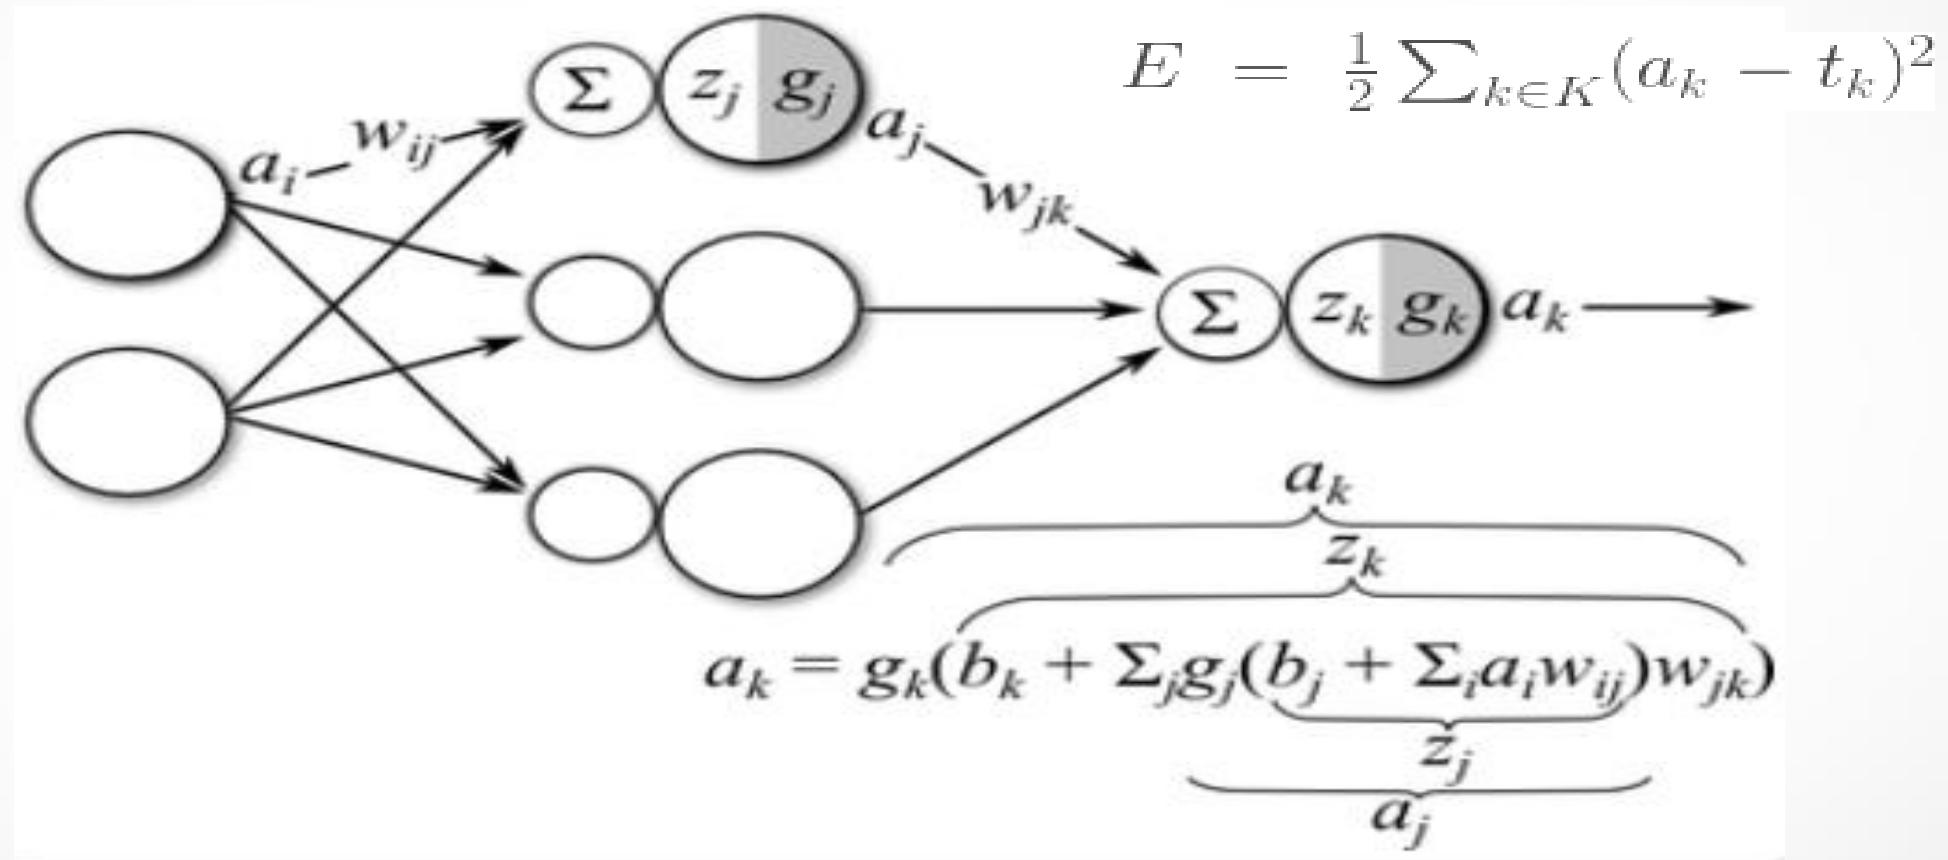


## Acknowledgement

The following backpropagation is adapted from

<https://theclevermachine.wordpress.com/2014/09/06/derivation-error-backpropagation-gradient-descent-for-neural-networks/>

# Backpropagation



# Backpropagation



$$E = \frac{1}{2} \sum_{k \in K} (a_k - t_k)^2$$

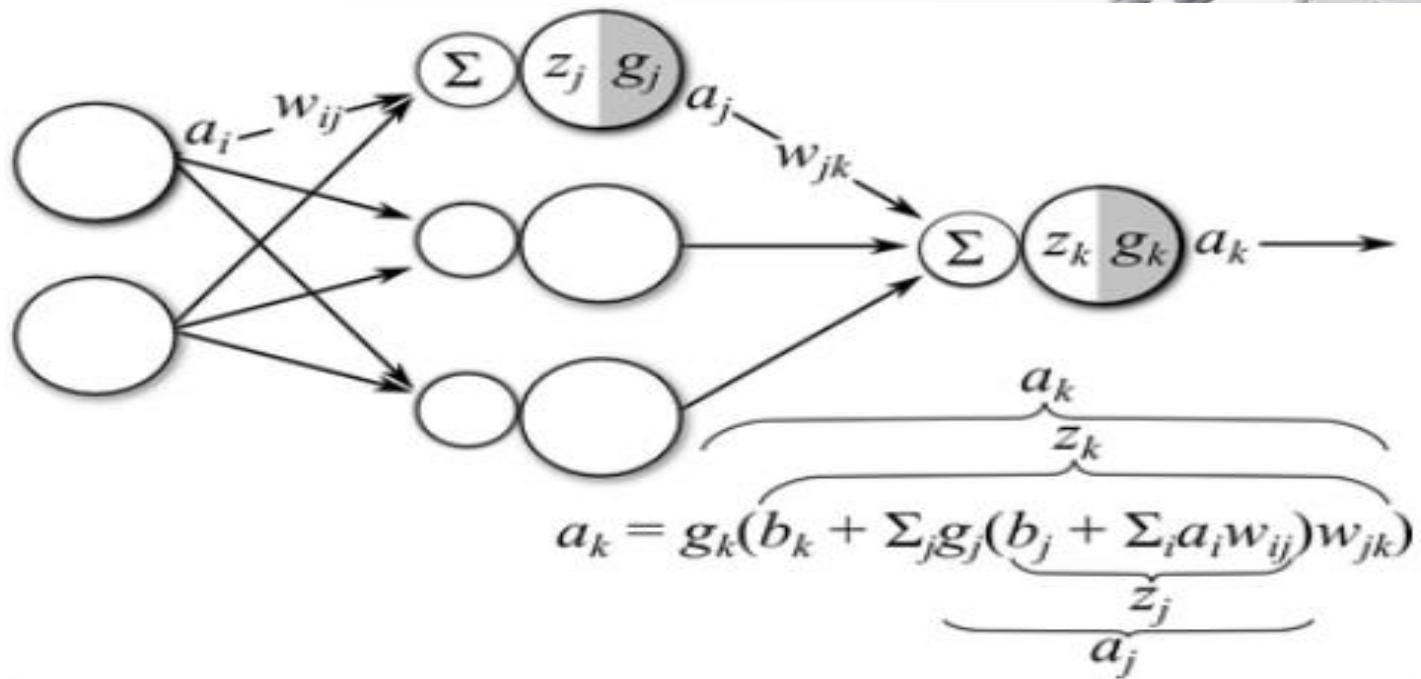
$$\begin{aligned}\frac{\partial E}{\partial w_{jk}} &= \frac{1}{2} \sum_{k \in K} (a_k - t_k)^2 \\ &= (a_k - t_k) \frac{\partial}{\partial w_{jk}} (a_k - t_k)\end{aligned}$$

$$\begin{aligned}\frac{\partial E}{\partial w_{jk}} &= (a_k - t_k) \frac{\partial}{\partial w_{jk}} a_k \\ &= (a_k - t_k) \frac{\partial}{\partial w_{jk}} g_k(z_k) \\ &= (a_k - t_k) g'_k(z_k) \frac{\partial}{\partial w_{jk}} z_k,\end{aligned}$$

$$z_k = b_j + \sum_j g_j(z_j) w_{jk}$$

$$\frac{\partial z_k}{\partial w_{j,k}} = g_j(z_j) = a_j$$

$$\frac{\partial E}{\partial w_{jk}} = (a_k - t_k)g'_k(z_k)a_j$$



$$\delta_k = (a_k - t_k)g'_k(z_k)$$

$$\frac{\partial E}{\partial w_{jk}} = \delta_k a_j$$

$$\begin{aligned}\frac{\partial E}{\partial b_k} &= (a_k - t_k)g'_k(z_k)(1) \\ &= \delta_k\end{aligned}$$

# Backpropagation

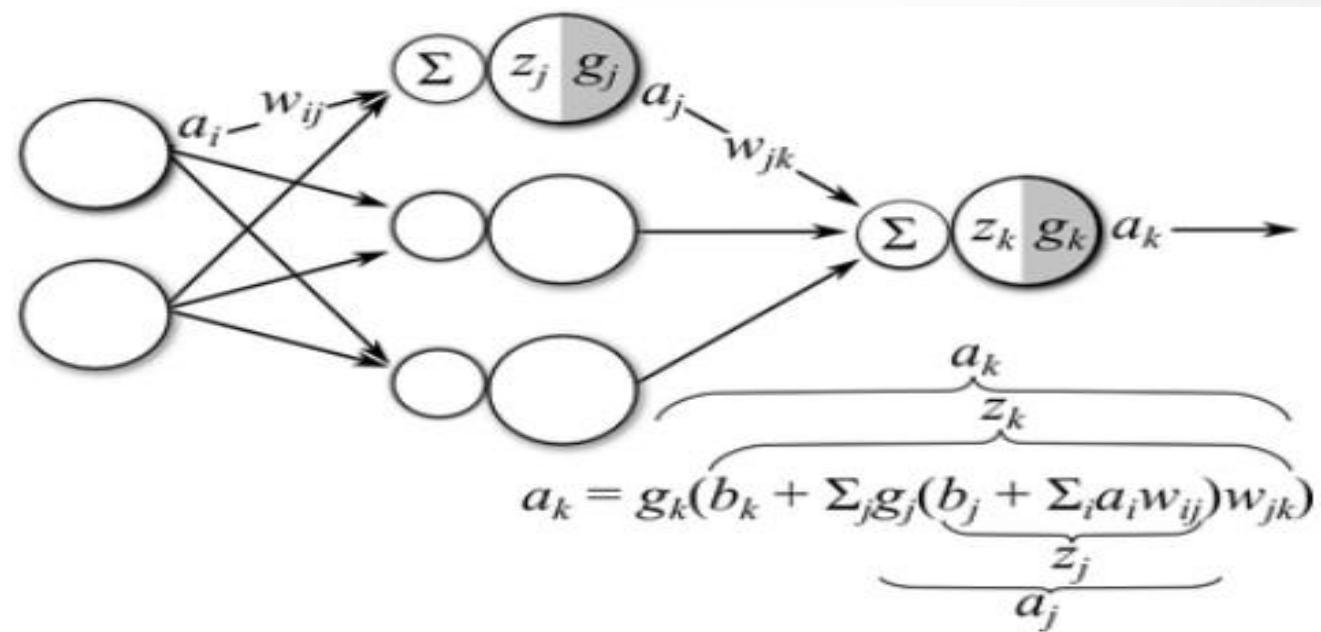


$$E = \frac{1}{2} \sum_{k \in K} (a_k - t_k)^2$$

$$\begin{aligned}\frac{\partial E}{\partial w_{ij}} &= \frac{1}{2} \sum_{k \in K} (a_k - t_k)^2 \\ &= \sum_{k \in K} (a_k - t_k) \frac{\partial}{\partial w_{ij}} a_k\end{aligned}$$

$$\begin{aligned}\frac{\partial E}{\partial w_{ij}} &= \sum_{k \in K} (a_k - t_k) \frac{\partial}{\partial w_{ij}} g_k(z_k) \\ &= \sum_{k \in K} (a_k - t_k) g'_k(z_k) \frac{\partial}{\partial w_{ij}} z_k\end{aligned}$$

$$\begin{aligned}z_k &= b_k + \sum_j a_j w_{jk} \\ &= b_k + \sum_j g_j(z_j) w_{jk} \\ &= b_k + \sum_j g_j(b_i + \sum_i z_i w_{ij}) w_{jk}\end{aligned}$$

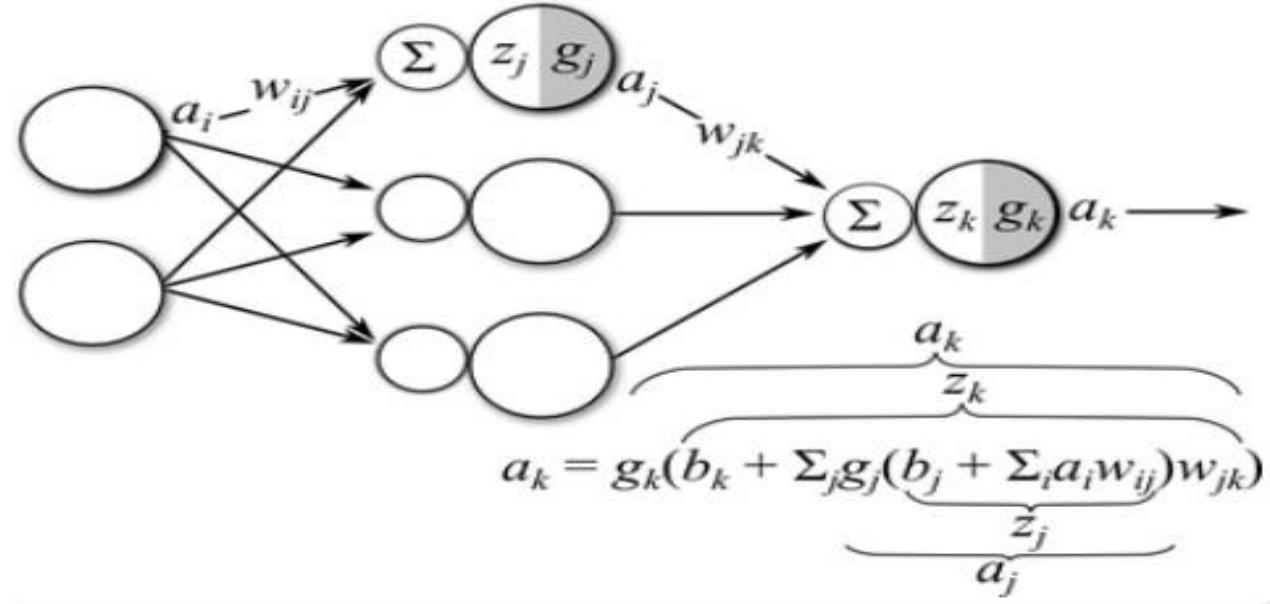


# Backpropagation



$$\begin{aligned}
 \frac{\partial z_k}{\partial w_{ij}} &= \frac{\partial z_k}{\partial a_j} \frac{\partial a_j}{\partial w_{ij}} \\
 &= \frac{\partial}{\partial a_j} a_j w_{jk} \frac{\partial a_j}{\partial w_{ij}} \\
 &= w_{jk} \frac{\partial a_j}{\partial w_{ij}} \\
 &= w_{jk} \frac{\partial g_j(z_j)}{\partial w_{ij}} \\
 &= w_{jk} g'_j(z_j) \frac{\partial z_j}{\partial w_{ij}} \\
 &= w_{jk} g'_j(z_j) \frac{\partial}{\partial w_{ij}} (b_i + \sum_i a_i w_{ij}) \\
 &= w_{jk} g'_j(z_j) a_i
 \end{aligned}$$

$$\begin{aligned}
 \frac{\partial E}{\partial w_{ij}} &= \sum_{k \in K} (a_k - t_k) g'_k(z_k) w_{jk} g'_j(z_j) a_i \\
 &= g'_j(z_j) a_i \sum_{k \in K} (a_k - t_k) g'_k(z_k) w_{jk} \\
 &= a_i g'_j(z_j) \sum_{k \in K} \delta_k w_{jk}
 \end{aligned}$$



$$\begin{aligned}
 \frac{\partial E}{\partial w_{ij}} &= a_i g'_j(z_j) \sum_{k \in K} \delta_k w_{jk} \\
 &= \delta_j a_i
 \end{aligned}$$

where

$$\delta_j = g'_j(z_j) \sum_{k \in K} \delta_k w_{jk}$$

# Backpropagation



$$\begin{aligned}\frac{\partial E}{\partial w_{ij}} &= a_i g'_j(z_j) \sum_{k \in K} \delta_k w_{jk} \\ &= \delta_j a_i\end{aligned}$$

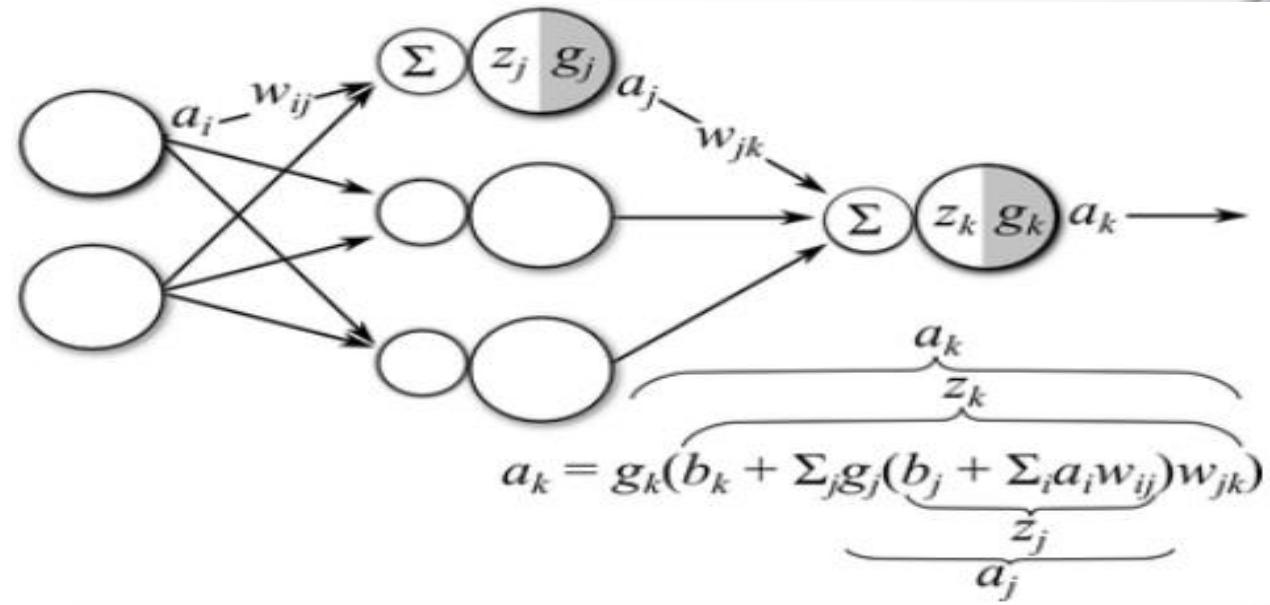
where

$$\delta_j = g'_j(z_j) \sum_{k \in K} \delta_k w_{jk}$$

$$\begin{aligned}\frac{\partial z_k}{\partial b_i} &= w_{jk} g'_j(z_j) \frac{\partial z_j}{\partial b_i} \\ &= w_{jk} g'_j(z_j) \frac{\partial}{\partial b_i} (b_i + \sum_i a_i w_{ij}) \\ &= w_{jk} g'_j(z_j)(1),\end{aligned}$$

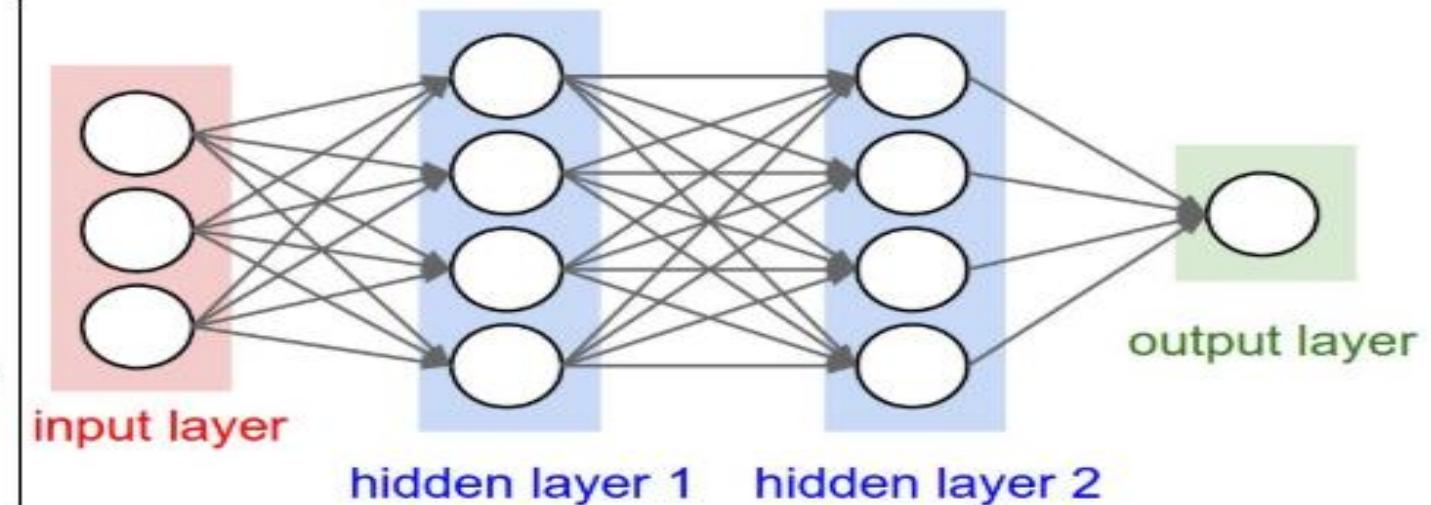
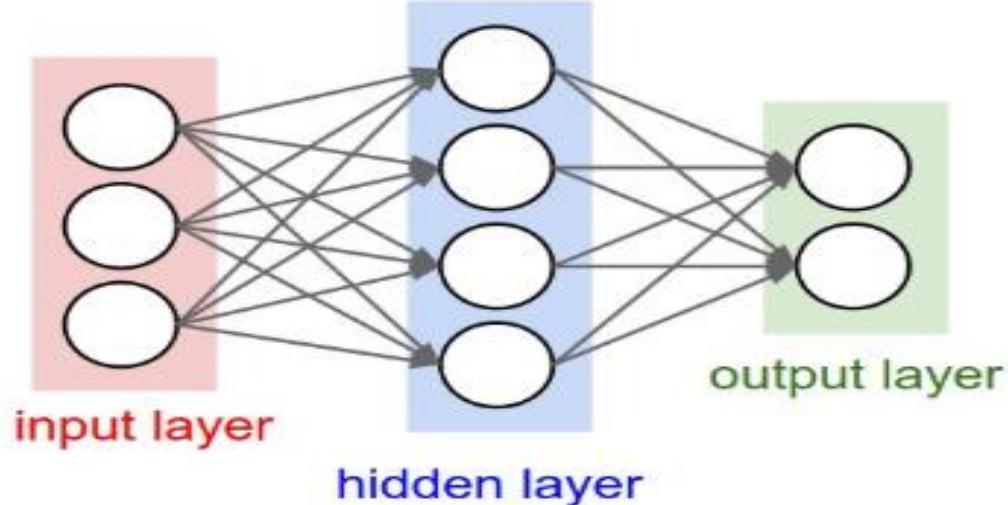
giving

$$\begin{aligned}\frac{\partial E}{\partial b_i} &= g'_j(z_j) \sum_{k \in K} \delta_k w_{jk} \\ &= \delta_j\end{aligned}$$

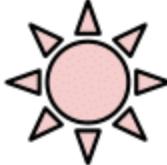
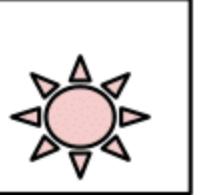
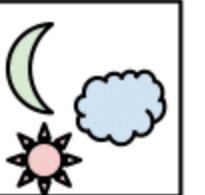


$$\begin{aligned}E &= \frac{1}{2} \sum_{k \in K} (a_k - t_k)^2 \\ \theta &\leftarrow \theta - \eta \frac{\partial E}{\partial \theta}\end{aligned}$$

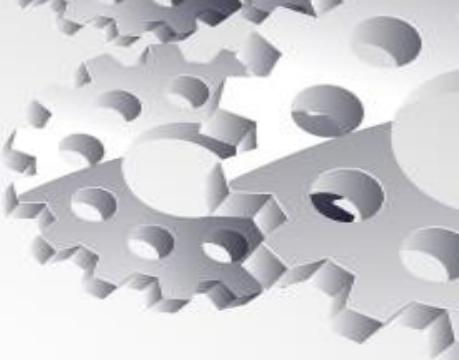
# Summary: Feedforward Neural Networks



# Summary: Multi-class / Multi-Label

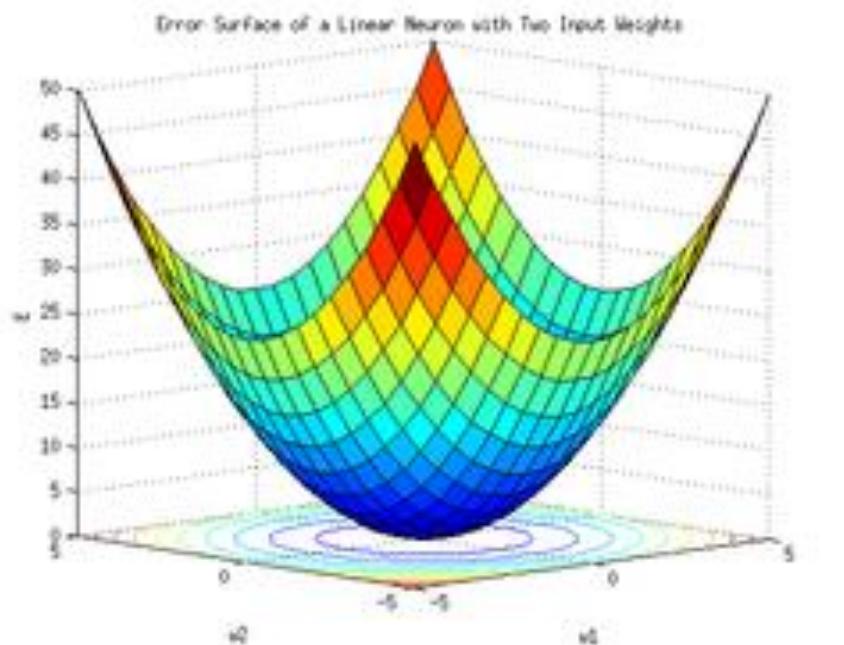
	Multi-Class	Multi-Label
$C = 3$	Samples	Samples
	  	  
	Labels (t)	Labels (t)
	$[0 \ 0 \ 1]$ $[1 \ 0 \ 0]$ $[0 \ 1 \ 0]$	$[1 \ 0 \ 1]$ $[0 \ 1 \ 0]$ $[1 \ 1 \ 1]$

# Square Error Loss

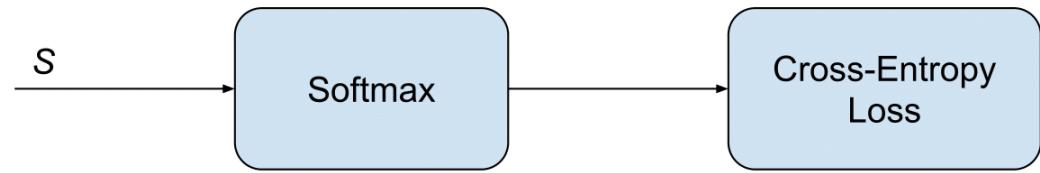


$$E_{\cdot} = (t - y)^2$$

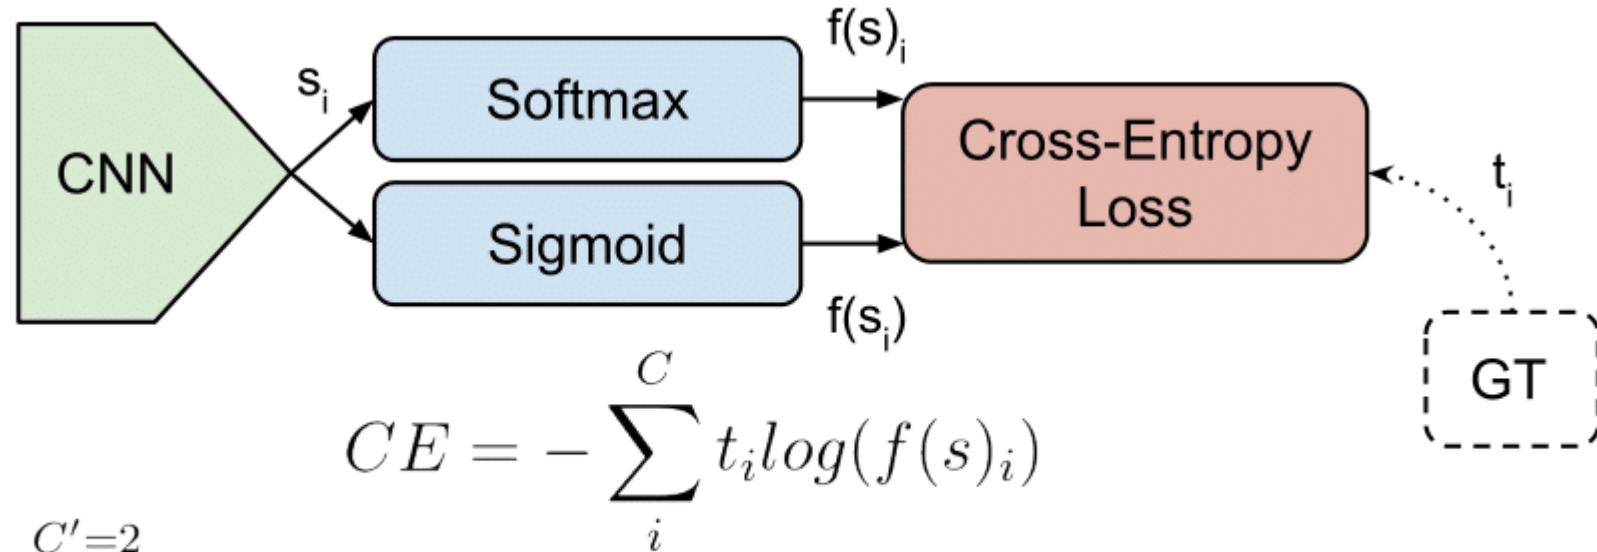
```
def MSE(y,yHat):  
    return np.sum((y-yHat )**2) / y.size
```



# Cross Entropy Loss



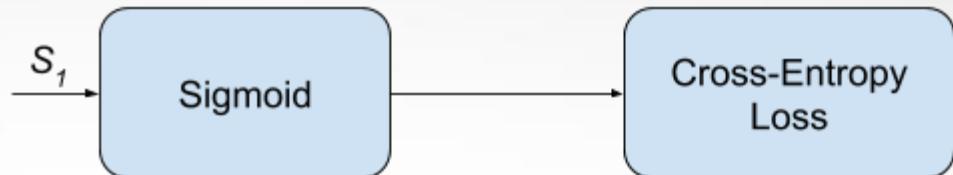
$$f(s)_i = \frac{e^{s_i}}{\sum_j^C e^{s_j}} \quad CE = - \sum_i^C t_i \log(f(s)_i)$$



$$CE = - \sum_i^C t_i \log(f(s)_i)$$

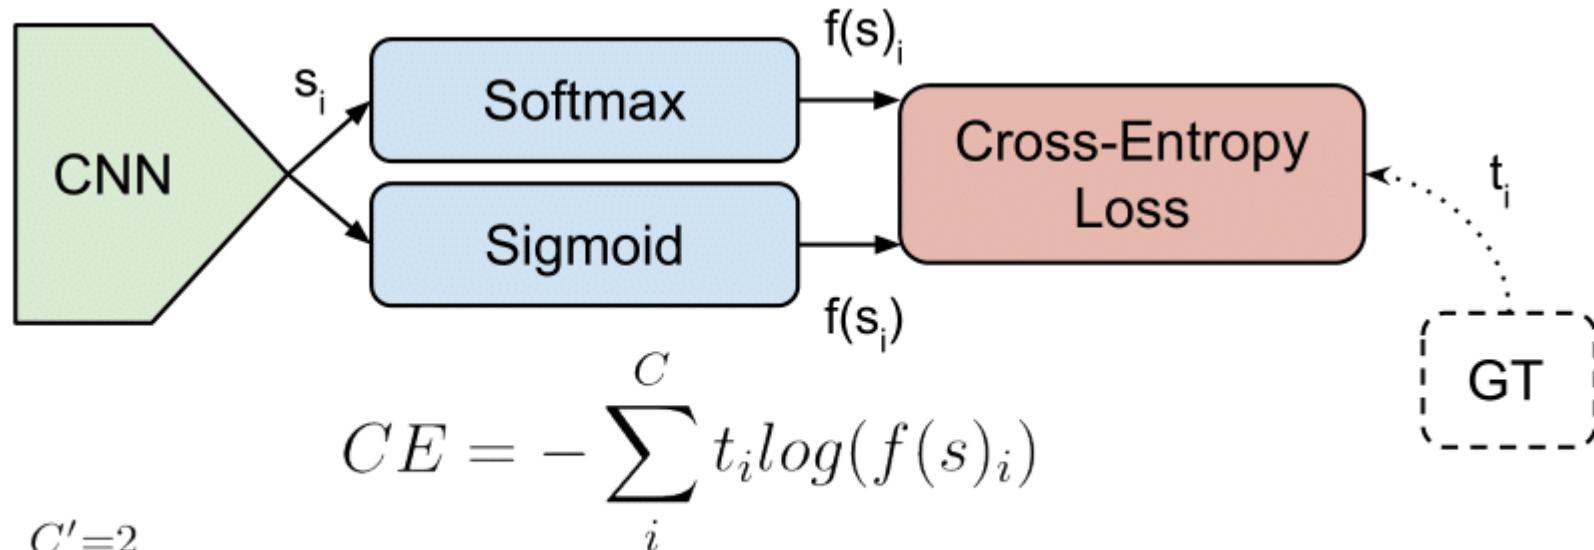
$$CE = - \sum_{i=1}^{C'=2} t_i \log(f(s_i)) = -t_1 \log(f(s_1)) - (1 - t_1) \log(1 - f(s_1))$$

# Cross Entropy Loss



$$CE = -t_1 \log(f(s_1)) - (1 - t_1) \log(1 - f(s_1))$$

$$f(s_i) = \frac{1}{1 + e^{-s_i}}$$



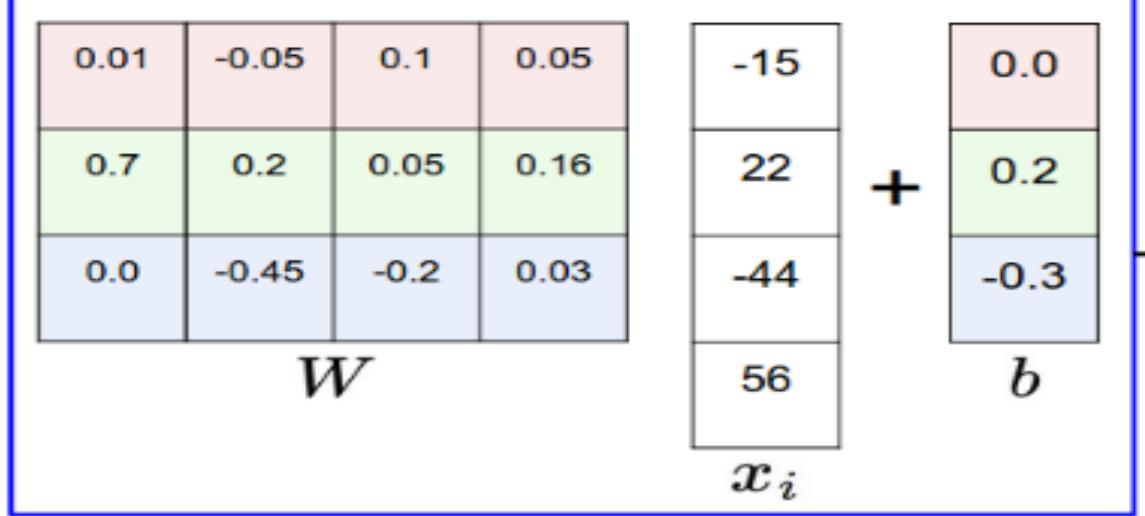
$$CE = - \sum_i^C t_i \log(f(s)_i)$$

$$CE = - \sum_{i=1}^{C'=2} t_i \log(f(s_i)) = -t_1 \log(f(s_1)) - (1 - t_1) \log(1 - f(s_1))$$

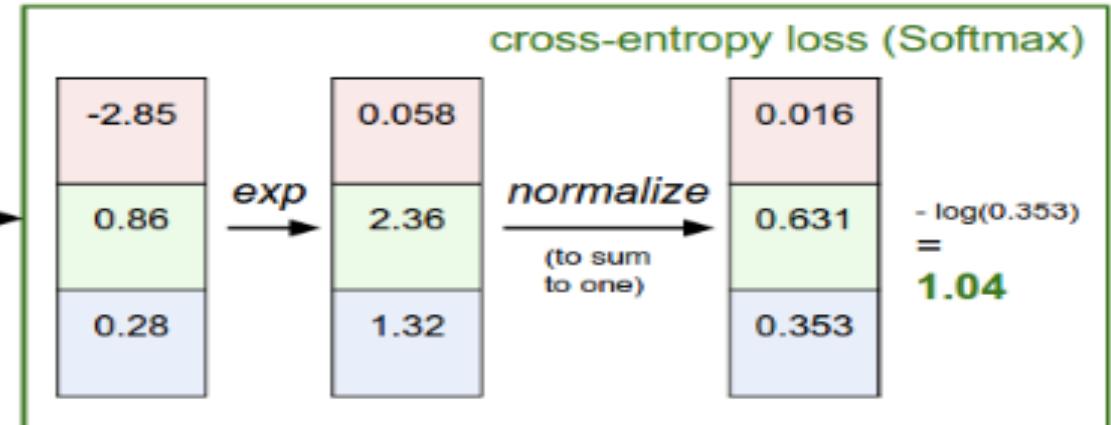
# Cross Entropy Loss



matrix multiply + bias offset



$y_i$  [2]



```
softmax = tf.exp(logits) / tf.reduce_sum(tf.exp(logits), axis)
```

```
def CrossEntropy(yHat, y):  
    if y == 1:  
        return -log(yHat)  
    else:  
        return -log(1 - yHat)
```

# Gradient and Stochastic Gradient Descent

- Stochastic gradient descent can be regarded as a stochastic approximation of gradient descent optimization, since it replaces the actual gradient (calculated from the entire data set) by an estimate thereof (calculated from a randomly selected subset of the data).
- Stochastic gradient descent is a popular algorithm for training a wide range of models in machine learning.
- When combined with the backpropagation algorithm, it is the de facto standard algorithm for training artificial neural networks.

# Stochastic Gradient Descent



- A compromise between computing the true gradient and the gradient at a single example is to compute the gradient against more than one training example (called a "mini-batch") at each step. This can perform significantly better than "true" stochastic gradient descent described.

$$Q(w) = \frac{1}{n} \sum_{i=1}^n Q_i(w), \quad Q_i(w) \text{ is the loss function}$$

a standard (or "batch") gradient descent method

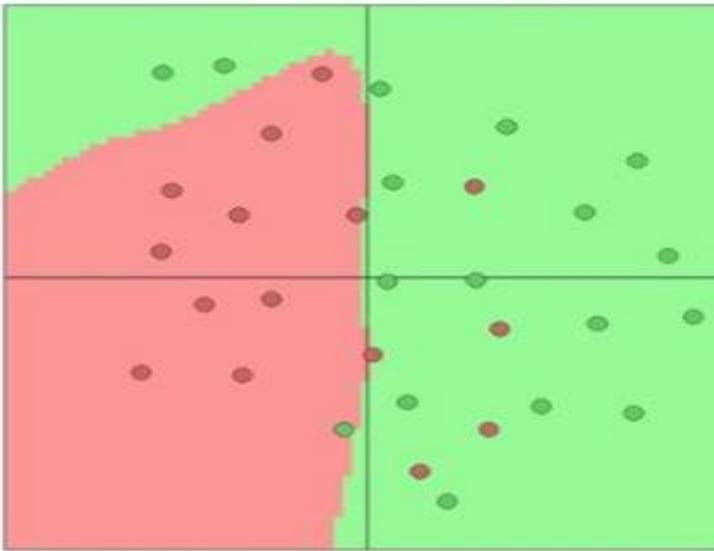
$$w := w - \eta \nabla Q(w) = w - \eta \sum_{i=1}^n \nabla Q_i(w) / n,$$

Wikipedia

# Effects of Network Structure



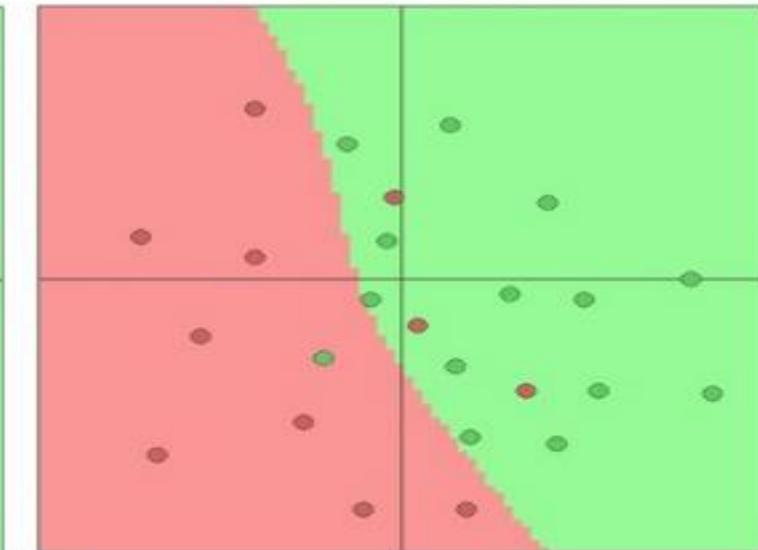
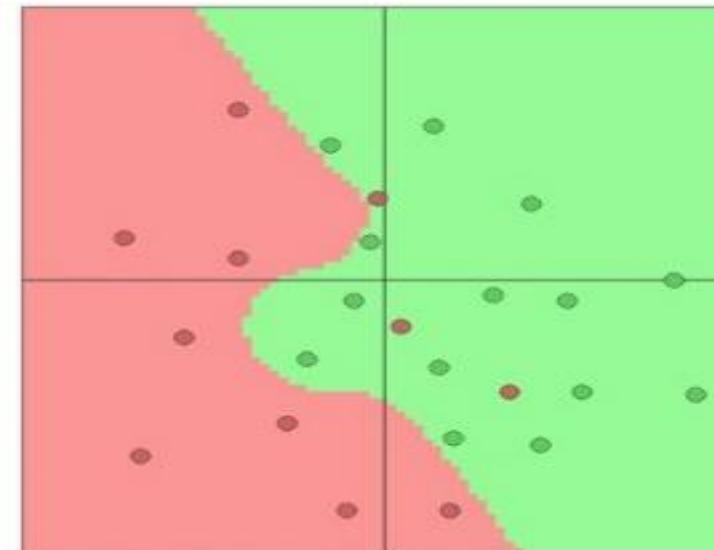
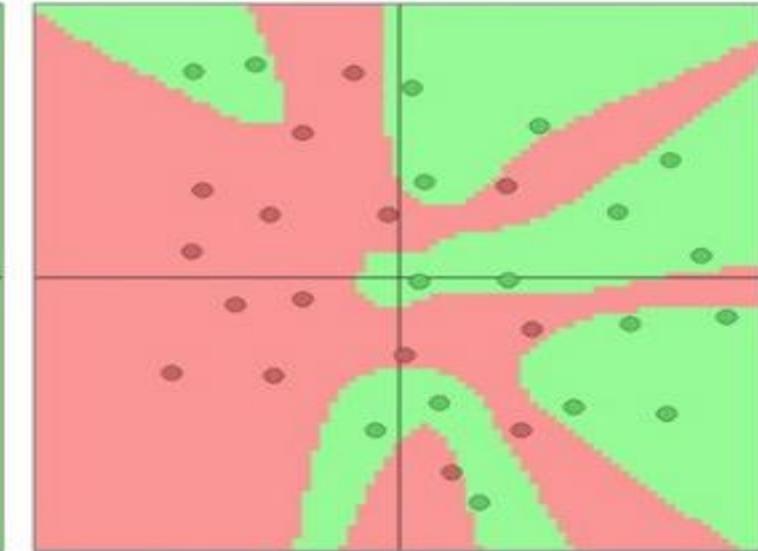
3 hidden neurons



6 hidden neurons

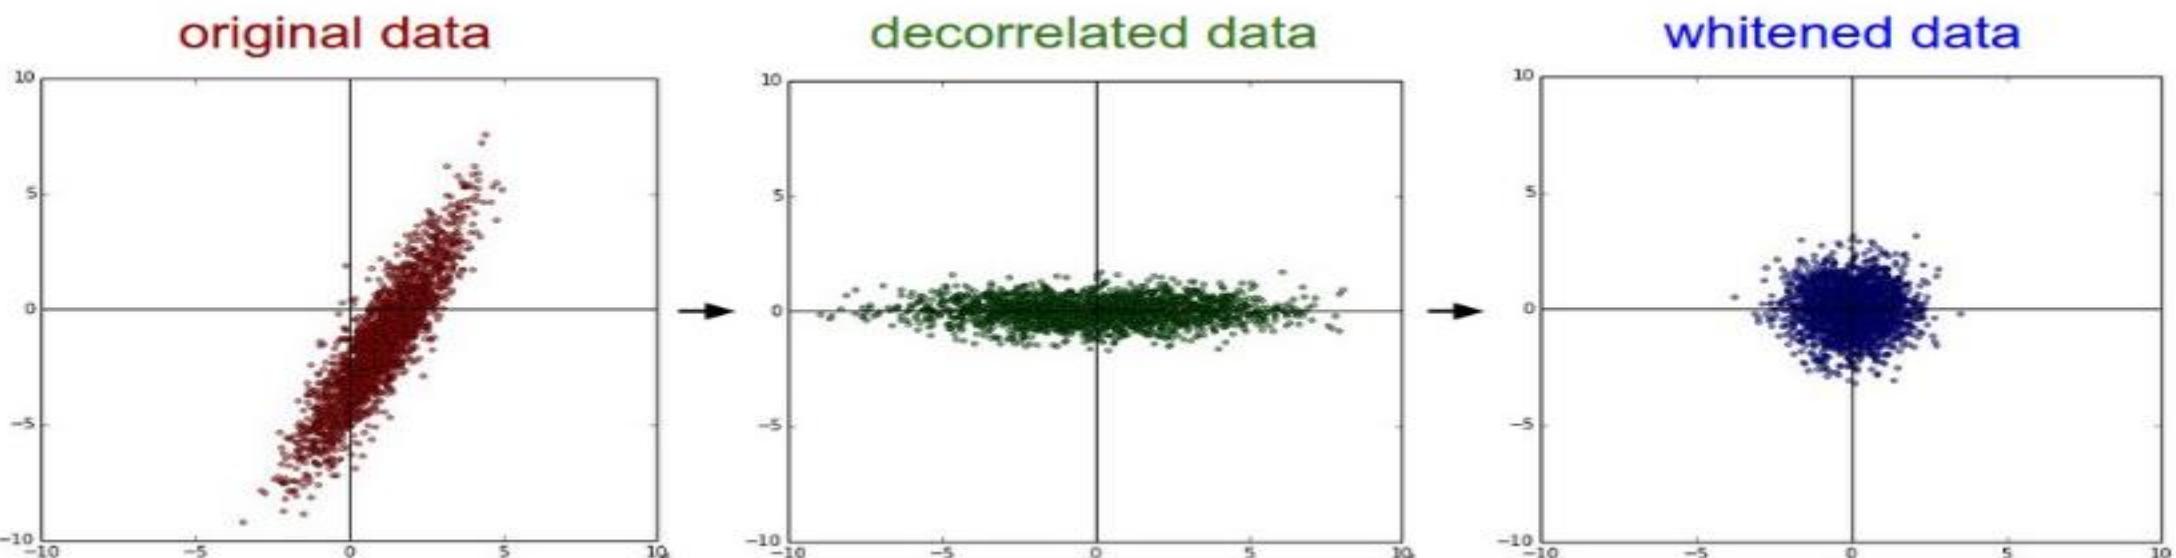
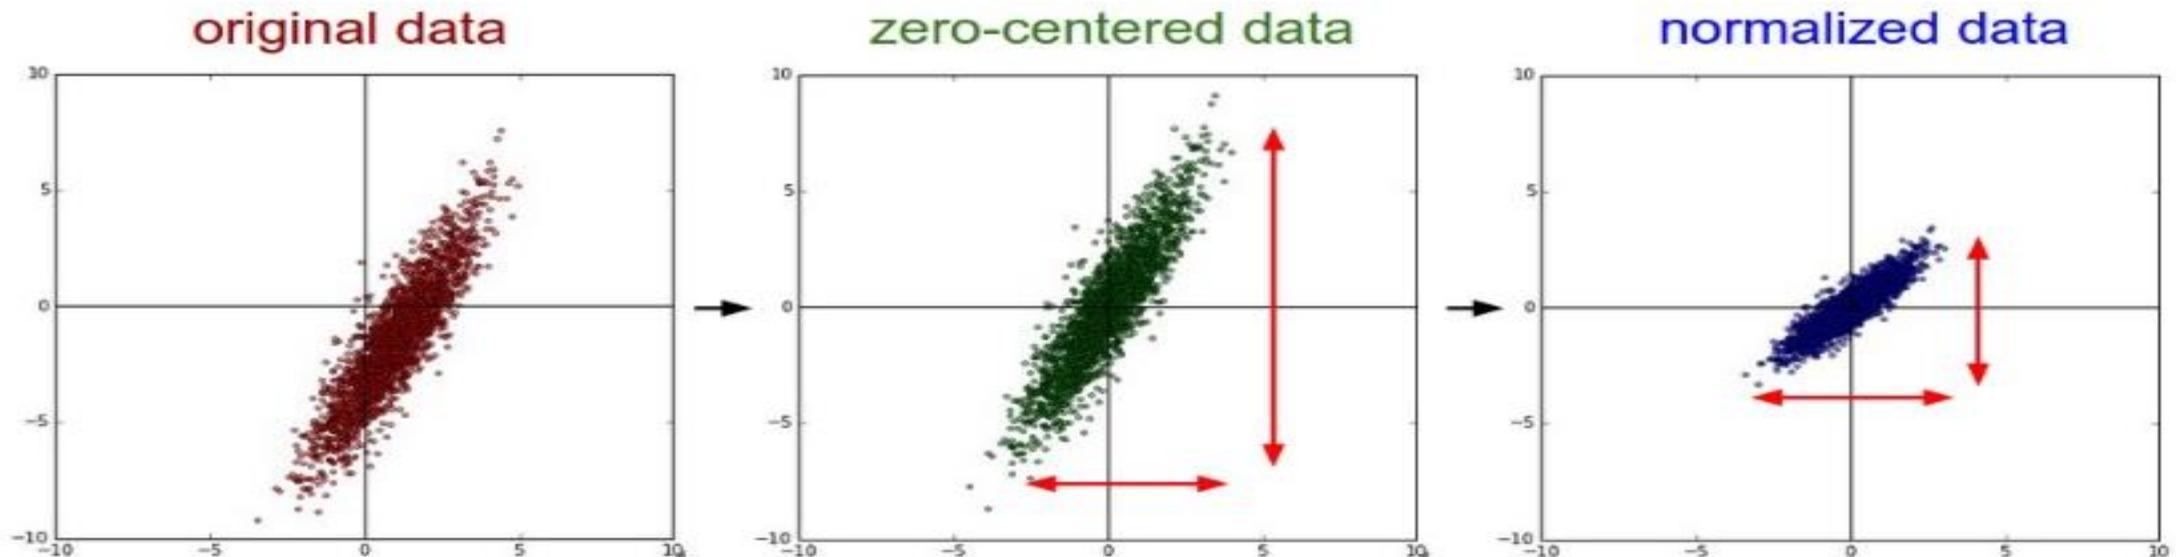


20 hidden neurons

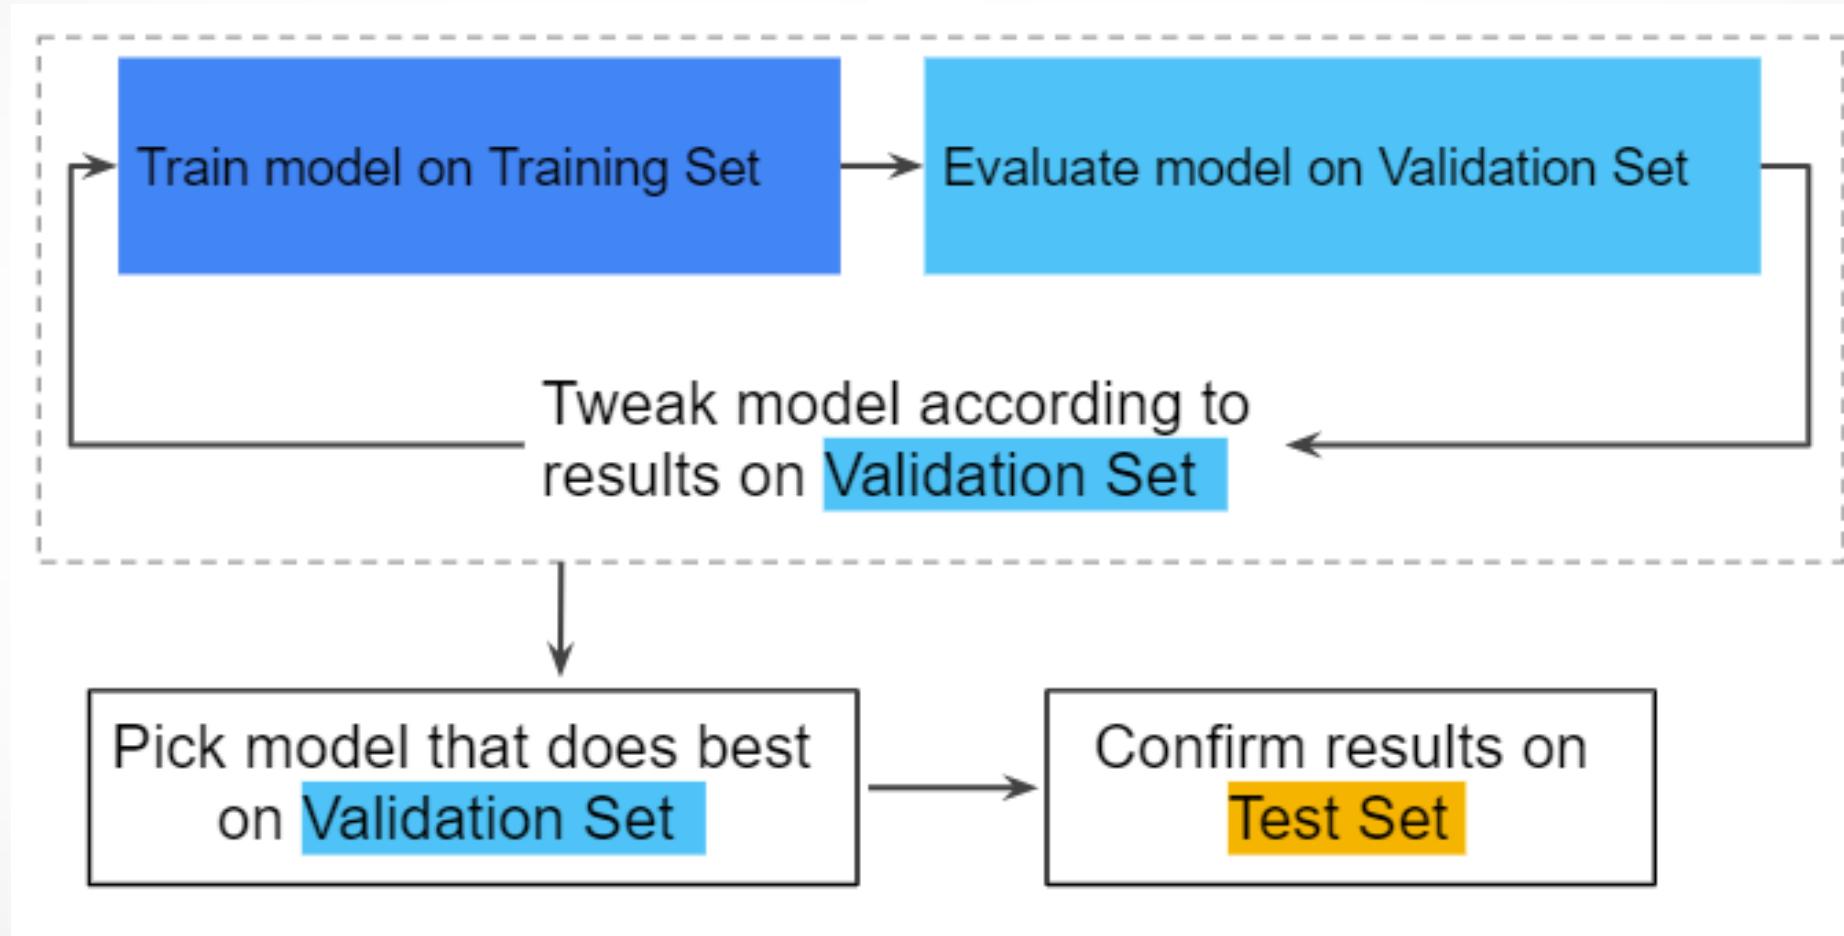


# Data Preprocessing

- Normalization and Dimensional reduction

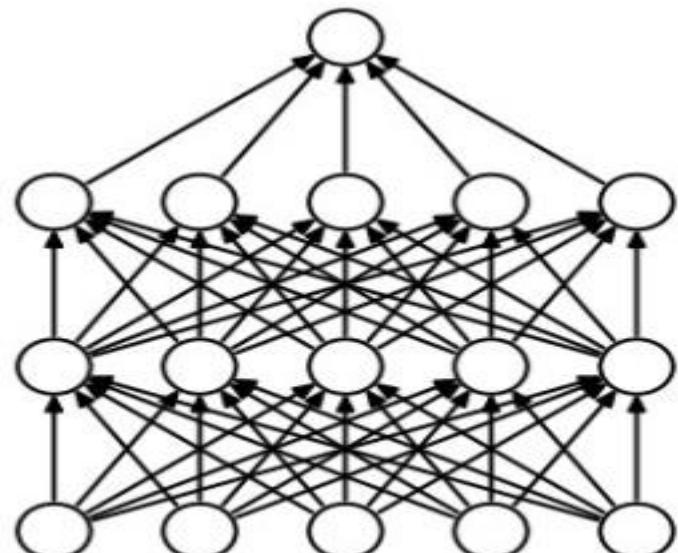


# Training, Validation & Test Set

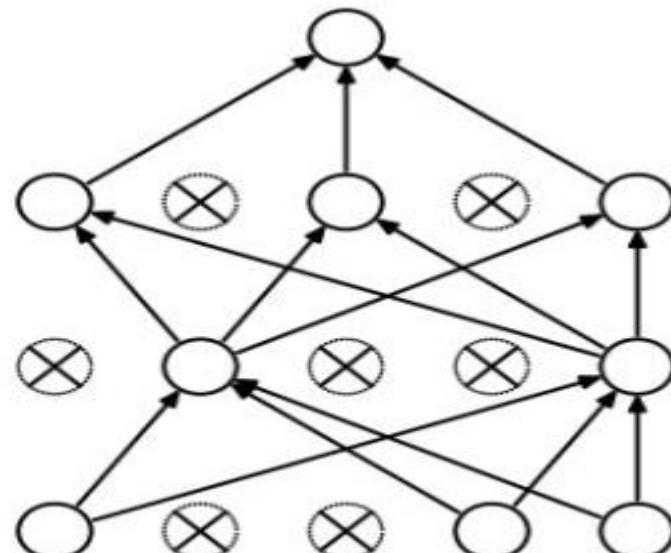


# Dropout

- Deal with overfitting with dropout



(a) Standard Neural Net



(b) After applying dropout.



# **Computational Graph**

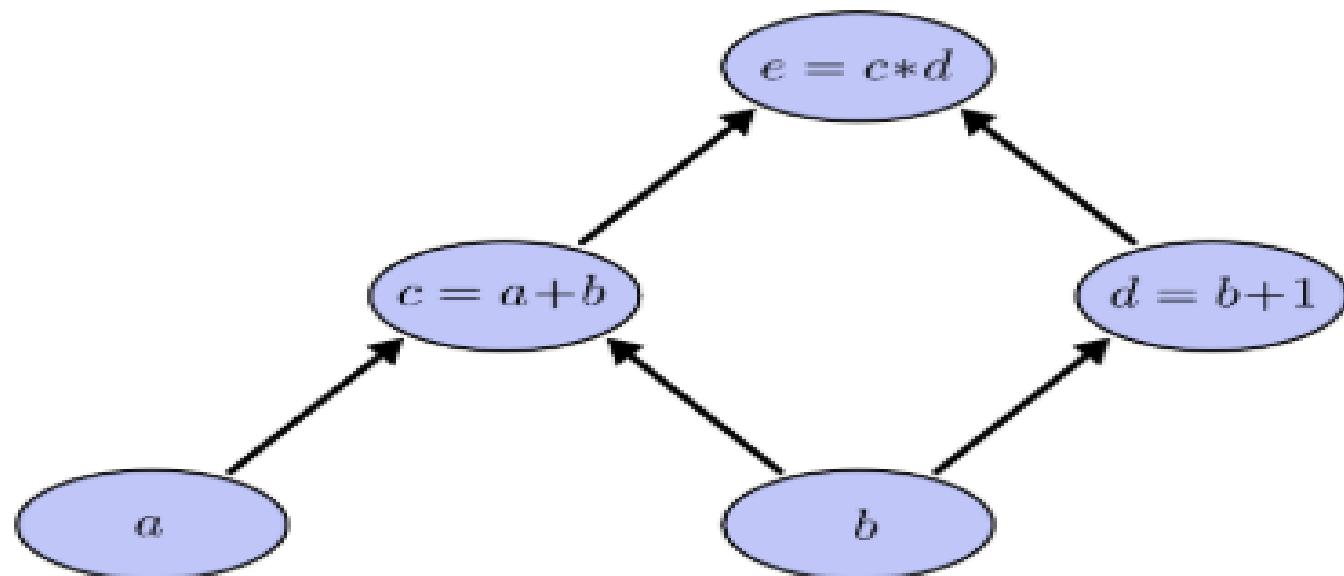
# **& TensorFlow**

AI workshop. Pre - Coding Conquest 2019 event

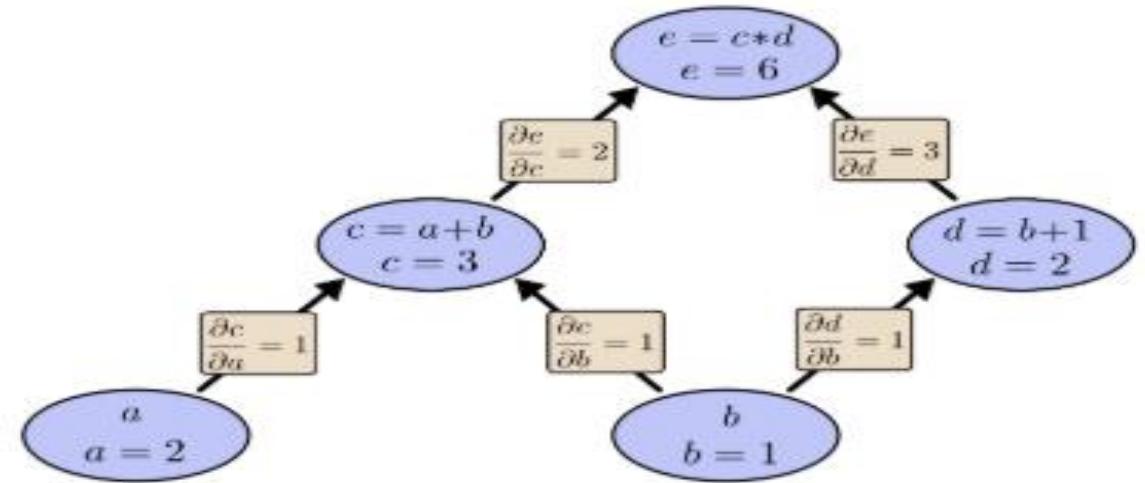
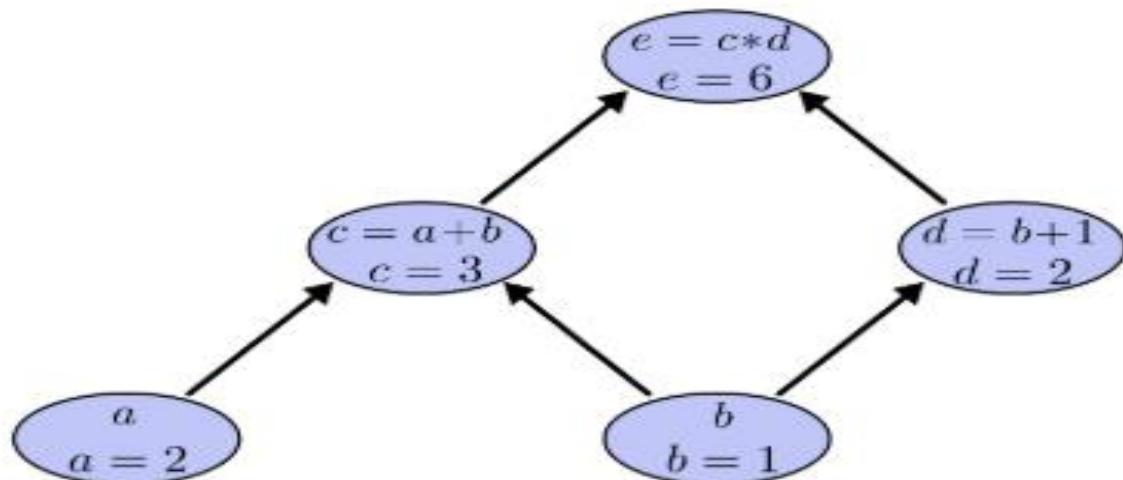
# Learning: Computational Graph



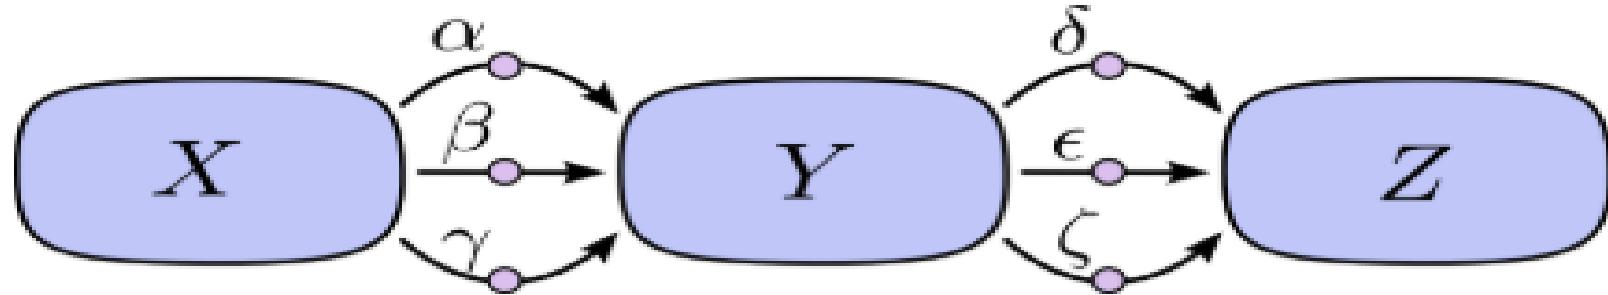
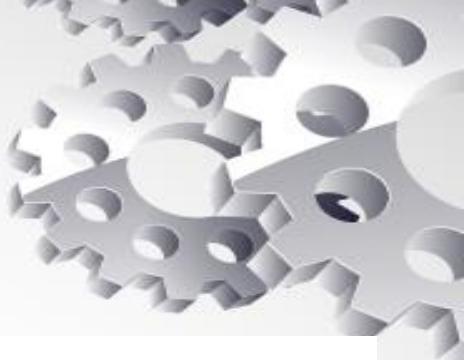
- Backpropagation can be computed via flow in the computational graph



# Computational Graph



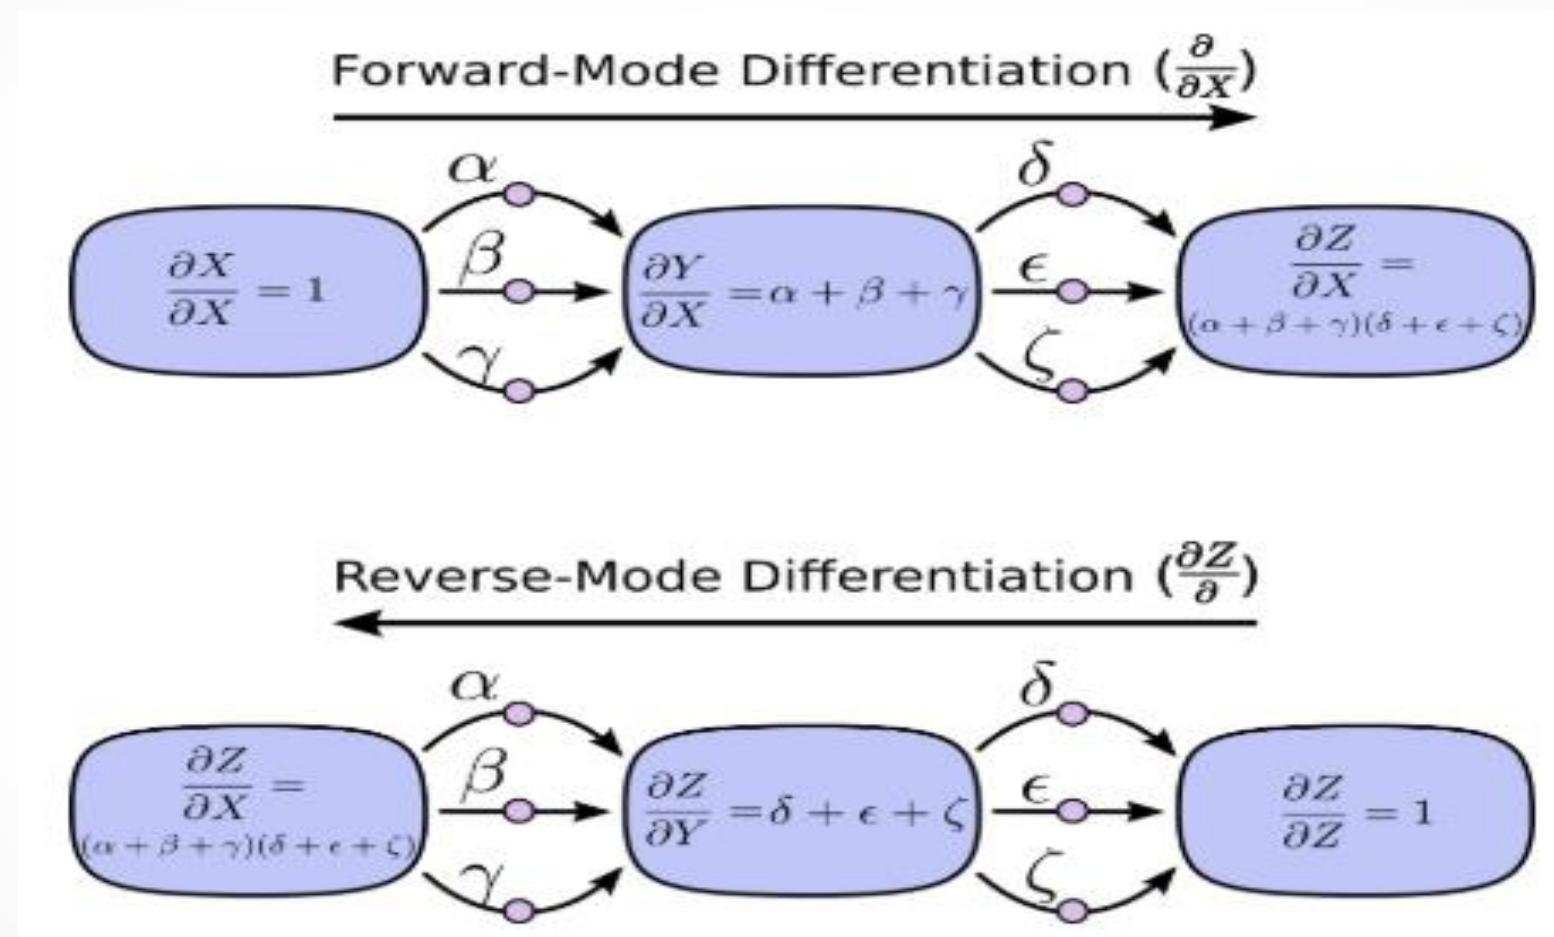
# Computational Graph



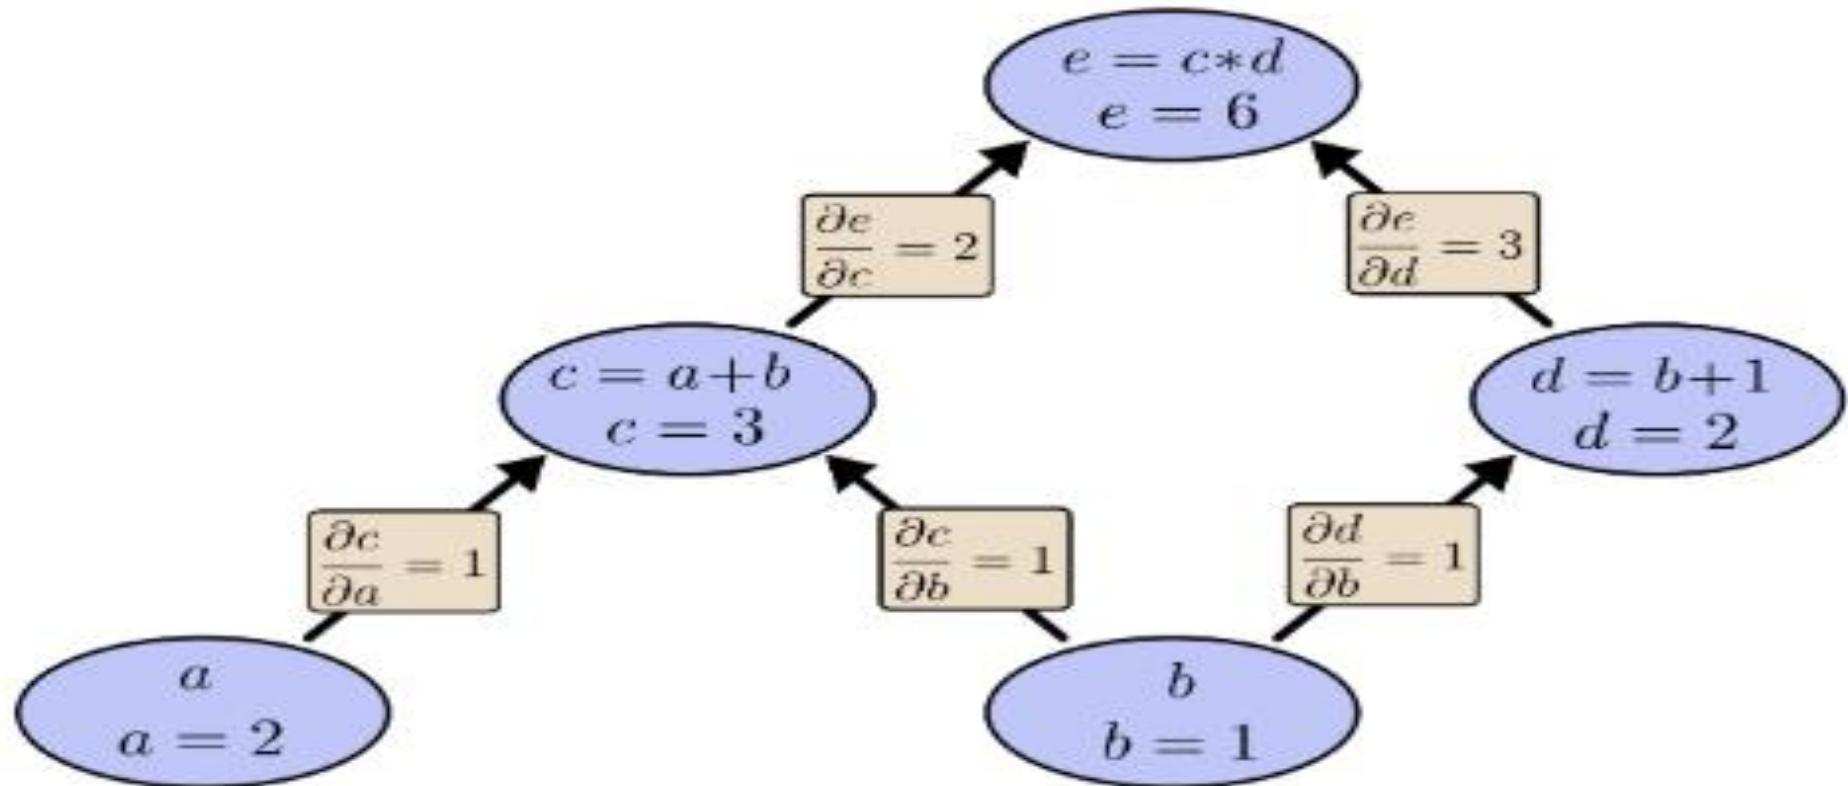
In the above diagram, there are three paths from  $X$  to  $Y$ , and a further three paths from  $Y$  to  $Z$ . If we want to get the derivative  $\frac{\partial Z}{\partial X}$  by summing over all paths, we need to sum over  $3 * 3 = 9$  paths:

$$\frac{\partial Z}{\partial X} = \alpha\delta + \alpha\epsilon + \alpha\zeta + \beta\delta + \beta\epsilon + \beta\zeta + \gamma\delta + \gamma\epsilon + \gamma\zeta$$

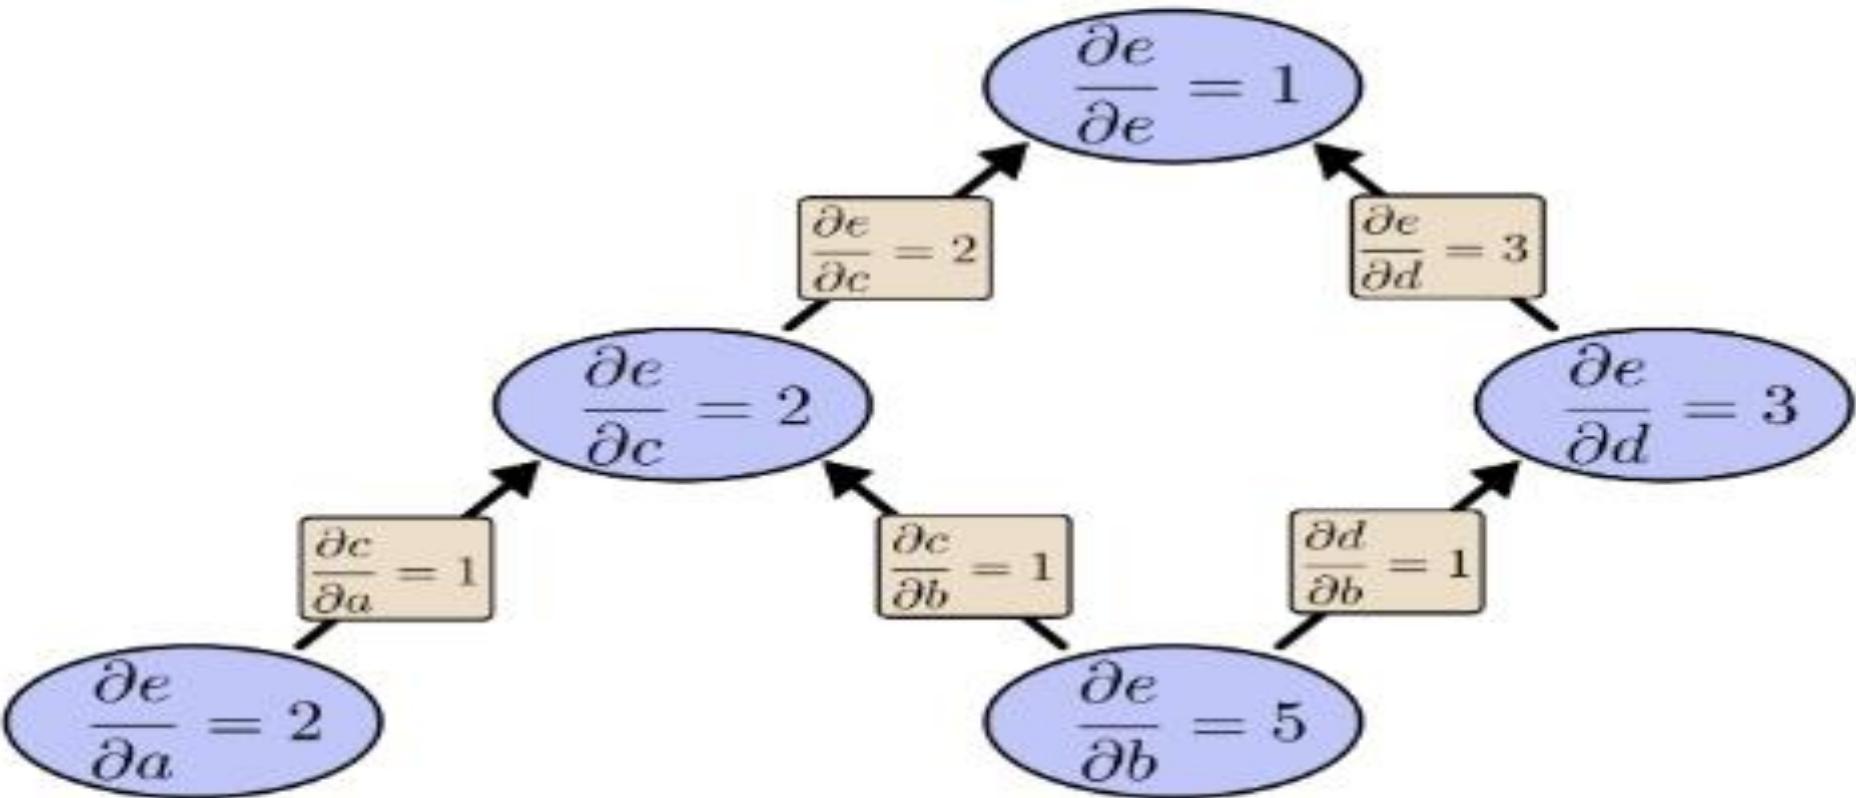
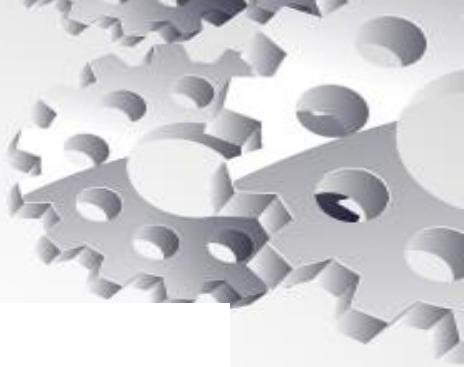
# Computational Graph



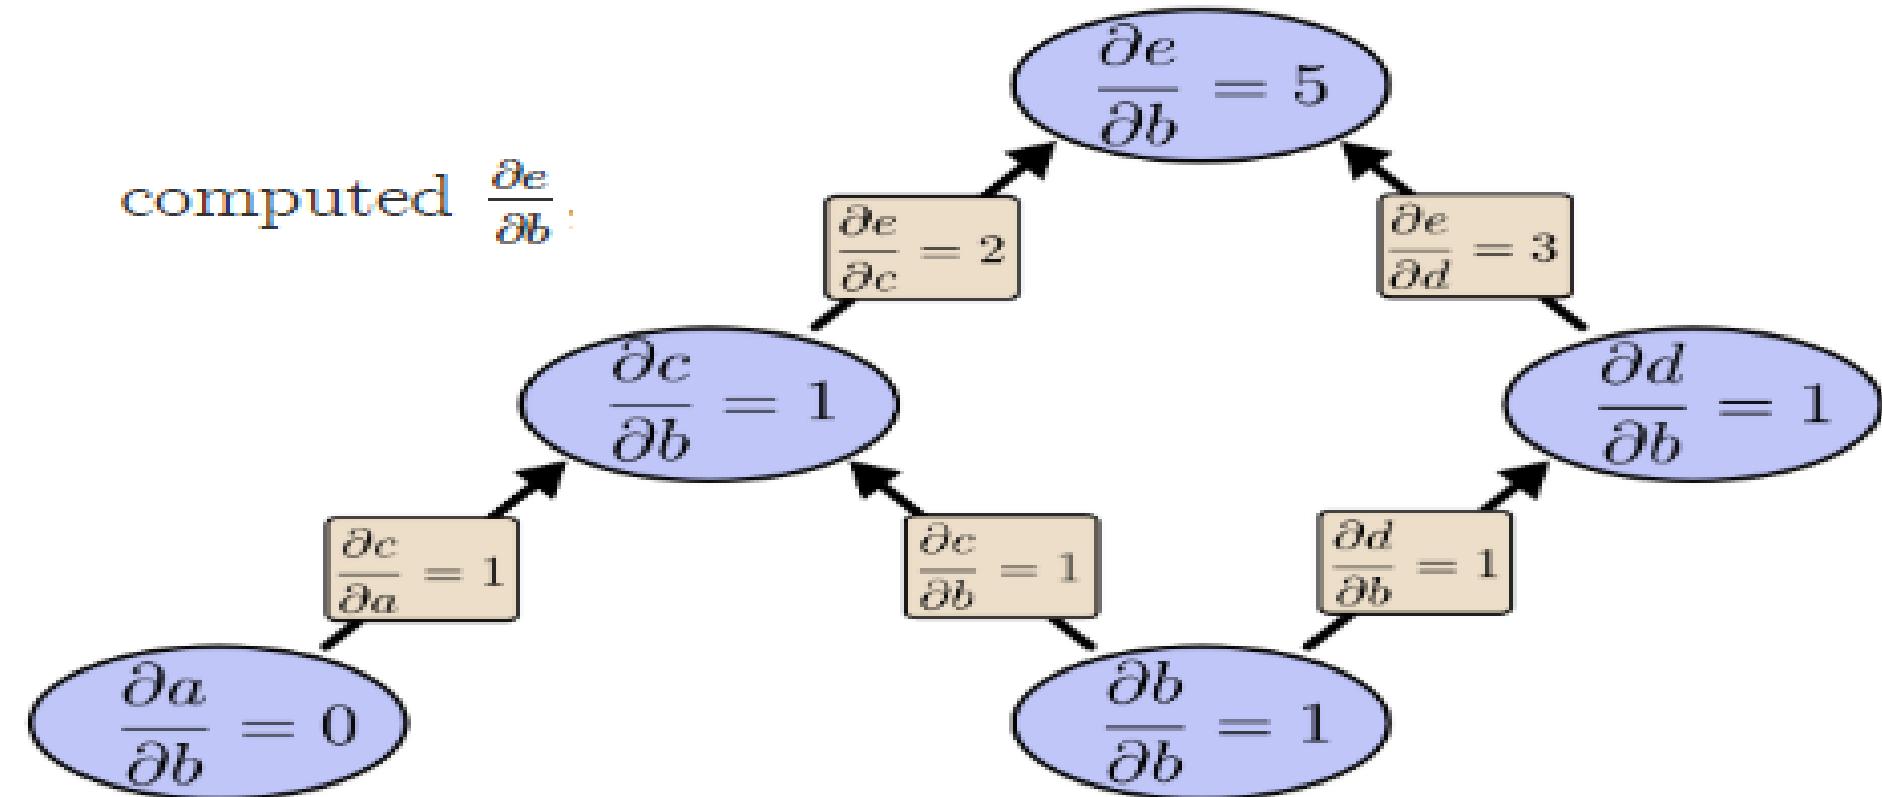
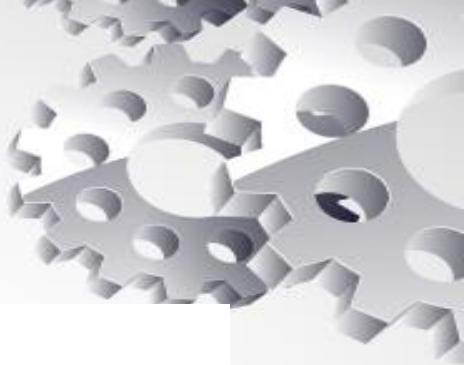
# Computational Graph



# Computational Graph

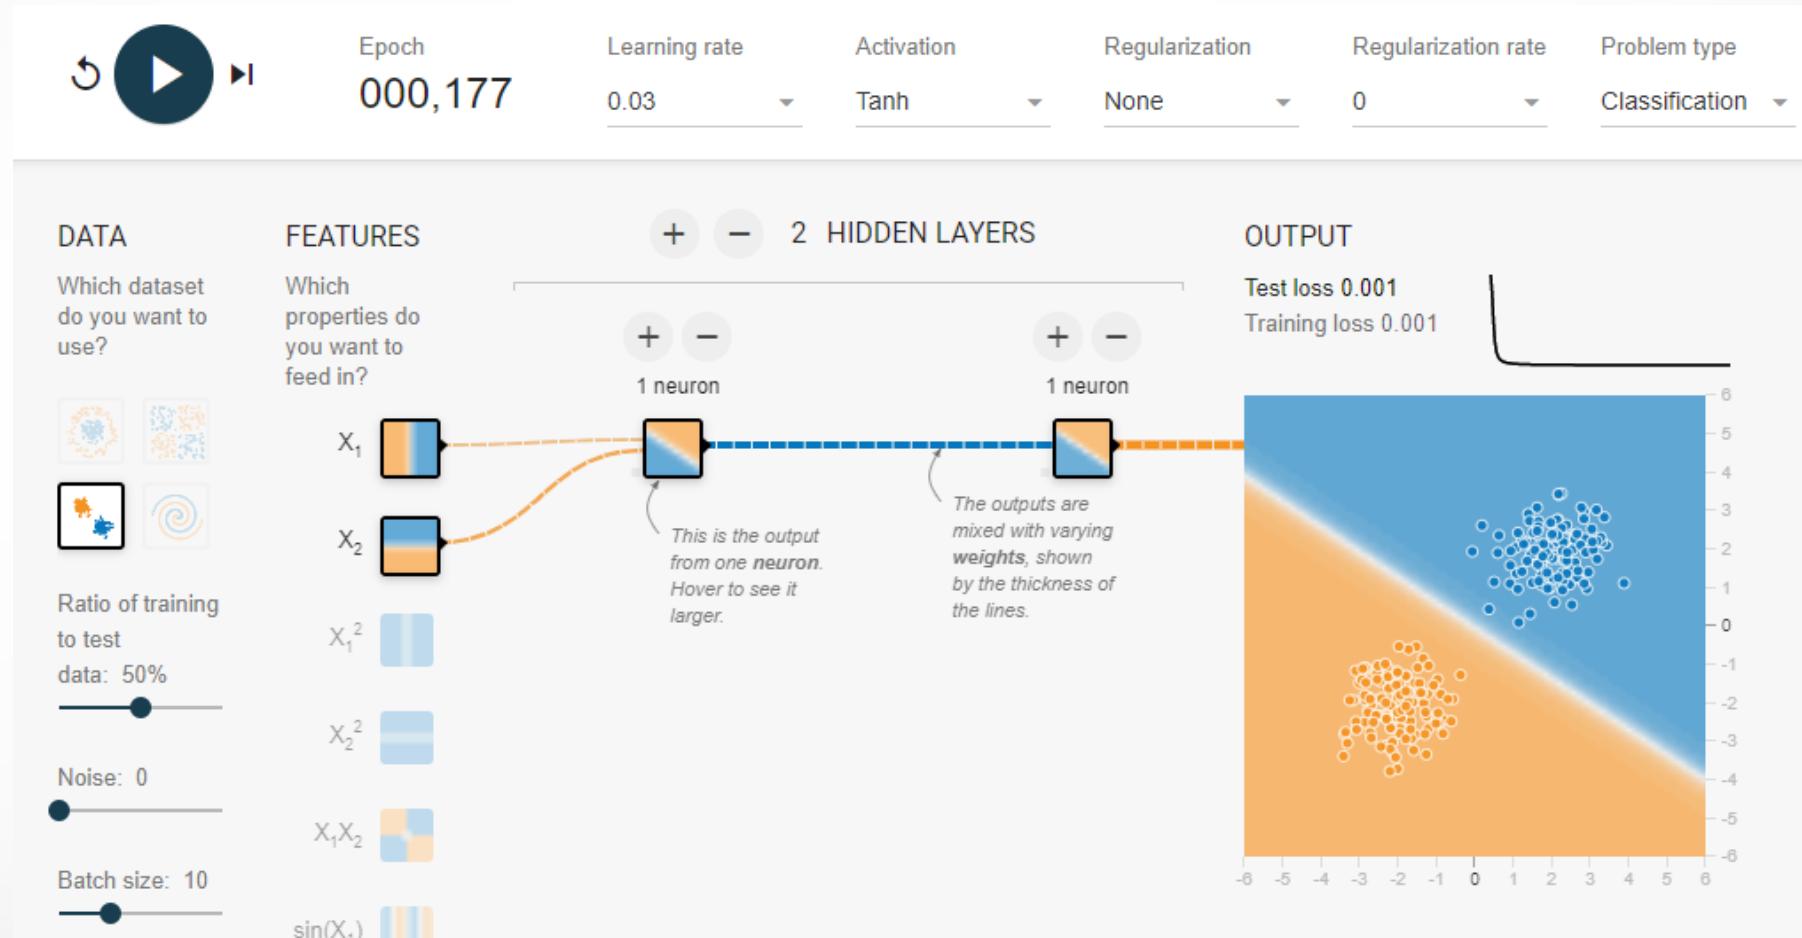


# Computational Graph



# TensorFlow ANN Playground

- <https://playground.tensorflow.org>



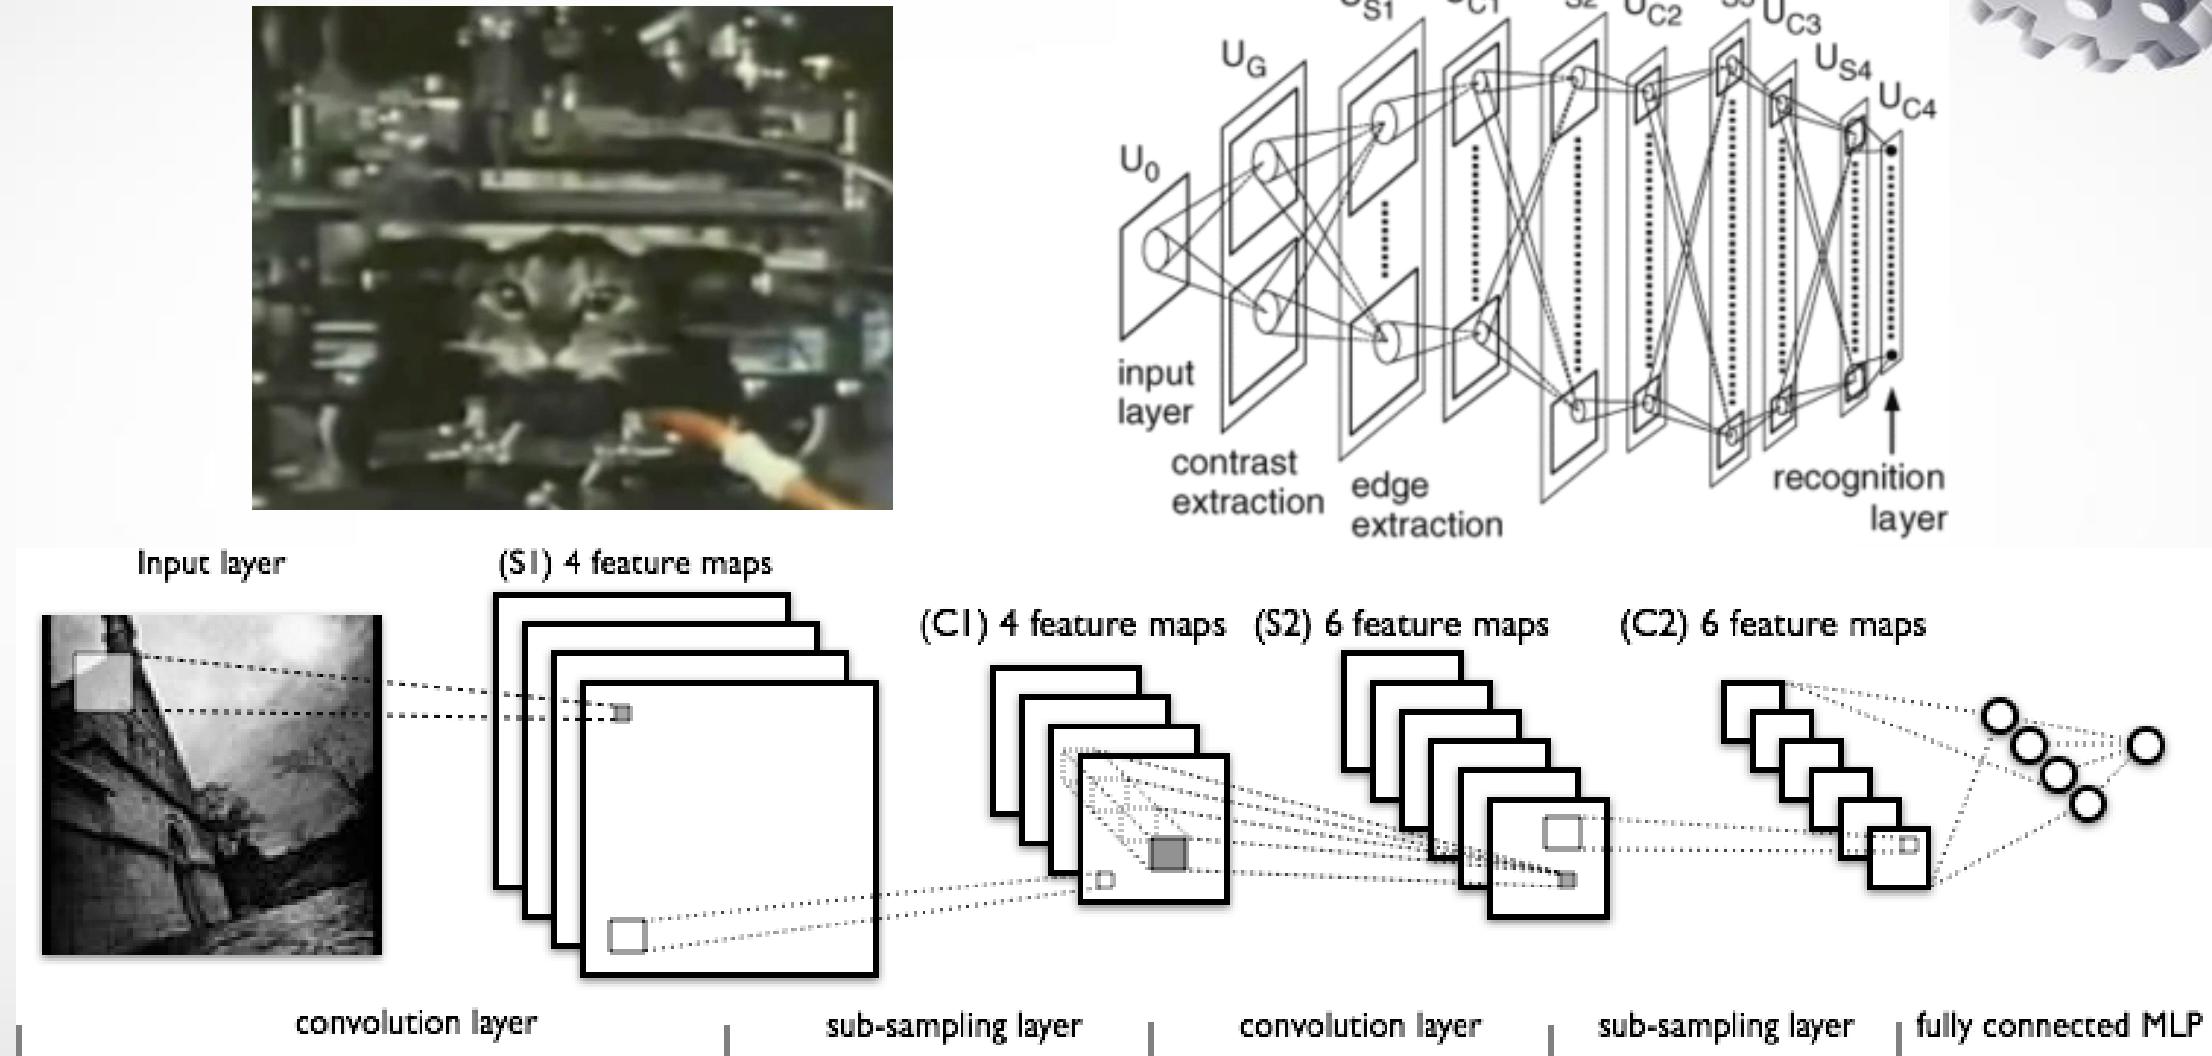


# **Deep Learning**

# **CNN based computer vision**

AI workshop. Pre - Coding Conquest 2019 event

# A Brief History of Deep Learning

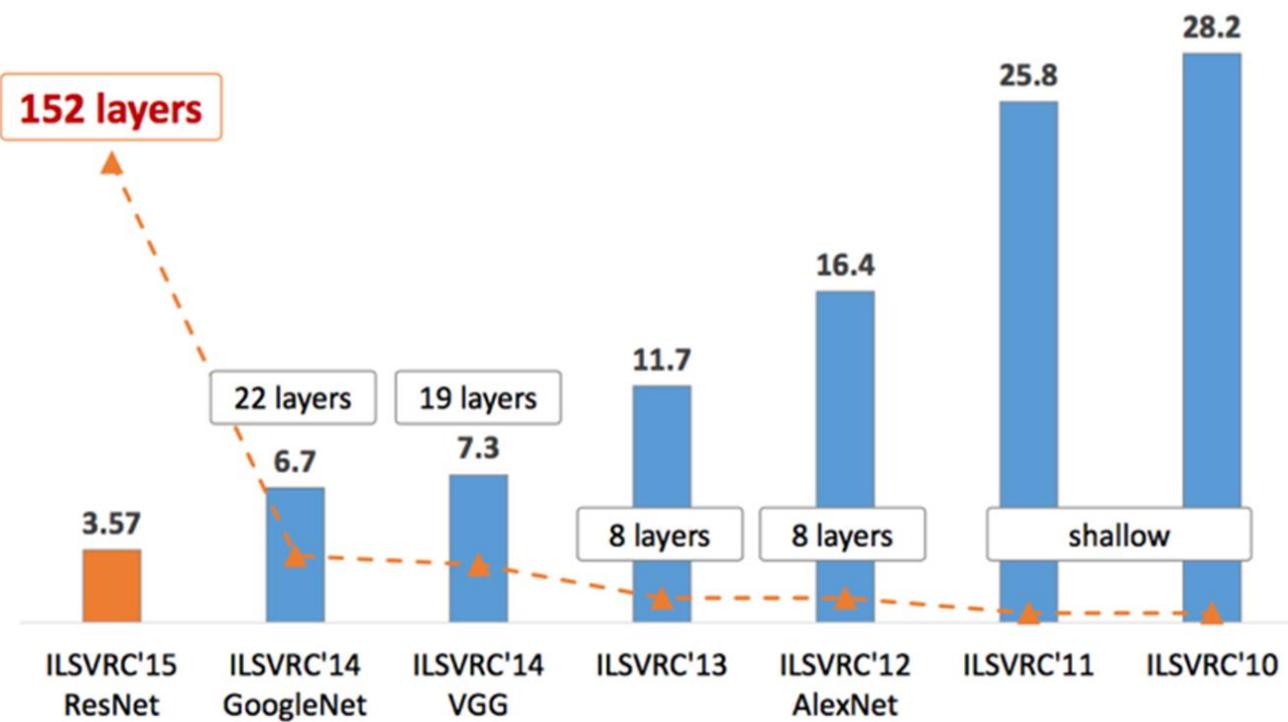
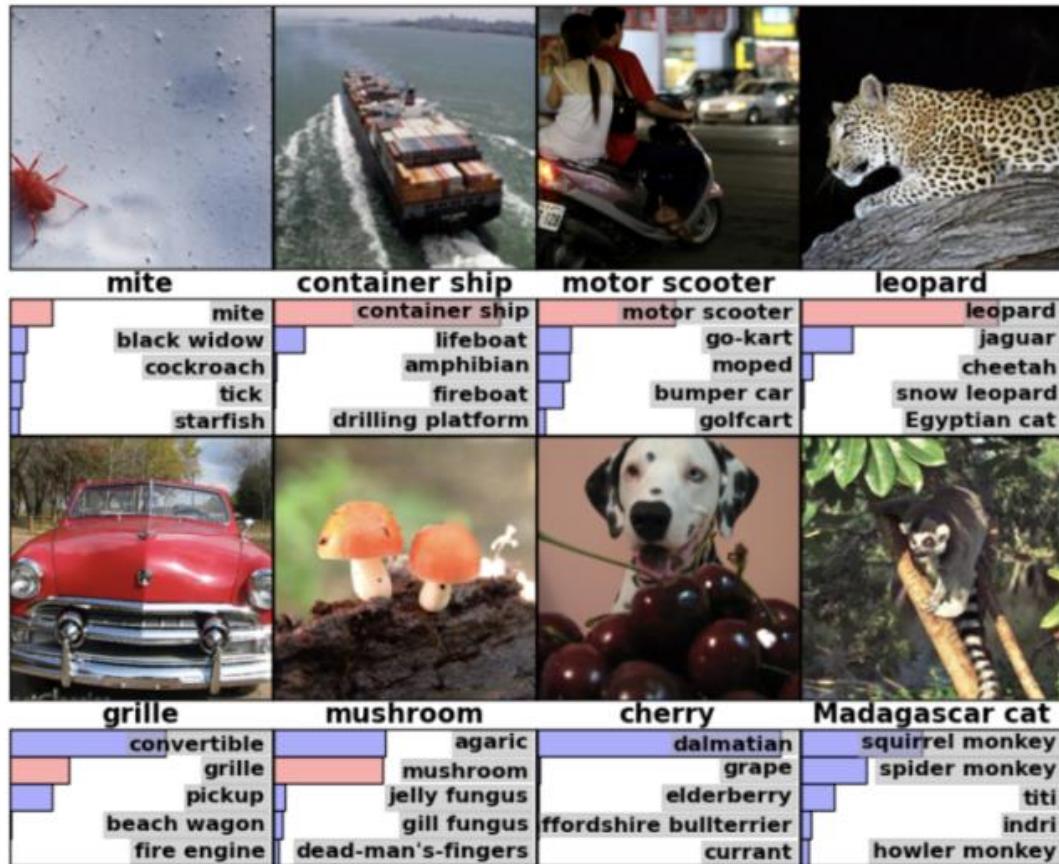


# A Brief History of Deep Learning

- Hubel and Wiesel (1959) found two types of important cells in the visual primary cortex: *simple cell* and *complex cell*.
- The neocognitron (Fukushima, 1979) was inspired by the model proposed by Hubel & Wiesel.
- Convolutional neural network (Yann LeCun, 1989) was also inspired by this line of concepts.

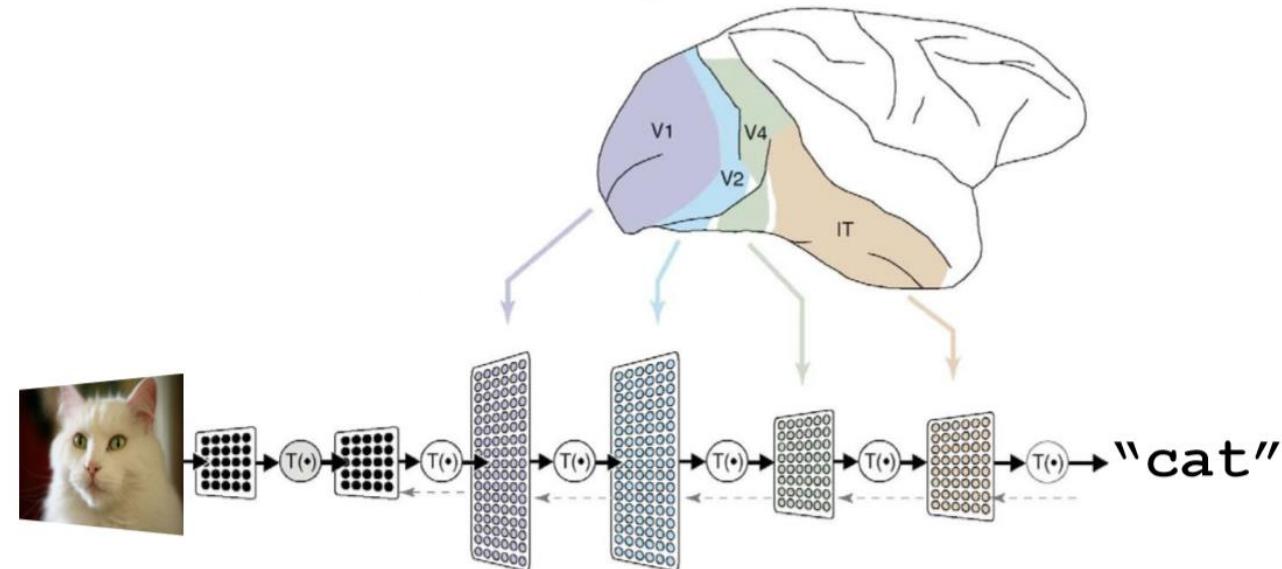


# ImageNet (Fei-Fei Li, 2009)

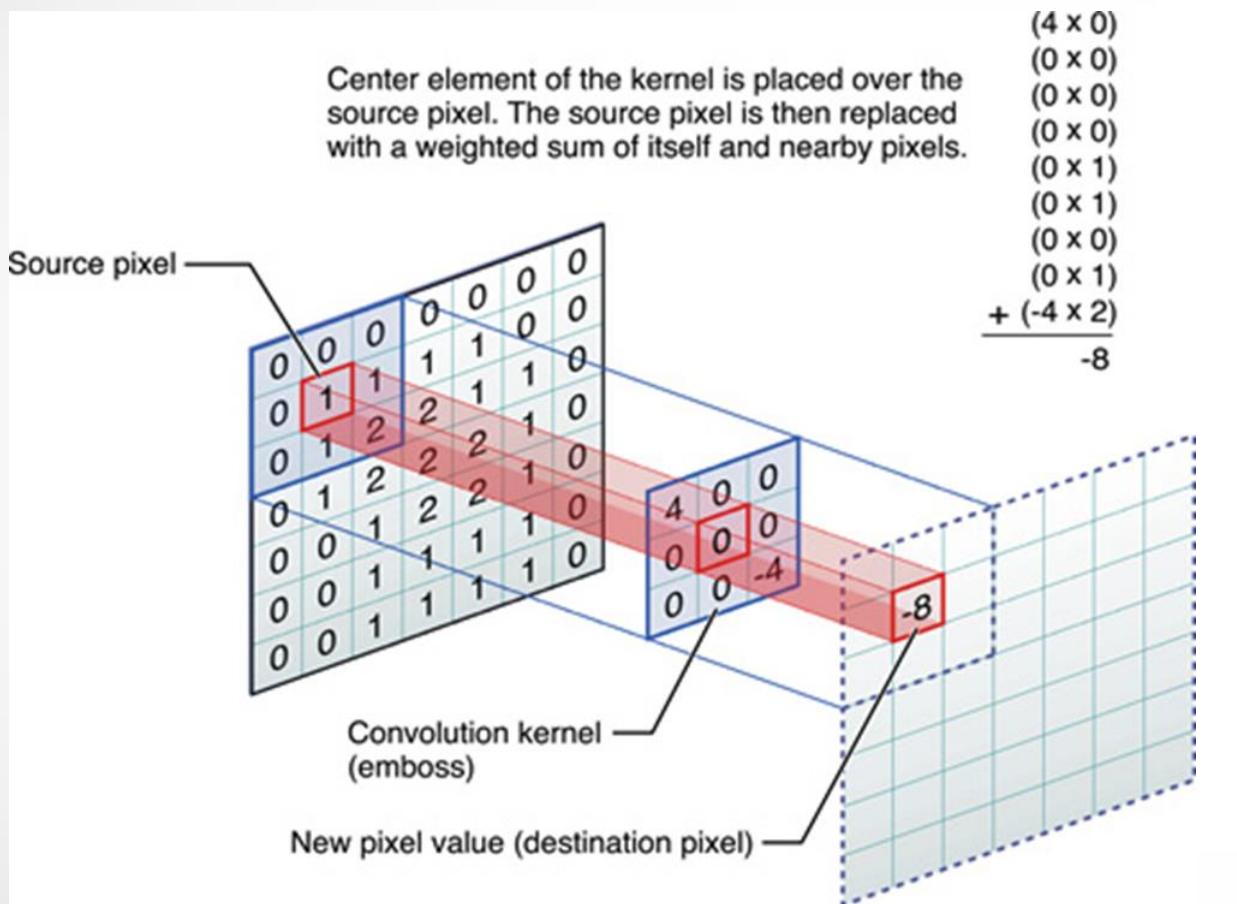


# Why does CNN Perform Better?

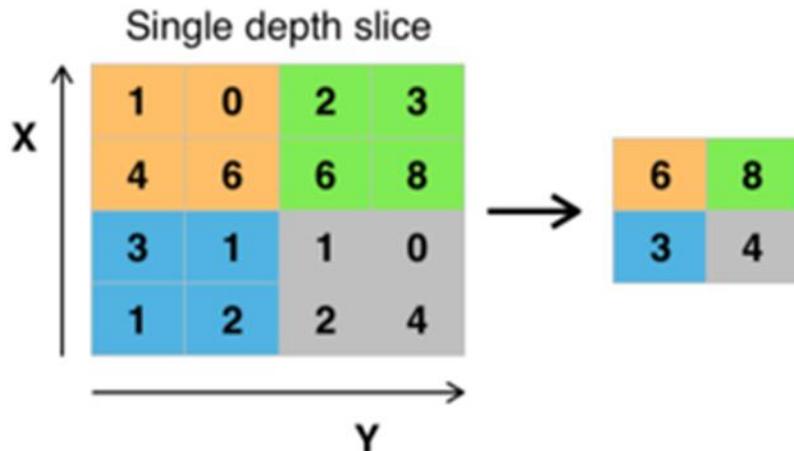
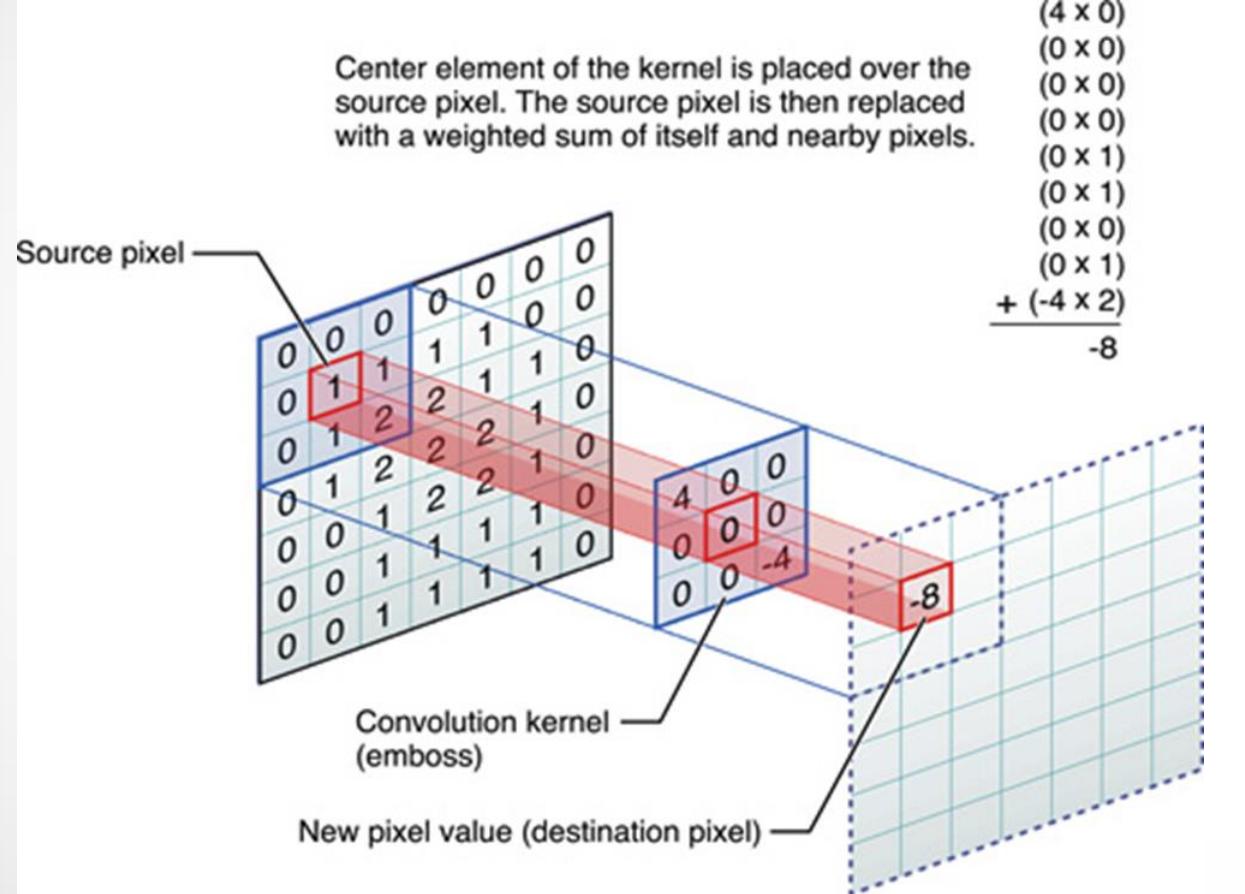
- CNN shares many commonalities with real brains.
  - Each neuron is connected to a small subset of neurons.
  - Each neuron perform simple function e.g., Sigmoid, ReLU.
  - Each subset of neurons learn specialised features (representation learning).



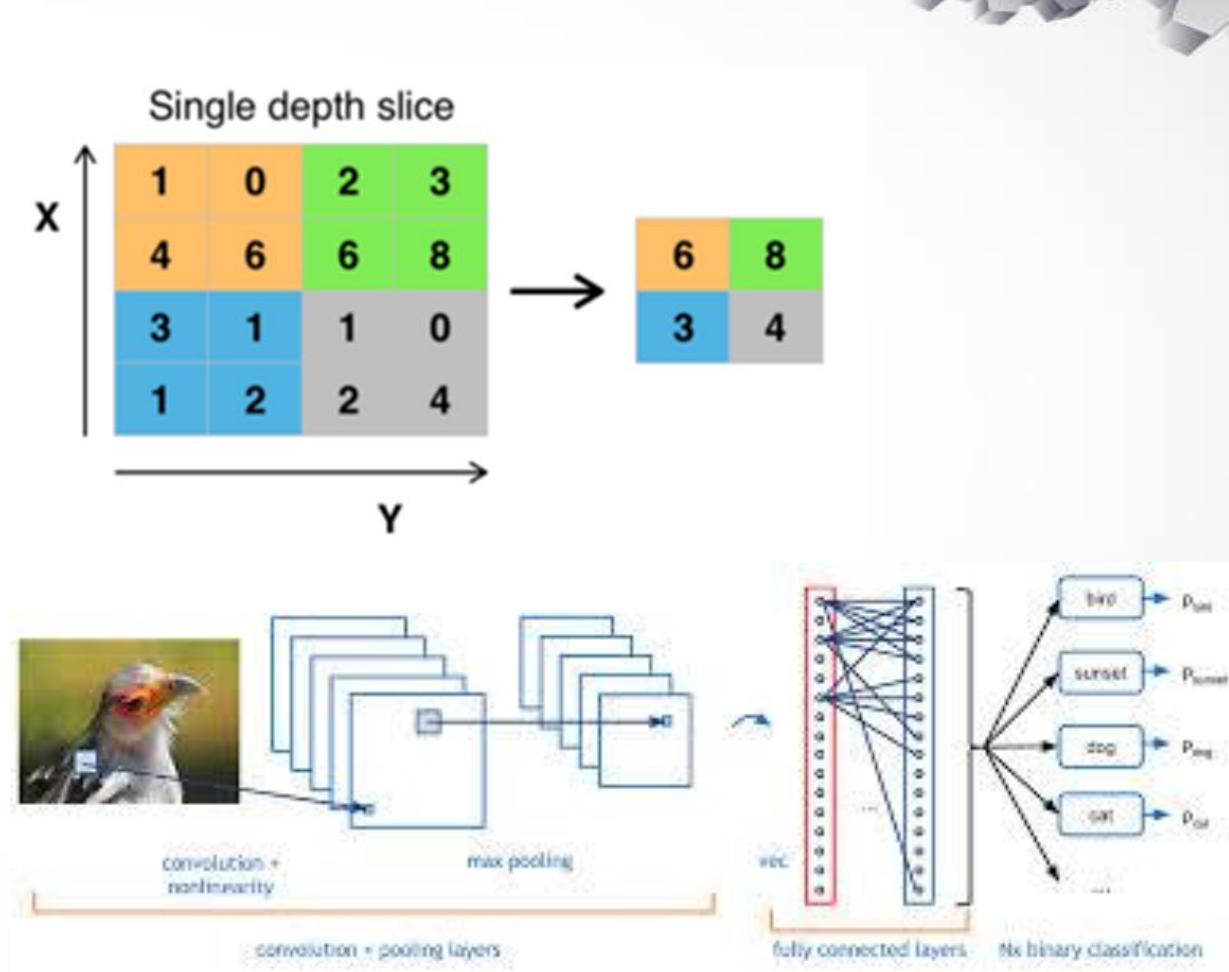
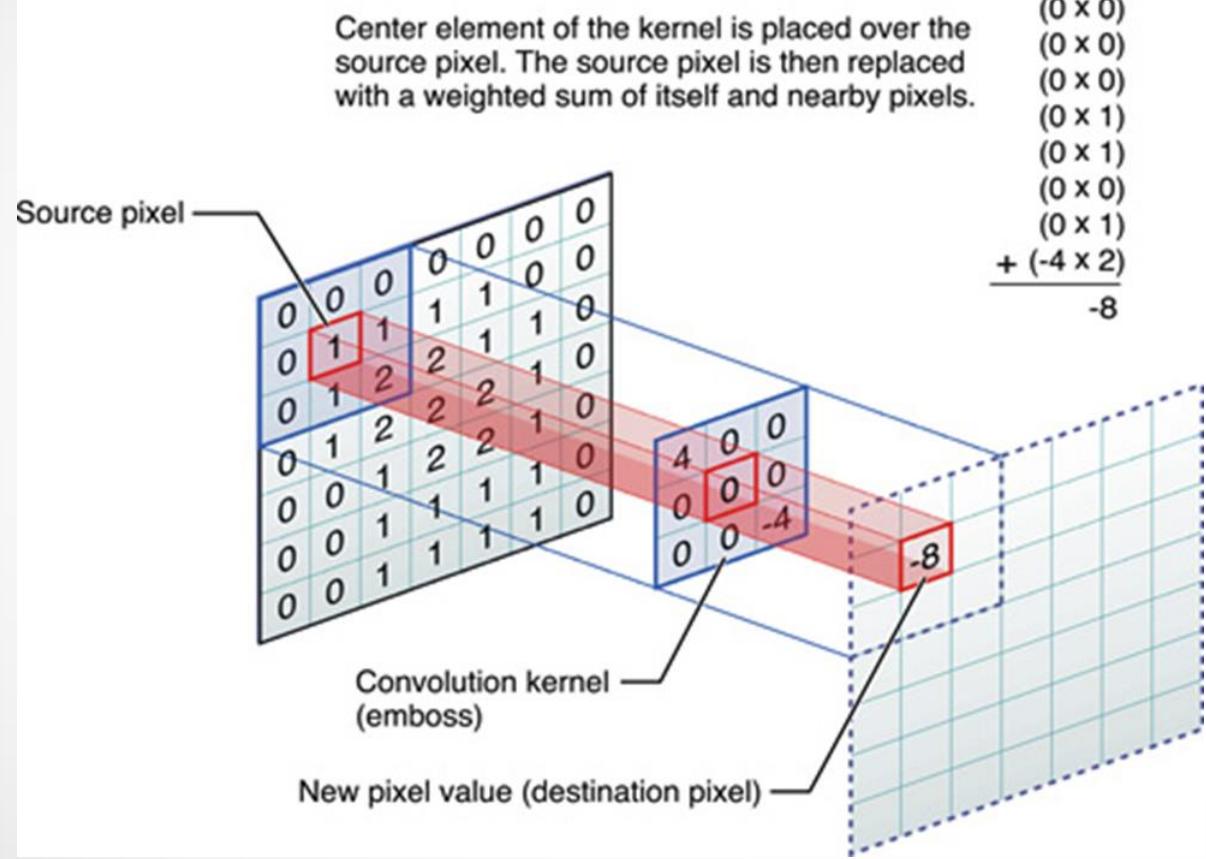
# Convolutional Neural Network



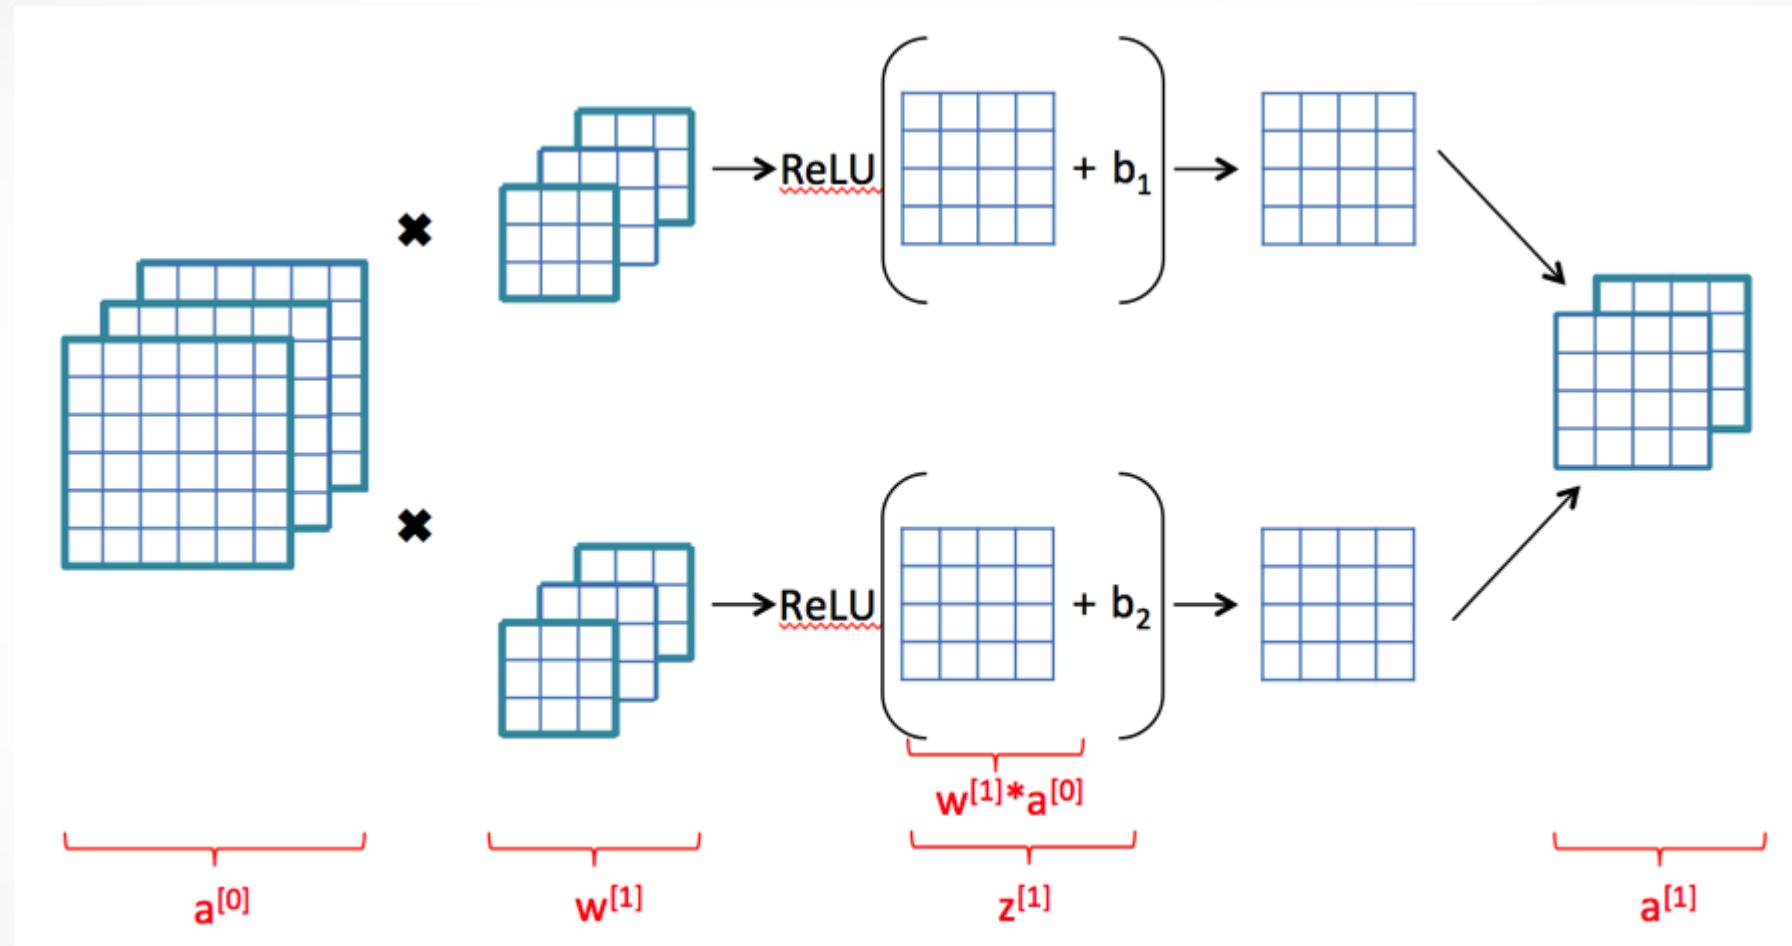
# Convolutional Neural Network



# Convolutional Neural Network

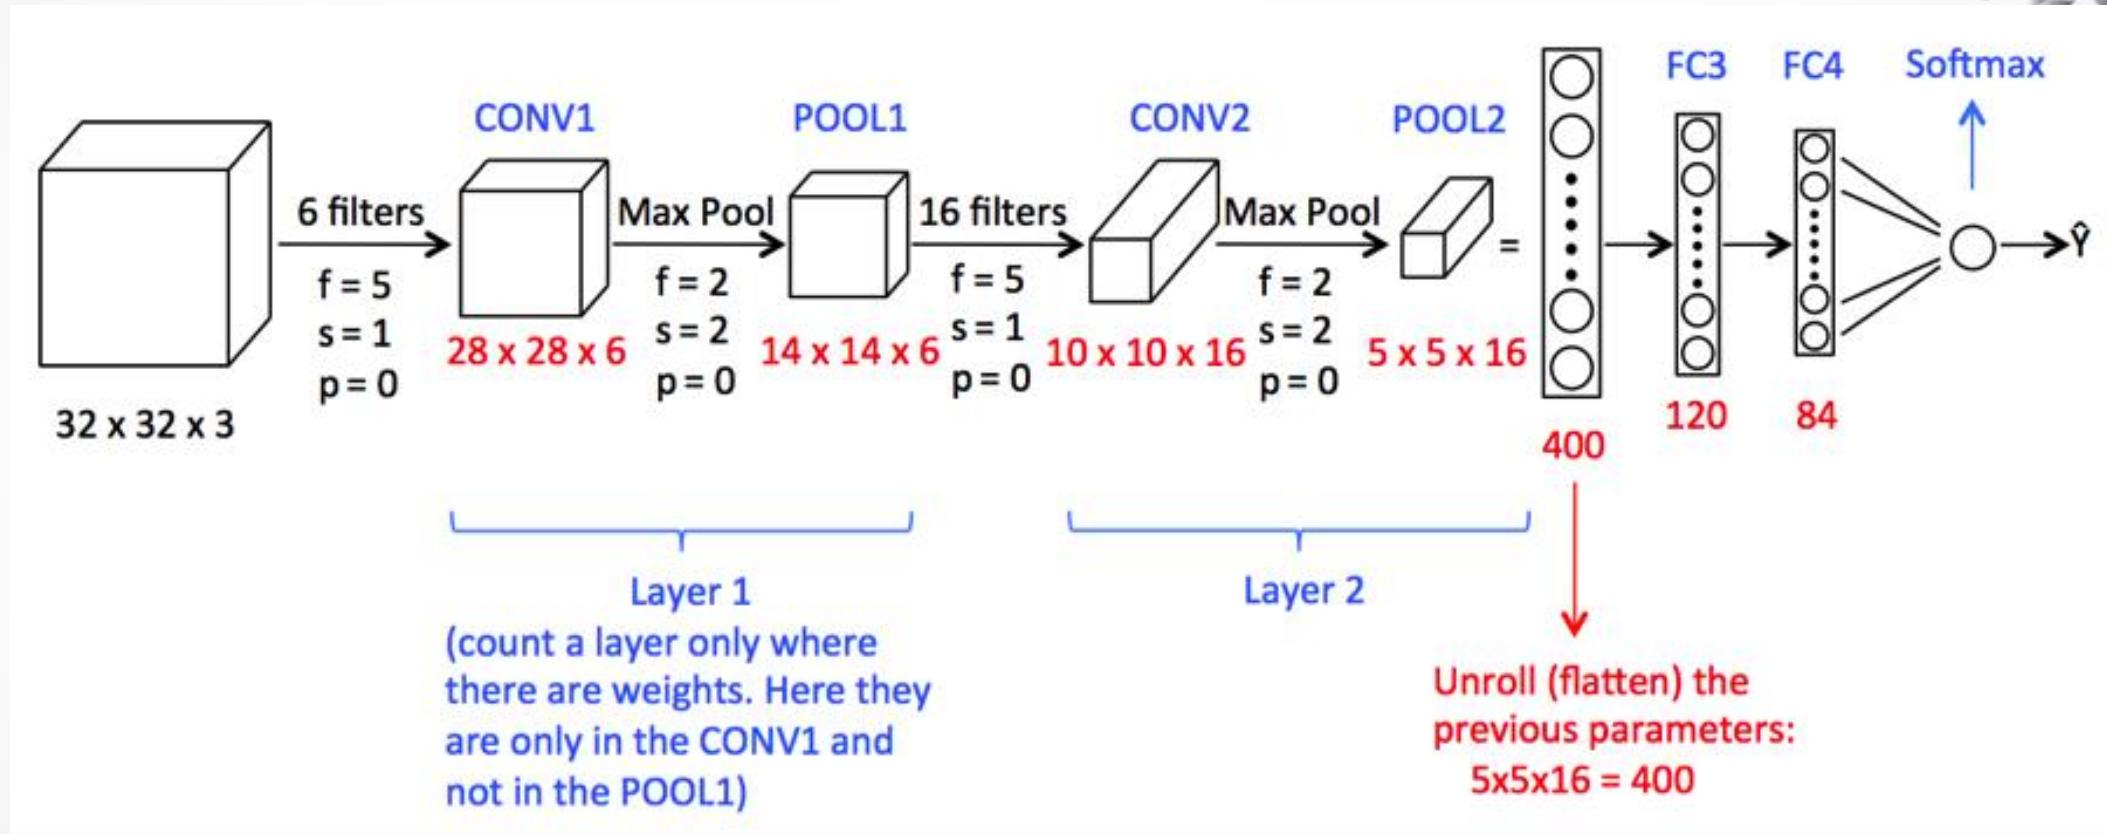


# Convolutional Neural Network



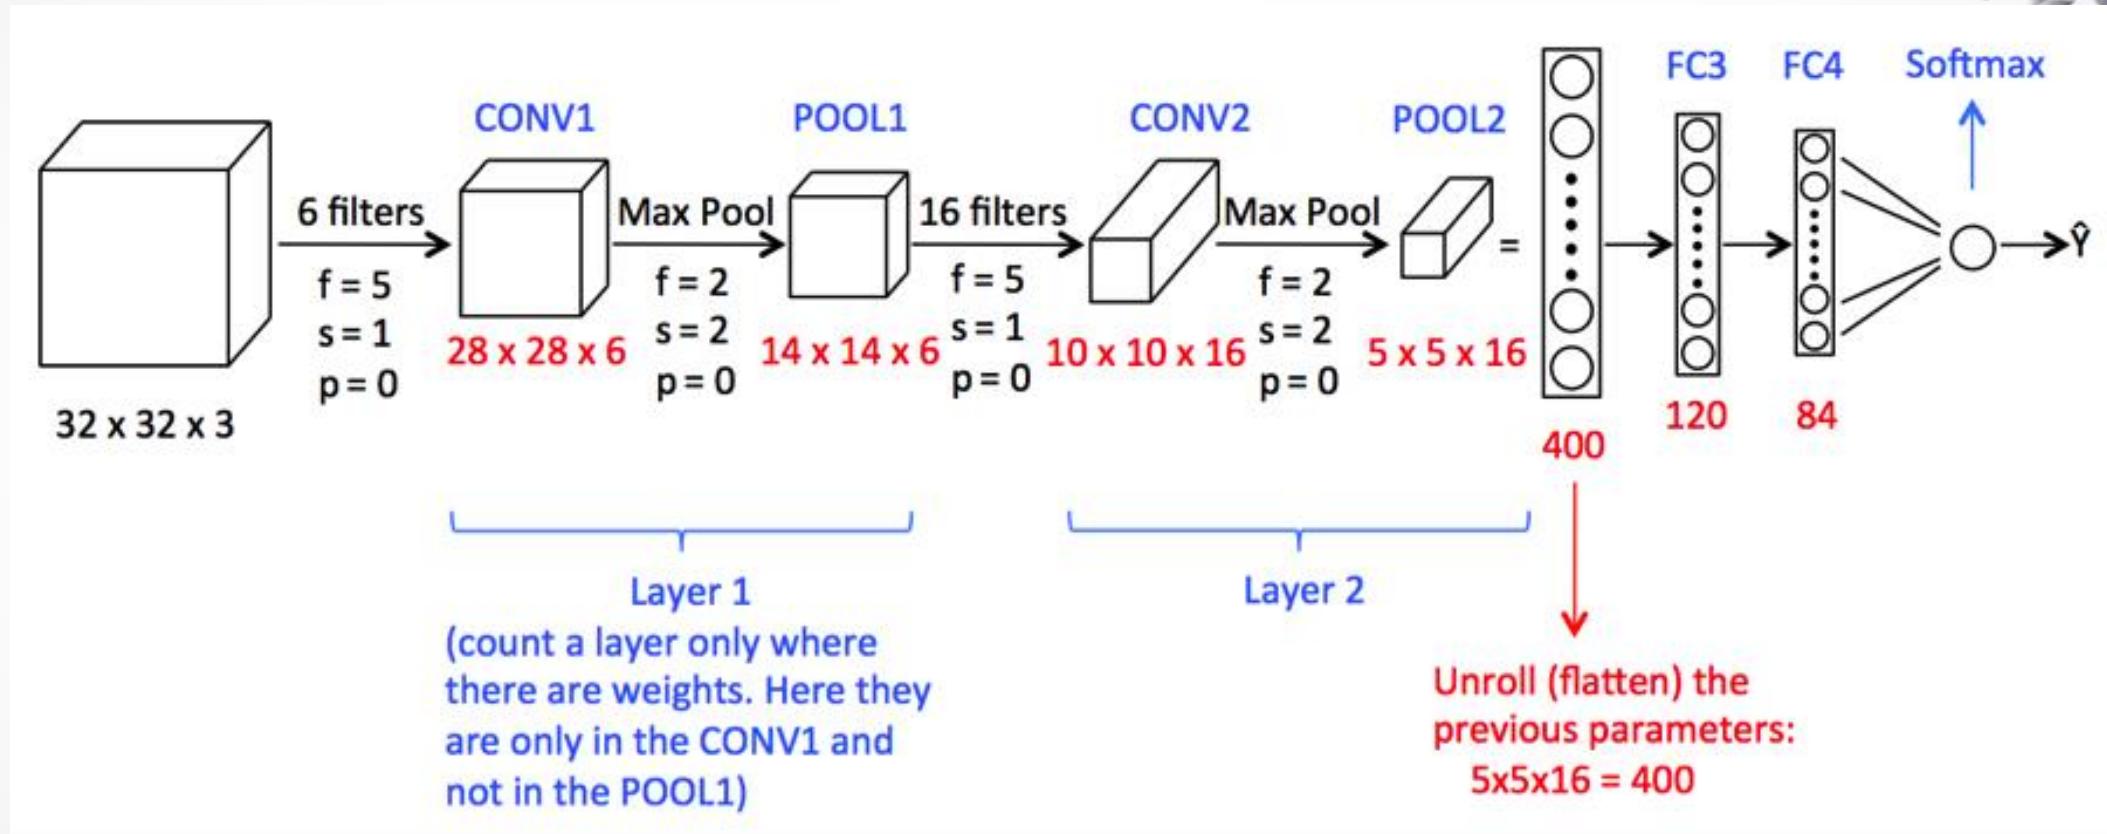
<https://medium.com/machine-learning-bites/deeplearning-series-convolutional-neural-networks-a9c2f2ee1524>

# Convolutional Neural Network



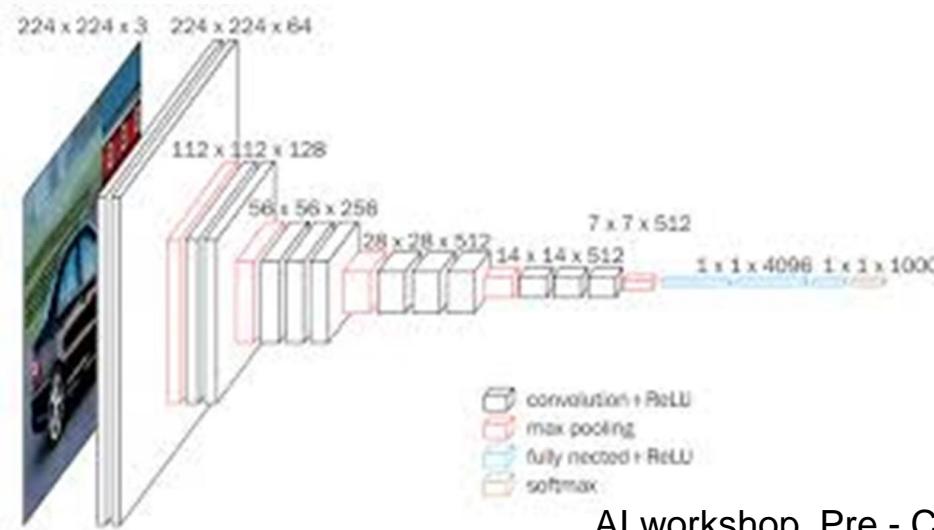
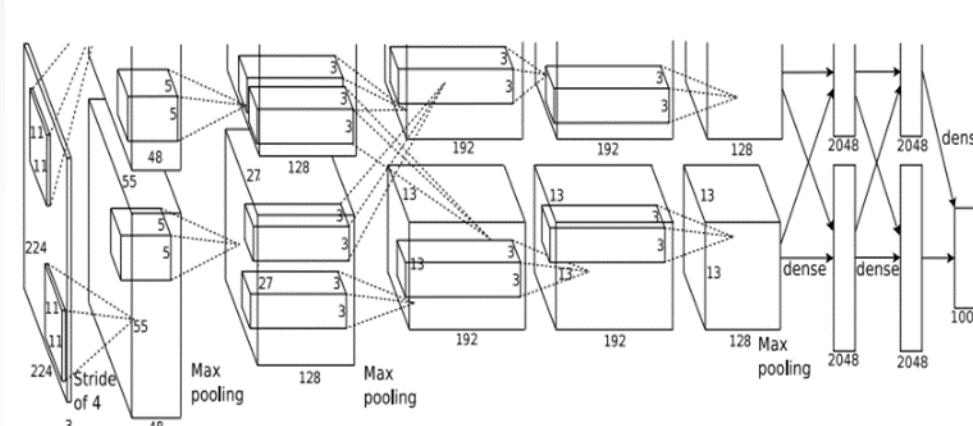
<https://medium.com/machine-learning-bites/deeplearning-series-convolutional-neural-networks-a9c2f2ee1524>

# Convolutional Neural Network

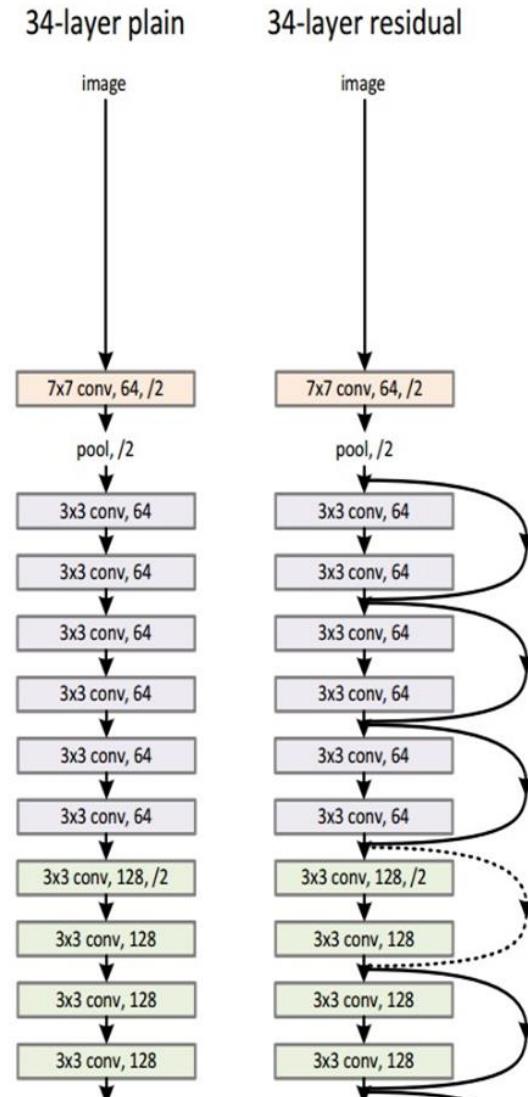
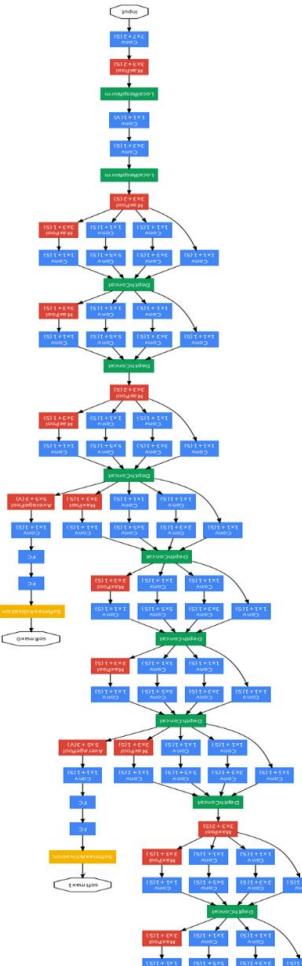


<https://medium.com/machine-learning-bites/deeplearning-series-convolutional-neural-networks-a9c2f2ee1524>

# Convolutional Neural Networks



AI workshop. Pre - Coding Conquest 2019 event



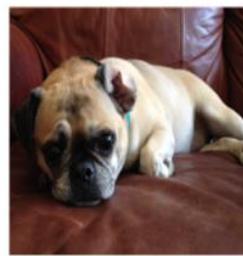
# End-to-End Approach (self learned features)



Preview

[Zeiler and Fergus 2013]

Visualization of VGG-16 by Lane McIntosh. VGG-16 architecture from [Simonyan and Zisserman 2014].

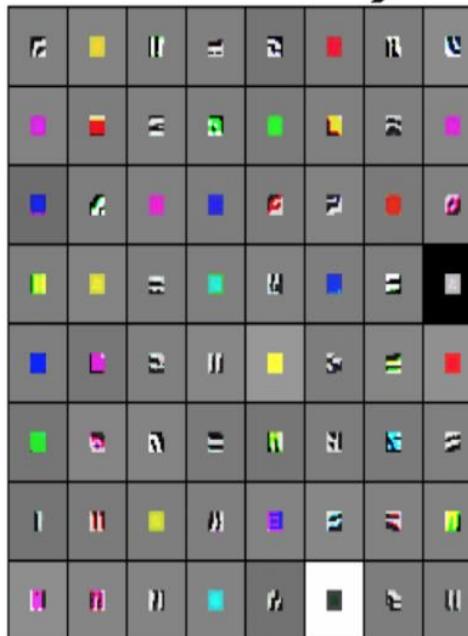


Low-level features

Mid-level features

High-level features

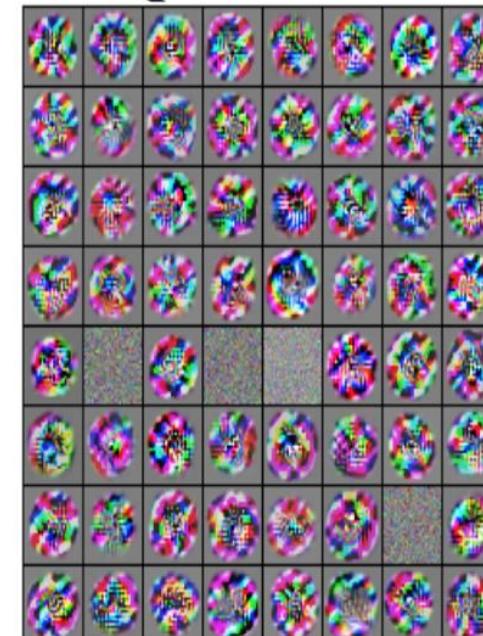
Linearly  
separable  
classifier



VGG-16 Conv1\_1



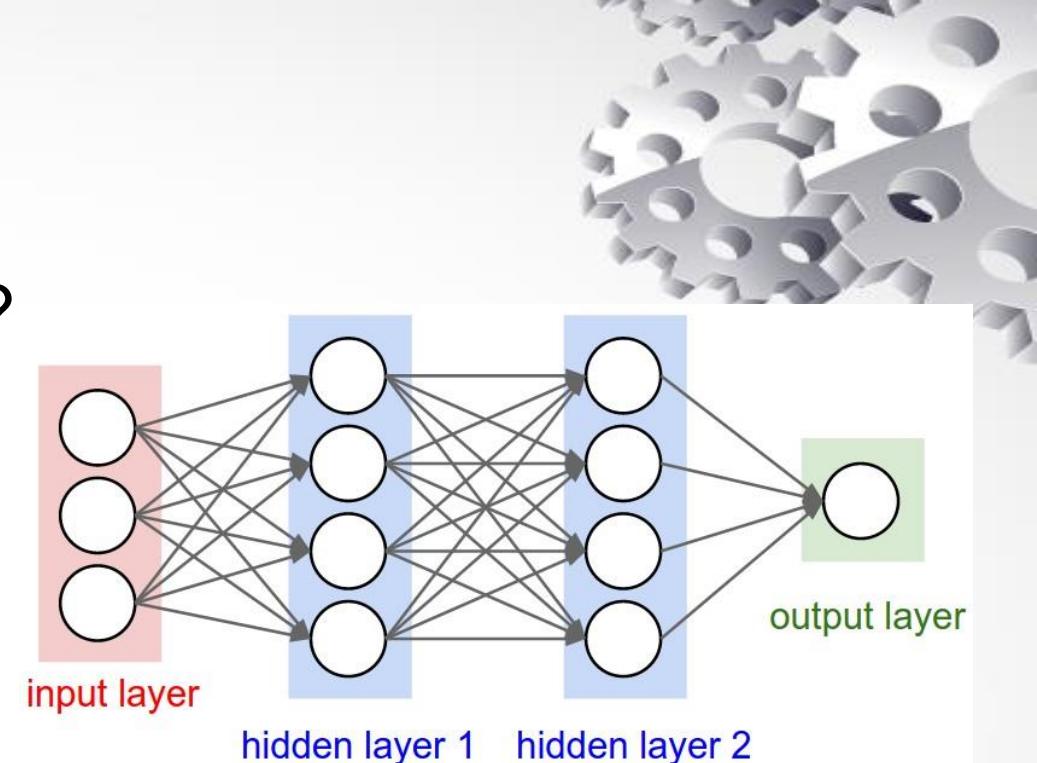
VGG-16 Conv3\_2



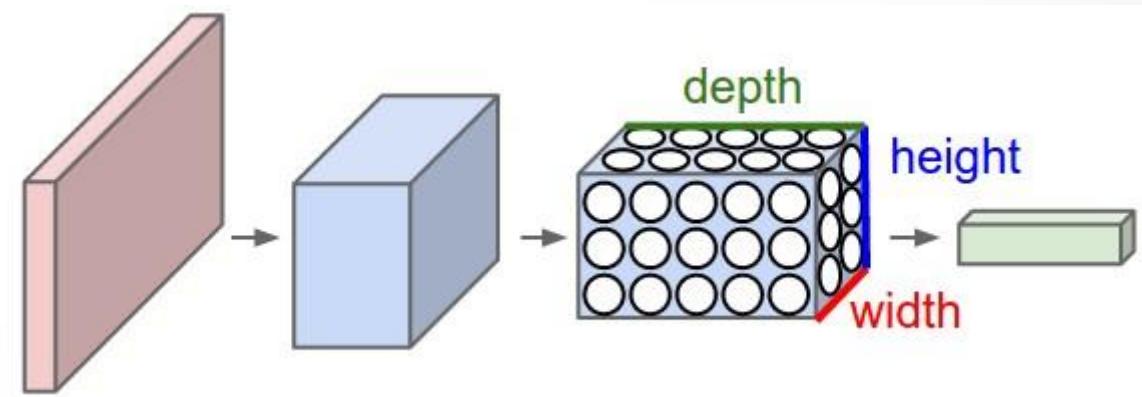
VGG-16 Conv5\_3

# Recap: MLP & CNN

- Should MLP be wide rather than deep?
- Vanishing Gradient Problem



- Convolutional layer
- Pooling layer
- Normalization layer
- Dropout
- Fully connected layer



# Deep Learning for Image Classification Tasks



[LeNet](#)

[AlexNet](#)

[ZFNet](#)

[VGGNet](#)

[SPPNet](#)

[GoogLeNet / Inception-v1](#)

[BN-Inception / Inception-v2](#)

[Inception-v3](#)

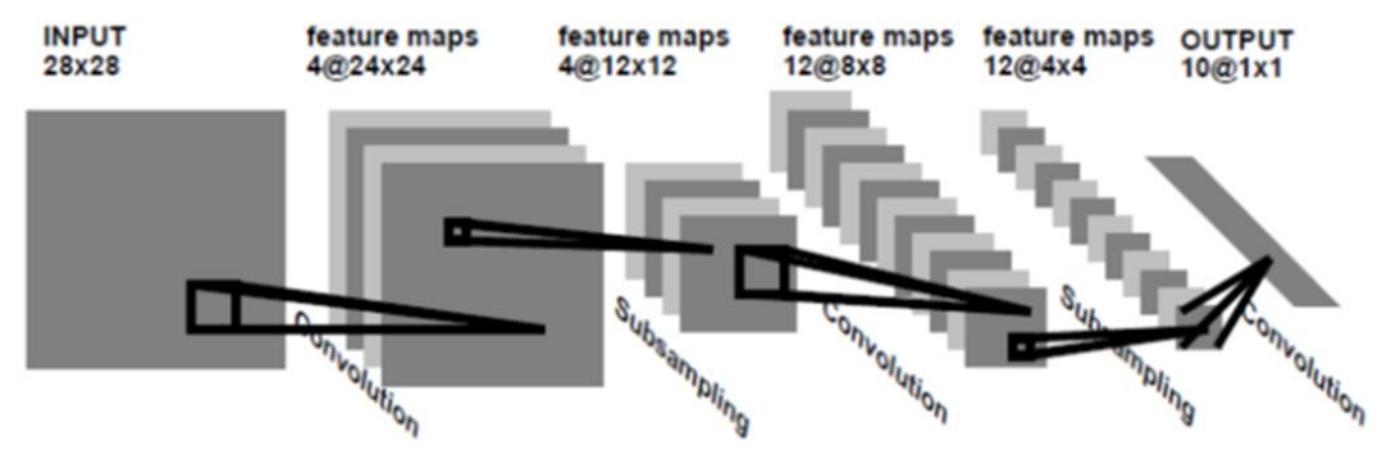
[Inception-v4](#)

[Xception](#)

[MobileNetV1](#)

[ResNet](#)

[DenseNet](#)



# Deep Learning for Image Classification Tasks



[LeNet](#)

[AlexNet](#)

[ZFNet](#)

[VGGNet](#)

[SPPNet](#)

[GoogLeNet / Inception-v1](#)

[BN-Inception / Inception-v2](#)

[Inception-v3](#)

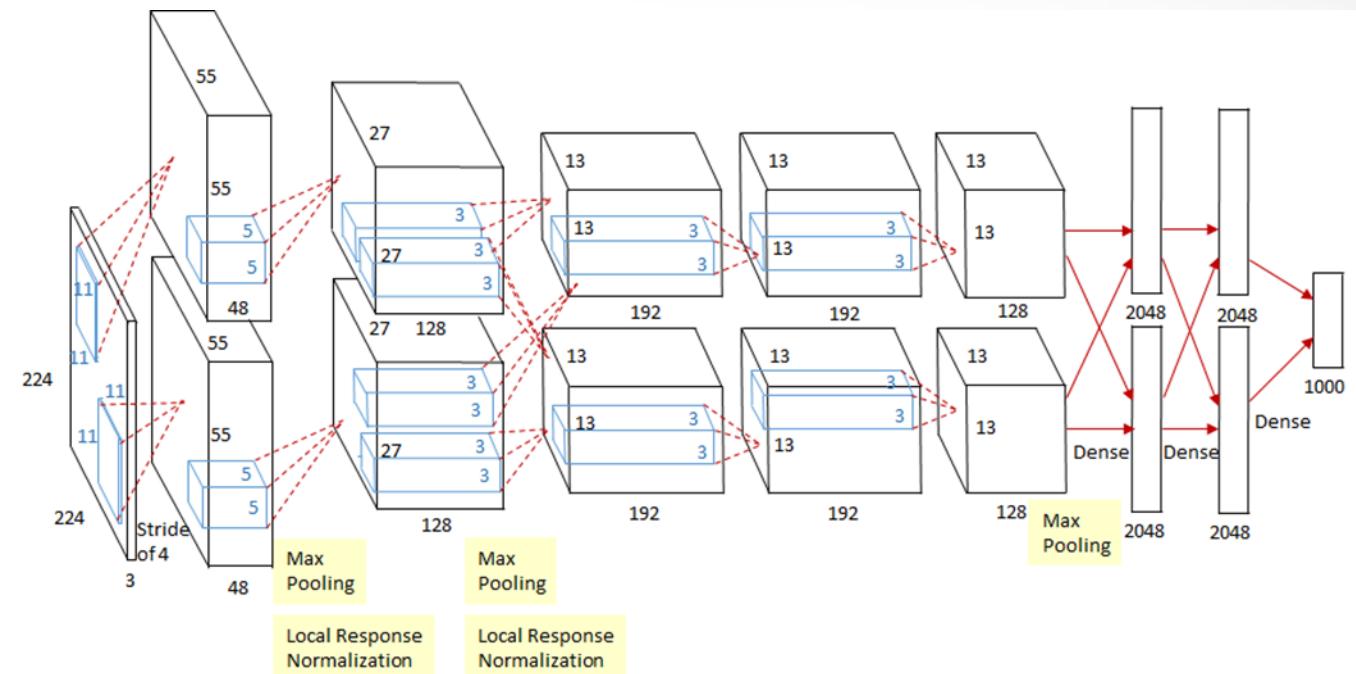
[Inception-v4](#)

[Xception](#)

[MobileNetV1](#)

[ResNet](#)

[DenseNet](#)



# Deep Learning for Image Classification Tasks



[LeNet](#)

[AlexNet](#)

[ZFNet](#)

[VGGNet](#)

[SPPNet](#)

[GoogLeNet / Inception-v1](#)

[BN-Inception / Inception-v2](#)

[Inception-v3](#)

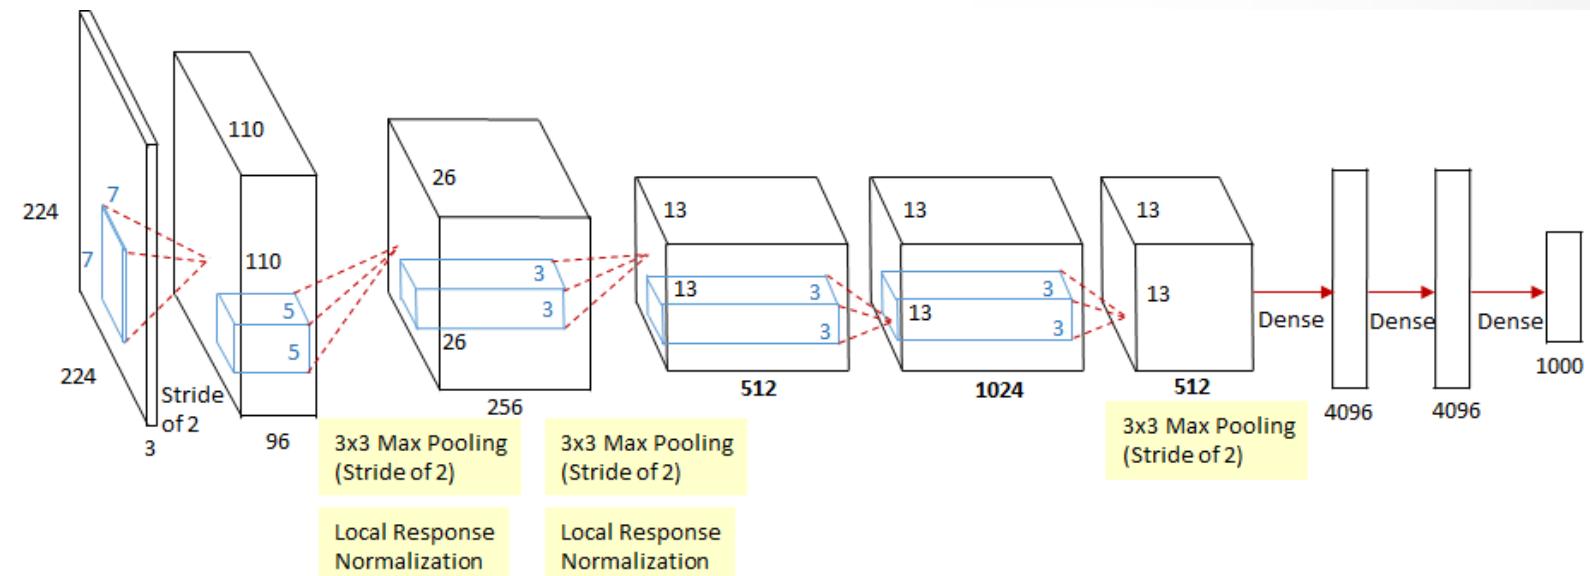
[Inception-v4](#)

[Xception](#)

[MobileNetV1](#)

[ResNet](#)

[DenseNet](#)



# Deep Learning for Image Classification Tasks



[LeNet](#)

[AlexNet](#)

[ZFNet](#)

[VGGNet](#)

[SPPNet](#)

[GoogLeNet / Inception-v1](#)

[BN-Inception / Inception-v2](#)

[Inception-v3](#)

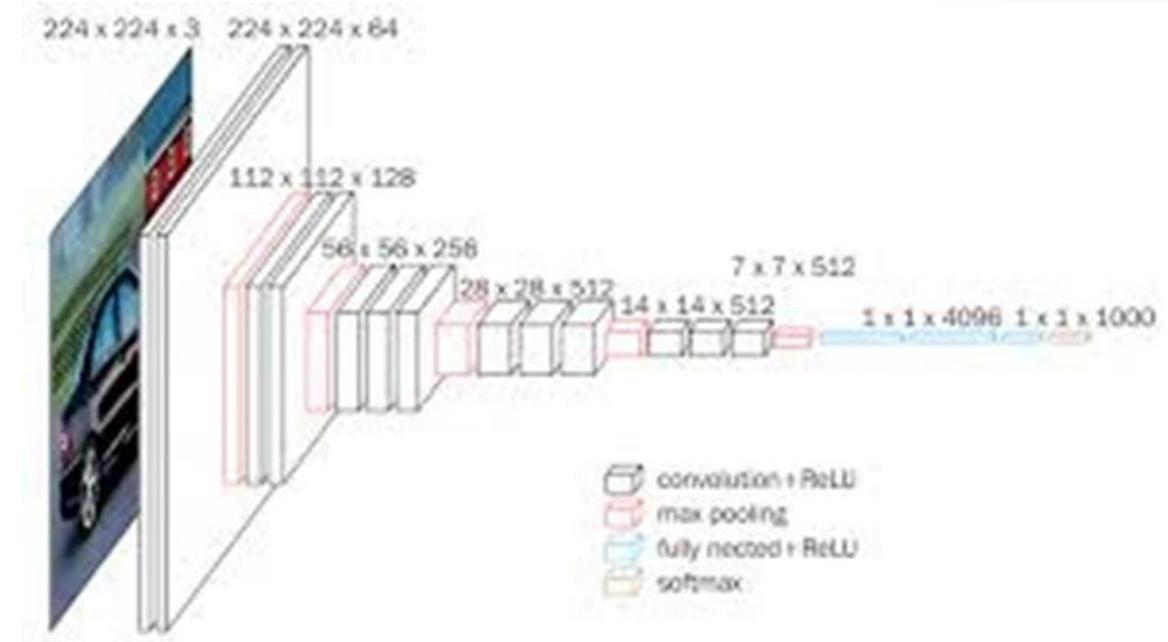
[Inception-v4](#)

[Xception](#)

[MobileNetV1](#)

[ResNet](#)

[DenseNet](#)



# Deep Learning for Image Classification Tasks



[LeNet](#)

[AlexNet](#)

[ZFNet](#)

[VGGNet](#)

[SPPNet](#)

[GoogLeNet / Inception-v1](#)

[BN-Inception / Inception-v2](#)

[Inception-v3](#)

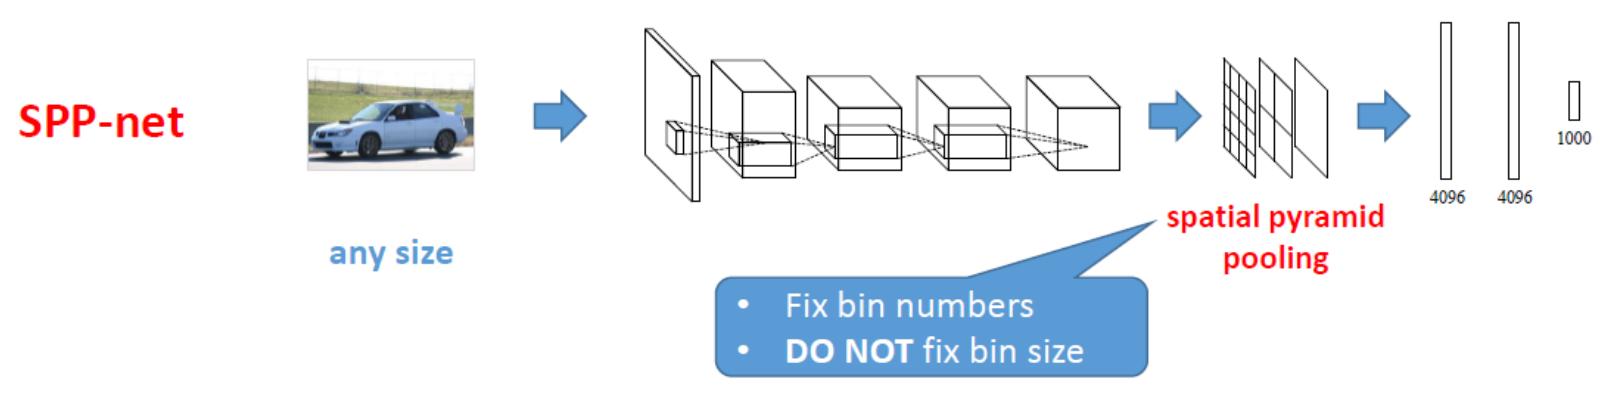
[Inception-v4](#)

[Xception](#)

[MobileNetV1](#)

[ResNet](#)

[DenseNet](#)



# Deep Learning for Image Classification Tasks

LeNet

AlexNet

ZFNet

VGGNet

SPPNet

**GoogLeNet / Inception-v1**

BN-Inception / Inception-v2

**Inception-v3**

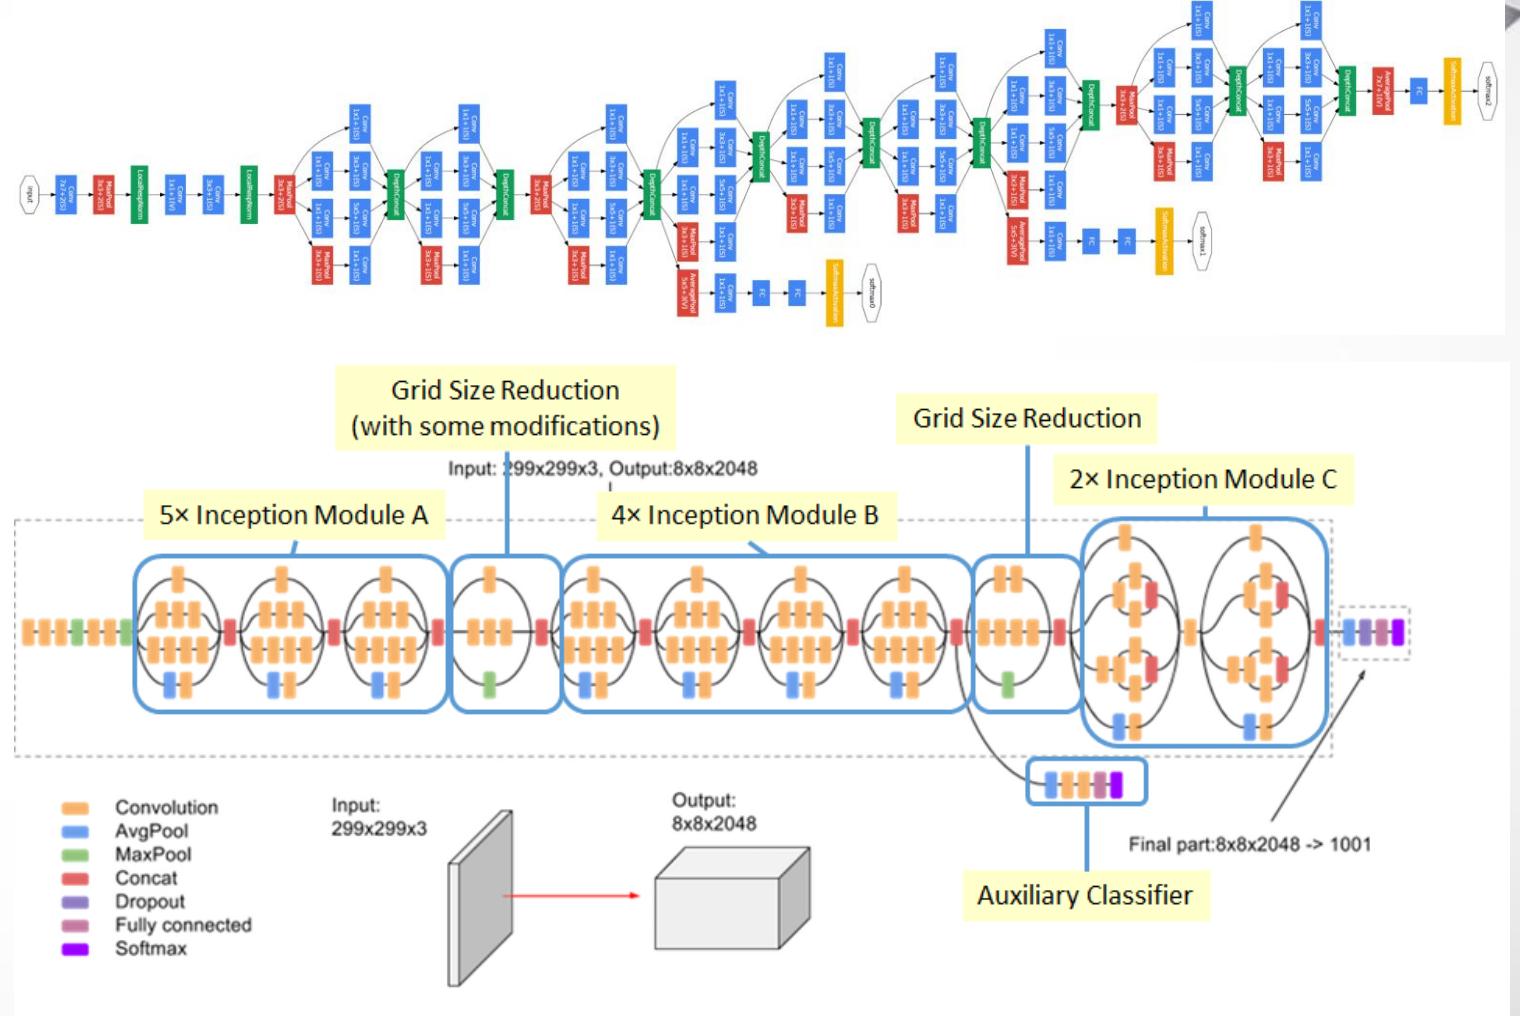
**Inception-v4**

Xception

MobileNetV1

ResNet

DenseNet



# Deep Learning for Image Classification Tasks



[LeNet](#)

[AlexNet](#)

[ZFNet](#)

[VGGNet](#)

[SPPNet](#)

[GoogLeNet / Inception-v1](#)

[BN-Inception / Inception-v2](#)

[Inception-v3](#)

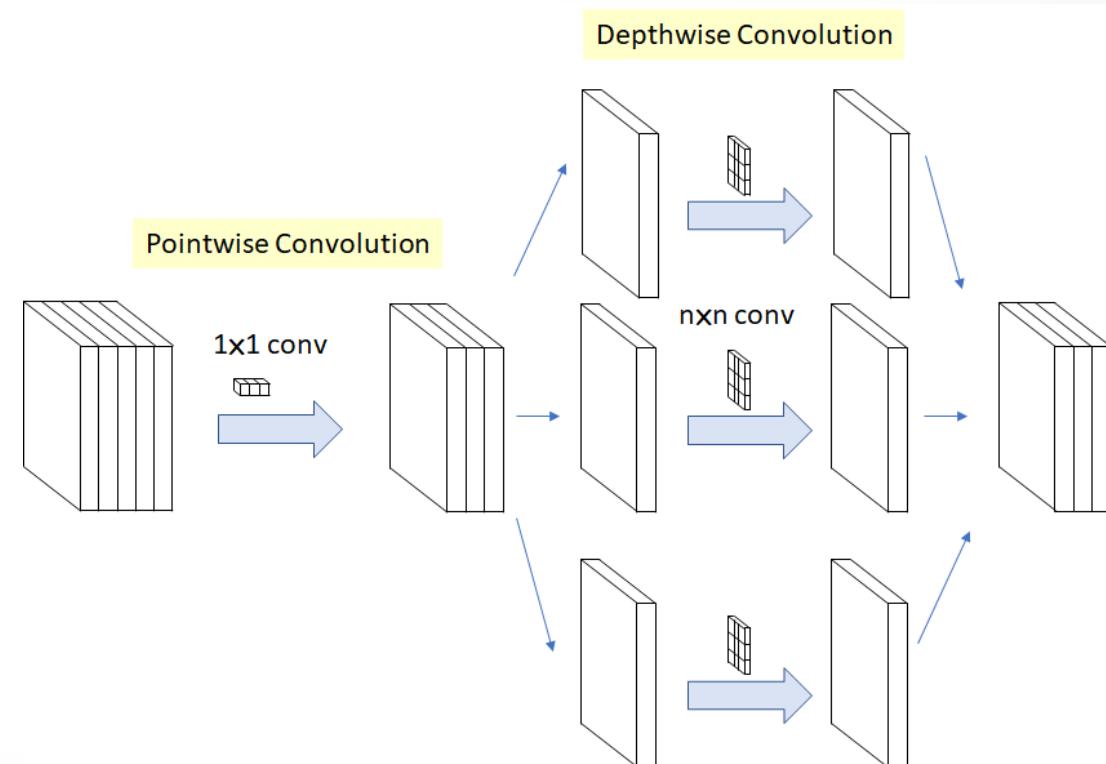
[Inception-v4](#)

[Xception](#)

[MobileNetV1](#)

[ResNet](#)

[DenseNet](#)



# Deep Learning for Image Classification Tasks



[LeNet](#)

[AlexNet](#)

[ZFNet](#)

[VGGNet](#)

[SPPNet](#)

[GoogLeNet / Inception-v1](#)

[BN-Inception / Inception-v2](#)

[Inception-v3](#)

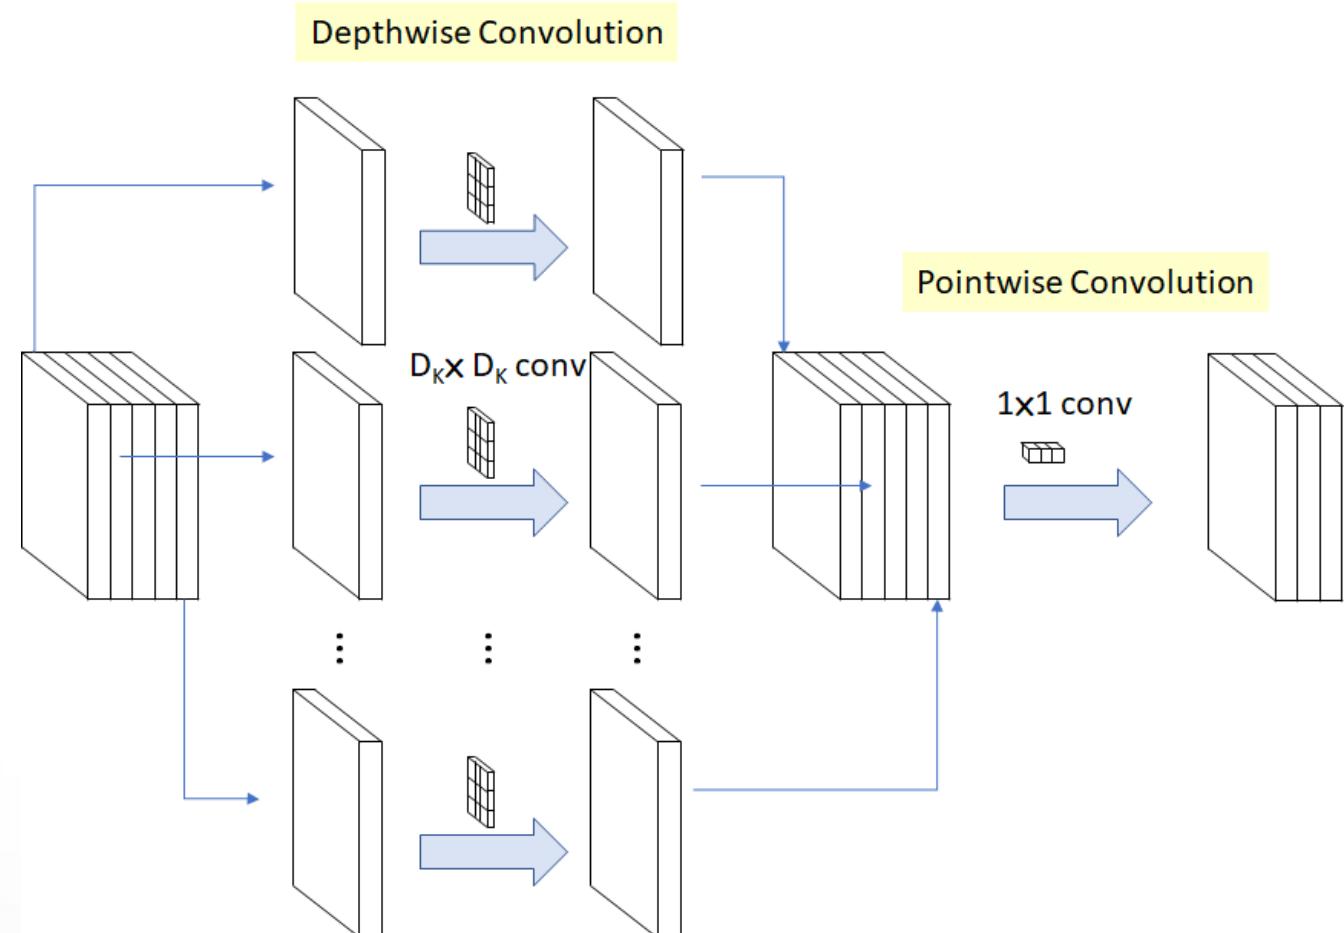
[Inception-v4](#)

[Xception](#)

[MobileNetV1](#)

[ResNet](#)

[DenseNet](#)



# Deep Learning for Image Classification Tasks



LeNet

AlexNet

ZFNet

VGGNet

SPPNet

GoogLeNet / Inception-v1

BN-Inception / Inception-v2

Inception-v3

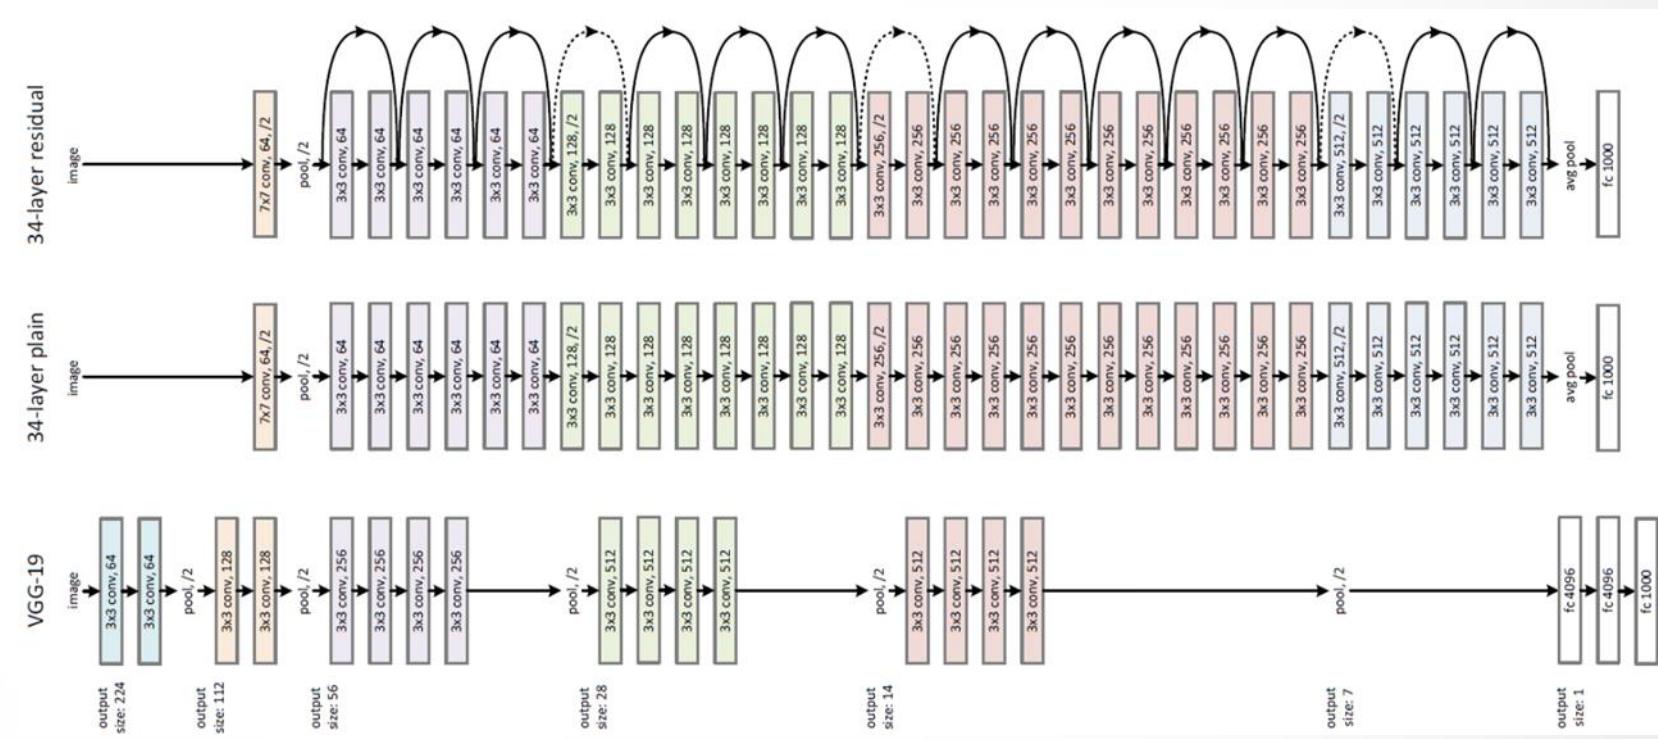
Inception-v4

Xception

MobileNetV1

ResNet

DenseNet



# Deep Learning for Image Classification Tasks



[LeNet](#)

[AlexNet](#)

[ZFNet](#)

[VGGNet](#)

[SPPNet](#)

[GoogLeNet / Inception-v1](#)

[BN-Inception / Inception-v2](#)

[Inception-v3](#)

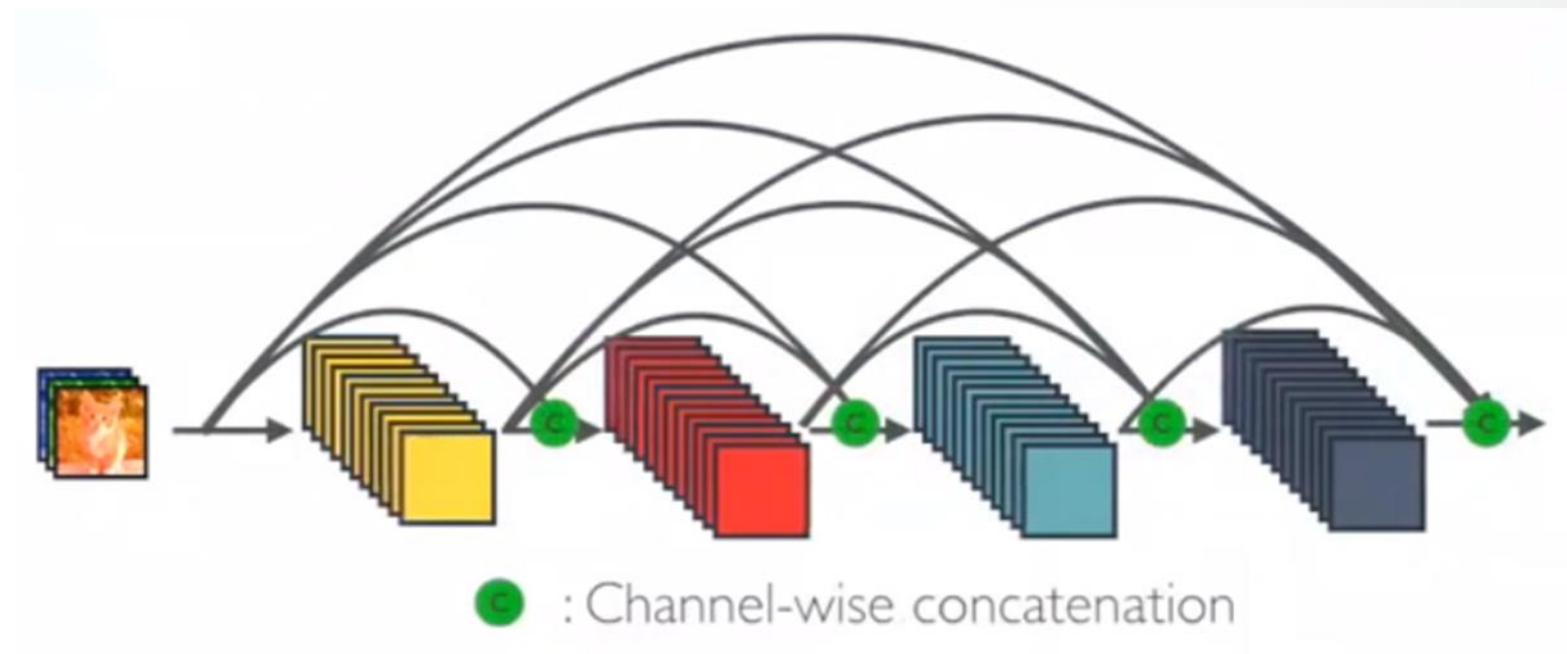
[Inception-v4](#)

[Xception](#)

[MobileNetV1](#)

[ResNet](#)

[DenseNet](#)



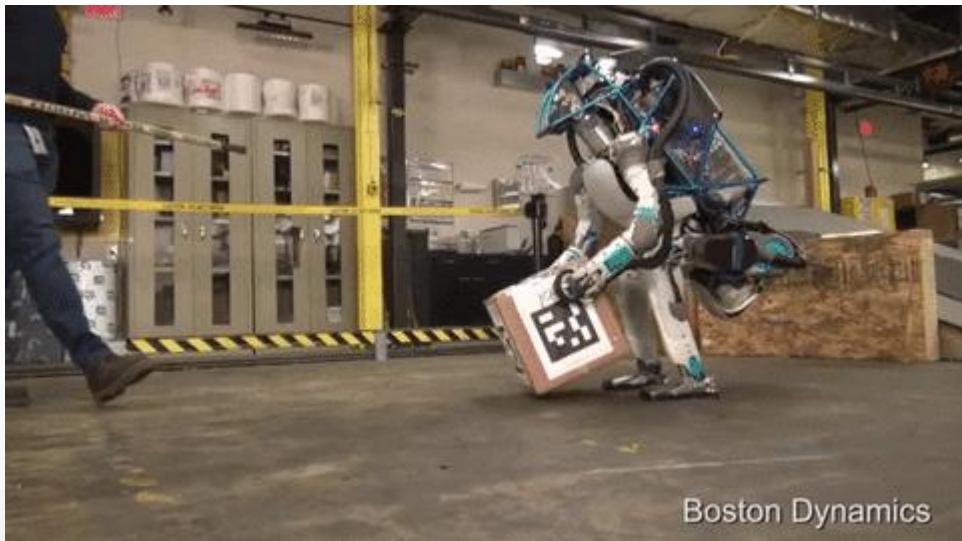
# Some Characteristics of CNN

- A deep network seems to have a better performance than a wide network, given the same amount of nodes.
- Skip connections between layers close to the input and those close to the output could improve training efficiency and performance of the CNN.
- The CNN architecture gets deeper and denser:
  - LeNet5 has 5 layers
  - VGG16 and VGG19 have 16 and 19 layers respectively
  - Residual Networks (ResNets) have more than 100 layers



# The Resurgence of Deep Learning

- Ecosystem
  - The availability of data
  - The progress of hardware, software
  - The progress of ICT infrastructure



Boston Dynamics



Matching of  
Problems & Solutions

# The Resurgence of Deep Learning

- Ecosystem
  - The availability of data
  - The progress of hardware, software
  - The progress of ICT infrastructure



Boston Dynamics



AI workshop. Pre - Coding Conquest 2019 event



## Deep Learning

- Good performance
- End-to-end approach (representation learning)

## Challenges

- Speed and accuracy in open-ended real life problems
- The whole ecosystem is still developing

# Deep Learning Frameworks

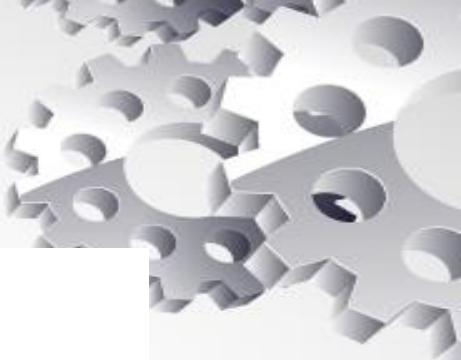


# Hardware Landscape



- There are four major types of technology that can be used to accelerate the training and use of deep neural networks:
- CPUs,
- GPUs,
- Field-programmable gate arrays (FPGAs), and
- Application-specific integrated circuits (ASICs).

# Hardware Landscape



## HARDWARE TECHNOLOGIES USED IN MACHINE LEARNING



# Q & A

