

# Introduction to Reinforcement Learning

## Session 3: RL Examples

Prof. Somnuk Phon-Amnuaisuk (online lecture August 19, 2022)  
ASEAN IVO-project workshop series  
Cover picture: credited <https://www.cyberpunk.net>  
Content: many images from the public domains; many  
slides from GameAI course UTB; CS188 Berkeley and  
AIMA book <http://aima.cs.berkeley.edu/instructors.html>



# meme

A meme (/ˈmiːm/ meem), a neologism coined by Richard Dawkins, is "an idea, behavior, or style that spreads from person to person within a culture". A meme acts as a unit for carrying cultural ideas, symbols, or practices that can be transmitted from one mind to another through writing, speech, gestures, rituals, or other imitable phenomena with a mimicked theme.



# Outline

- Designing of states and actions, examples from OpenAI Gym and Unity Game Engine
- Look at an implementation from Dung-Yi, Chao, Purdue University: Reinforcement Learning on Route Planning through Google map for Self-driving System



# Q Learning

$$Q^{new}(s_t, a_t) \leftarrow \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \underbrace{\left( \underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} - \underbrace{Q(s_t, a_t)}_{\text{old value}} \right)}_{\text{new value (temporal difference target)}}$$

temporal difference

In 2014 Google DeepMind patented an application of Q-learning to deep learning, titled "deep reinforcement learning" or "deep Q-learning" that can play Atari 2600 games at expert human levels.

Initialized

Q-Table		Actions					
		South (0)	North (1)	East (2)	West (3)	Pickup (4)	Dropoff (5)
States	0	0	0	0	0	0	0
	.	.	.	.	.	.	.
	.	.	.	.	.	.	.
	.	.	.	.	.	.	.
	327	0	0	0	0	0	0
.	.	.	.	.	.	.	.
.	.	.	.	.	.	.	.
.	.	.	.	.	.	.	.
499	0	0	0	0	0	0	0

Training

Q-Table		Actions					
		South (0)	North (1)	East (2)	West (3)	Pickup (4)	Dropoff (5)
States	0	0	0	0	0	0	0
	.	.	.	.	.	.	.
	.	.	.	.	.	.	.
	.	.	.	.	.	.	.
	328	-2.30108105	-1.97092096	-2.30357004	-2.20591839	-10.3607344	-8.5583017
.	.	.	.	.	.	.	.
.	.	.	.	.	.	.	.
.	.	.	.	.	.	.	.
499	9.96984239	4.02706992	12.96022777	29	3.32877873	3.38230603	

# Deep Q Learning

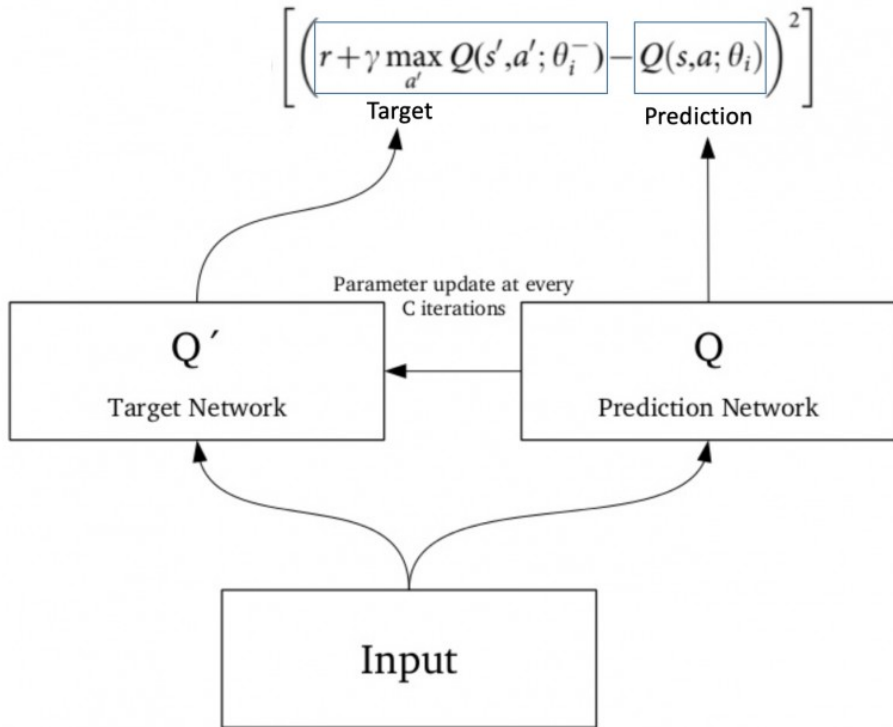
- The DeepMind system used a deep convolutional neural network, with layers of tiled convolutional filters to mimic the effects of receptive fields.
- **Reinforcement learning is unstable** or divergent when a nonlinear function approximator such as a neural network is used to represent  $Q$ . This instability comes from the correlations present in the sequence of observations, the fact that small updates to  $Q$  may significantly change the policy and the data distribution, and the correlations between  $Q$  and the target values.
- The technique used **experience replay**, a biologically inspired mechanism that uses a random sample of prior actions instead of the most recent action to proceed. This removes correlations in the observation sequence and smooths changes in the data distribution. Iterative updates adjust  $Q$  towards target values that are only periodically updated, further reducing correlations with the target.

# Deep Q Learning

$$Q^{new}(s_t, a_t) \leftarrow \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \left( \underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} - \underbrace{Q(s_t, a_t)}_{\text{old value}} \right)$$

$$\underbrace{\left( \underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} - \underbrace{Q(s_t, a_t)}_{\text{old value}} \right)}_{\text{new value (temporal difference target)}}$$

1. Select an action from possible Q-values actions, select using the epsilon-greedy policy.
2. Perform this action **a** in a state **s** and move to a new state **s'** to receive a reward **r**.
3. Record this transition in our replay buffer as **<s,a,r,s'>**
4. After C iterations, sample data randomly from the replay-buffer and train the ANN using fix target.
5. Update the target Q network
6. Repeat




# OpenAI Gym

```
import gym

env = gym.make("CartPole-v1")
observation, info = env.reset(seed=42, return_info=True)

for _ in range(1000):
    action = env.action_space.sample()
    observation, reward, done, info = env.step(action)

    if done:
        observation, info = env.reset(return_info=True)
env.close()
```




# OpenAI Gym

```
# CART POLE  
env = gym.make('CartPole-v0')  
print(env.observation_space.low)  
print(env.observation_space.high)
```

```
[-4.8000002e+00 -3.4028235e+38 -4.1887903e-01 -3.4028235e+38]  
[4.8000002e+00 3.4028235e+38 4.1887903e-01 3.4028235e+38]
```

```
print(len(env.observation_space.low, env.action_space.n)
```


```
4 2
```






# Balance Ball Example

- Game-play
- States
- Actions
- Rewards



# Electrical Vehicle Route Planning on Google Map with RL


- Traditional Google API provides the best route in term of traveling duration.
- In this work, the authors find the best route with minimum energy cost and acceptable duration.
- Use elevation to approximate energy consumption among possible routes.



# Problem Formulation


Other assumptions:

- Each action is based on the energy required to travel with 750m displacement on the grid map.
- To compute the energy required between position A and position 1, for example, we use the duration and distance between position A and position 1 to calculate the average velocity  $V$ . Combine  $V$  with the elevation, we can get the angle  $\theta$  of the road and consider the height of the road as linear increasing or decreasing.
- Because we don't take regenerative braking into account in our experiment, we treat the downhill road flat.




# States and Actions

- State: latitude and longitude
- Actions: N, E, S, W



(a)



(b)

# Rewards

- The fundamental concept of defining the reward is based on the energy consumption in one stride from the current position to the next position.
- For example, from A to B shown in the figure. The energy is calculated by the method provided previous section. We then divide the energy by 10000 and times -1. In order to minimize the number of total steps during training, we add -0.1 to each transition if the next position is reachable. In other words, the reward  $r$  for taking any reachable step will be
$$r = -0.1 - (\text{energy consumption} / 10000).$$



# Q & A